

From Agriculture to Services: Development without Industrialization

Tancredi Rapone[†]

UAB · Barcelona School of Economics

September 17, 2026

Abstract Does industrialization leave a lasting legacy after manufacturing declines? I argue that it does. Large manufacturing firms create demand for managerial and technical capabilities that remain valuable as economies transition toward services. I study this mechanism using cross-country evidence, U.S. county data from 1850 to 2000, and a quantitative growth model in which firm scale raises the return to developing reusable solutions to recurring organizational problems that can later diffuse across firms and sectors. In the data, countries with larger historical manufacturing peaks develop service sectors more intensive in high-skill producer services and grow faster even after manufacturing has declined. Across counties in the United States, a 10pp higher manufacturing peak employment share driven by national growth in industries where large scale coordination was technologically important raises the share of high-skill producer services in total private-service employment by 2.8 percentage points in 2000, or 45% of the cross-county average. The calibrated model shows that this mechanism goes a long way toward explaining cross-country differences in service composition, skill accumulation, and growth, and identifies firm scale and the diffusion of organizational knowledge as central forces. Despite manufacturing’s pervasive association with favorable development outcomes, subsidizing manufacturing itself is ineffective; persistent gains instead come from policies that expand the supply of skills while creating demand for them by enabling firms to grow and formalize production.

Keywords: structural transformation, industrialization, producer services, firm scale, organizational capital, service-led growth.

JEL Classification: O14, O41, O33, O15, J24.

[†] First and foremost, I am deeply grateful to my advisors, Raül Santaaulàlia-Llopis and André Gröger, for their continued support, patience, and guidance throughout this project. A special thanks also goes to my PhD evaluation committee members for their valuable insight and continued guidance: Jordi Caballe, Marti Mestieri, Pau Roldan, and Luis Rojas. I also thank Tasso Adamopoulos, Tristany Armangué, Gabriela Araujo, Sophie Brochet, Alejandro Gobierno, Jacob Hess, Ivan La Franca, Luca Looser, Enoch Ntsiful, Nicolas Oviedo, Michael Peters, Luis Felipe Sáenz, Alberto Salio, Etienne Wasmer, and seminar participants at the BSE PhD Jamboree, Bellaterra Macro Club, EWMES 2025 and LAGV 2026 for their valuable comments and suggestions. All errors are my own. For comments, suggestions, or critiques, you can reach me at tancredi.rapone@icloud.com.

1 Introduction

For most of economic history, development meant industrialization. Labor moved out of agriculture, manufacturing expanded, and services became dominant only after a substantial industrial phase (Clark, 2007; Herrendorf et al., 2014). That sequence is no longer typical. Many developing economies now move directly from agriculture into services, reaching lower manufacturing-employment peaks and deindustrializing at lower income levels than earlier industrializers (Rodrik, 2016). This raises a natural question: does a service economy that emerges directly from agriculture inherit the same productive capabilities as one that develops after a substantial manufacturing phase?

I argue that the answer depends on whether the development path passes through production at sufficient scale to induce organizational investment. Historically, industrialization concentrated production in large establishments, where coordinating large numbers of workers and transactions generated recurring organizational problems (Pollard, 1965; Chandler, 1990; Hounshell, 1984; Yates, 1993). These problems—including managing inventories, controlling quality, collecting information, and evaluating investment decisions—are not specific to manufacturing, but the payoff to developing systematic methods for handling them rises when they recur at large scale. Large industrial firms therefore employed managers, engineers, accountants, clerks, lawyers, and other specialized workers to organize and coordinate production; I refer to these activities as producer services and to the skilled workers performing them as high-skill producer-service workers (HSPS henceforth). The historical record reflects this pattern: many major innovations in modern management originated inside large industrial firms—including return-on-investment accounting at DuPont, the multidivisional form at General Motors, and statistical quality control—and subsequently became tools for organizing firms well beyond manufacturing (Mol and Birkinshaw, 2014; McKenna, 2006). The central argument of this paper is therefore that large-scale manufacturing builds managerial and technical capabilities that remain useful as the economy moves into services.

I provide evidence for this mechanism in two steps. First, across countries since 1990, larger historical manufacturing peaks predict a service transition tilted toward high-skill producer services rather than trade and other consumer-facing services, and faster subsequent GDP-per-capita growth even after manufacturing has declined. Second, using U.S. county data from 1850 to 2000, I show that counties with larger manufacturing peaks accumulated more high-skill producer-service employment within manufacturing and later developed service sectors more specialized in high-skill producer-service industries. A natural concern is that manufacturing simply made counties larger and more urban, independently generating subsequent demand for high-skill producer services. To address this concern, I draw on the historical literature distinguishing high-throughput industries, in which large-scale plants and intensive coordination were technologically important, from low-throughput industries that could operate at smaller scale (Chandler, 1990). While both types of manufacturing predict population growth and urbanization, only high-throughput manufacturing is associated with subsequent growth in HSPS services. To sharpen identification, I construct shift-share instruments that interact counties' 1880 manufacturing-industry composition with subsequent leave-county-out national industry growth. In the preferred IV specification, a ten-percentage-point increase in peak manufacturing employment induced by high-throughput national growth raises the high-skill producer-service share of private-service employment in 2000 by 2.83 percentage points, or about 45% of the cross-county average. By contrast, an instrument pooling national growth across all manufacturing industries also strongly predicts the manufacturing peak but produces no statistically detectable long-run effect on service composition.

The empirical evidence links large-scale manufacturing to the composition of services, but it cannot quantify how much this mechanism matters for aggregate development, how it propagates through skill accumulation and growth, or what would happen under alternative development paths. I therefore formalize the mechanism in a quantitative growth model. Firms repeatedly face organizational problems which can either be handled case by case or formalized by paying a fixed skilled-labor cost to develop a reusable procedure. Because larger firms encounter the same problems more often, scale raises the return to formalization and the demand for skilled producer-service workers. Newly formalized procedures are initially proprietary but can subsequently enter a public stock from which other firms and sectors adapt them at lower cost, creating a knowledge externality. The model generates a transition from agriculture to services in which the long-run growth rate of the economy is fully endogenous and determined by the limiting share of skilled workers. Manufacturing can permanently raise this growth rate when production at scale induces firms to formalize organizational tasks, creating persistent demand for skilled labor.¹

I calibrate the model to the U.S. transition from 1850 to 2019, initializing the economy in 1790, and discipline structural change, firm scale, skills, producer services, fertility, and productivity growth. As an untargeted validation of the organizational mechanism, the model’s annual additions to public organizational knowledge have a correlation of 0.75 with contemporaneous Library of Congress publications on business organization over 1906–1977, significantly higher than for 11 placebo knowledge categories. I then conduct a model-based cross-country accounting exercise, comparing the model-generated relationships between manufacturing peaks and observable outcomes with the corresponding gradients in the data. Varying the relative pull of manufacturing via a preference shifter generates 98 percent of the observed cross-country gradient in the HSPS share of services, 32 percent of the gradient in tertiary education, 17 percent of the GDP-per-capita gradient, and 40 percent of the gradient in average GDP-per-capita growth over 2010–2019. Through the lens of the model, countries that pass through a more pronounced manufacturing phase therefore emerge not only richer and more skill intensive, but also continue to grow faster after manufacturing has declined. Mechanism-removal exercises identify firm scale and knowledge diffusion as central: equalizing the sectoral entry environment reduces the HSPS gradient to 12.7 percent of the baseline, while eliminating knowledge diffusion cuts it by more than half. These results carry directly into policy implications. In an economy calibrated to a representative lower-middle-income country, removing distortions to firm growth raises scale and income but not sufficiently to induce firms to formalize production and demand HSPS labor. Temporary support for organizational formalization, by contrast, raises persistent skill demand and growth, with effects that are much larger when distortions preventing firm growth are removed. Education subsidies temporarily expand skill supply but leave little persistent change in the skilled share or organizational knowledge without reforms which remove barriers to firm growth. The policy implication is therefore one of complementarity: sustained growth requires expanding the supply of skills together with reforms that generate demand for them by allowing firms to grow and formalize production.

¹ A secondary contribution concerns the determination of long-run growth. In standard semi-endogenous growth models such as [Jones \(1995\)](#), diminishing returns to the existing stock of knowledge make sustained productivity growth depend on exogenous population growth. Because the model follows a centuries-long development transition in which fertility and productivity growth evolve jointly ([Galor, 2005](#)), I combine a Jones-inspired productivity technology with the endogenous fertility mechanism of [Galor and Weil \(2000\)](#). This generates an equilibrium in which fertility converges to replacement and population to a finite endogenous level, while productivity growth remains positive and is determined by the economy’s limiting skilled share. Section 4 states the result formally.

The paper is most closely related to recent work connecting firm scale, skills, and the organization of production. [Engbom et al. \(2025\)](#) also build directly on Chandler’s account of the modern enterprise, emphasizing how the supply of educated white-collar workers enables firms to delegate, reorganize, and grow. My mechanism runs in the complementary direction: large-scale production creates demand for managerial and technical workers who solve recurring organizational problems systematically. I then study the long-run consequence of this process: whether the organizational routines and specialized expertise developed inside manufacturing persist and support the emergence of producer services after manufacturing declines. The paper is also closely related to [Peters et al. \(2026\)](#), who study fast-growing developing economies that have expanded services without undergoing broad industrialization and ask whether productivity growth in consumer services can sustain such a development path. I instead ask whether a substantial manufacturing phase changes the service economy that subsequently emerges, and show that large-scale manufacturing promotes high-skill producer services with persistent consequences for skill demand and growth.

More broadly, the paper connects three literatures. First, it contributes to work on structural transformation and service-led development. This literature explains movements across agriculture, manufacturing, and services through income effects, relative prices, productivity growth, and trade, and studies both cross-country differences in industrialization and the rise of skill-intensive and service-led development ([Kongsamut et al., 2001](#); [Ngai and Pissarides, 2007](#); [Herrendorf et al., 2014](#); [Boppart, 2014](#); [Comin et al., 2021](#); [Rodrik, 2016](#); [Huneus and Rogerson, 2024](#); [Fujiwara and Matsuyama, 2024](#); [Buera and Kaboski, 2012](#); [Buera et al., 2022](#); [Fan et al., 2023](#); [Chen et al., 2023](#); [Sposi et al., 2026](#)). I shift attention from what determines the manufacturing phase to how that phase shapes the composition of the service economy that follows. Second, the paper relates to recent work on the persistent effects of industrial policy. [Juhász \(2018\)](#), [Incoronato and Lattanzio \(2025\)](#), and [Choi and Levchenko \(2025\)](#) show that temporary industrial interventions can generate lasting changes in technology adoption, industrial structure, and firm performance. I provide a mechanism through which an industrial expansion can leave effects beyond manufacturing: large-scale production induces firms to formalize organizational tasks, and the resulting capabilities can subsequently diffuse into services. Finally, I connect historical accounts of the emergence and diffusion of modern management ([Pollard, 1965](#); [Chandler, 1977, 1990](#); [Yates, 1993](#); [McKenna, 2006](#); [Mol and Birkinshaw, 2014](#)) to models of organization capital ([Prescott and Visscher, 1980](#); [Atkeson and Kehoe, 2005](#)) by making the creation and diffusion of reusable organizational routines endogenous to firm scale.

The remainder of the paper proceeds as follows. Section 2 presents the cross-country evidence. Section 3 studies the mechanism by comparing the development path of U.S. counties. Section 4 develops the model, Section 5 describes its calibration and historical validation, and Section 6 reports the quantitative results. Section 7 studies policy implications and Section 8 concludes.

2 Cross-Country Facts

This section establishes two facts connecting industrialization to the service economy that follows. First, countries that experienced larger manufacturing-employment peaks developed service economies more concentrated in high-skill producer services. Second, larger historical manufacturing peaks are associated with faster GDP-per-capita growth, even decades after the manufacturing peak.




Figure 1: Manufacturing peaks and the subsequent service economy. The figure includes all 51 ETD countries observed annually over 1990–2018. Countries are grouped by the median historical peak manufacturing-employment share, calculated over their full available GGDC 10-sector–ETD history, with ETD taking precedence in overlapping years. The above-median group includes the country at the median and contains 26 countries; the below-median group contains 25. Panel A reports employment in high-skill producer-service (HSPS) industries—business services, finance, and real estate—as a share of total employment. Panel B reports employment in generic market services, defined as the ETD trade and other-services sectors, using the same denominator. Trade includes accommodation and food services. Transport and government services are excluded from both numerators. Each country’s share is rebased to zero in 1990 before taking unweighted group means. Shaded regions show the mean plus or minus 1.96 cross-country standard errors.

The two exercises use complementary country samples. To study the composition of services, I use the GGDC/UNU–WIDER Economic Transformation Database (ETD), which reports annual employment in 12 broad industries for 51 economies from 1990 through 2018. I define high-skill producer-service (HSPS) industries as business services, finance, and real estate, while generic market services comprise trade and other market services. Manufacturing peaks are computed by combining the GGDC 10-Sector Database with ILOSTAT. For each country, I define the manufacturing peak as the maximum employment share over its longest available harmonized series, which begins between 1950 and 1990 depending on the country.² GDP per capita is from the World Bank and Gapminder. Appendix B describes the data, concordances, and sample construction in detail.

FACT 1: *Larger historical manufacturing peaks are associated with greater expansion of high-skill producer-service employment.*

I begin by comparing service-employment trajectories across economies with different industrial histories. I split the 51-country ETD sample according to whether each country’s historical manufacturing peak lies above or below the cross-country median and compare the evolution of employment in HSPS industries—business services, finance, and real estate—and in generic market services—trade and other market services—as shares of total employment. To remove initial level differences, I subtract each country’s own 1990 employment share before computing unweighted group means. Because most countries

² The GGDC 10-Sector database provides data going back to the late 1940s or early 1950s depending on the country. ILOSTAT provides data over 1990–2018. The majority of countries reach their manufacturing peak in the early 1990s.




Figure 2: Growth in service economies by peak manufacturing exposure. The vertical axis is average annual growth in PPP real GDP-per-capita over 2010–2019; the horizontal axis is the average service-employment share over the same period. Countries are divided at the median historical manufacturing peak. Data sources are ILO-STAT and the World Bank via Gapminder.

reached their manufacturing peak by the early 1990s, the exercise compares how service economies subsequently evolved across countries with different depths of prior industrialization. Figure 1 shows substantially greater HSPS expansion in economies with higher manufacturing peaks. Between 1990 and 2018, the HSPS share of total employment rises by 5.75 percentage points in the higher-peak group, compared with 2.77 points in the lower-peak group. The difference is 2.98 percentage points and is statistically significant in a two-sided Welch test ($p = 0.001$). Countries in the lower peak group instead expanded in generic market services which are primarily consumer facing and include retail and trade. The share of generic market services in total employment rises by 9.59 percentage points in the lower-peak group, compared with 5.90 points in the higher-peak group. This approximately 3.7-percentage-point difference is also statistically significant ($p = 0.016$). Thus, the two groups differ in the service industries that absorb employment: higher-peak economies experience greater HSPS expansion, while lower-peak economies experience greater expansion of trade and other services.³ Appendix B.2 shows that this relationship survives controls for contemporaneous GDP-per-capita growth, income-group and region fixed effects, and the overall size of the service sector, and is driven primarily by business services, followed by finance and real estate.⁴

³ Other services does not include government services.

⁴ In the ETD data, business services are the service activities most closely connected to firms' information and specialized professional needs—principally information and communication and professional, scientific, and technical activities.

FACT 2: *Countries with larger manufacturing peaks grow faster as they transition into services.*

The rise of services is often associated with slower aggregate productivity growth. In the standard Baumol cost-disease logic, labor reallocates toward sectors with relatively slow productivity growth, reducing aggregate growth as the service share rises (Baumol, 1967; Duernecker and Sanchez-Martinez, 2023; Duernecker et al., 2024). Fact 1 suggests an additional dimension: service economies may differ in their internal composition depending on the development path through which they emerged. Figure 2 plots average annual PPP GDP-per-capita growth over 2010–2019 against the average service-employment share over the same period in a sample of 173 countries. Countries are divided according to whether their historical manufacturing peak lies above or below the cross-country median. Service employment is negatively associated with growth in both groups; however, conditional on the service-employment share, countries with above-median manufacturing peaks systematically grow faster.

The relationship survives more demanding comparisons. Appendix Table 26 shows that the manufacturing peak remains positively associated with subsequent GDP-per-capita growth after including region and income-group fixed effects and controlling for institutional quality. Moreover, the association remains positive, with a slightly larger point estimate, among countries where services account for more than half of employment. These facts are consistent with a persistent legacy of industrialization rather than an effect operating only while manufacturing itself remains large.

Economies with larger manufacturing peaks develop service economies concentrated in high-skill producer service sectors and grow faster. The cross-country evidence does not, however, identify the source of this persistence. Section 3 turns to the U.S. historical transition to test whether it reflects the scale and organizational demands of manufacturing and the managerial and technical capabilities it creates.

3 U.S. County Evidence

U.S. counties provide a setting in which the proposed mechanism can be observed over the full industrial transition. I first trace the buildup of HSPS capabilities within manufacturing around county-specific industrial peaks and their later emergence in specialized service industries. I then ask whether establishment scale and manufacturing HSPS employment account for the long-run relationship between manufacturing exposure and composition of services. Finally, I exploit counties' 1880 industry composition and subsequent national industry growth to test whether this legacy is concentrated in high-throughput manufacturing rather than manufacturing expansion more generally. Together, these exercises distinguish the scale-dependent organizational channel from broader consequences of industrialization such as population growth and urbanization.

3.1 Data and measurement

To study the mechanism over the full U.S. structural transformation, I combine county tabulations from the Census of Manufactures with full-count population censuses and postwar administrative employment data, harmonizing counties to fixed 1850 boundaries. The analysis distinguishes three objects: the depth of industrialization, measured by the peak manufacturing employment share; managerial and technical capability within manufacturing, measured by HSPS employment in manufacturing relative to

total classified employment; and the subsequent expansion of specialized HSPS services, measured relative to the same employment base. For the long-run cross-sectional analysis, I additionally measure the HSPS share within private-service employment to isolate the association between industrialization and the composition of the service sector rather than its overall size. The peak manufacturing employment share is manufacturing employment obtained from the Census of Manufactures relative to total classified employment obtained from the Census of Population. Let L_{ct}^C denote total classified employment, defined as the number of employed individuals whose industries can be classified as primary, manufacturing, or services.⁵ I define county c 's manufacturing peak as

$$M_c^{\text{peak}} = \max_{t \in \{1880, 1900, 1920, 1930, 1940\}} \left\{ \frac{L_{ct}^M}{L_{ct}^C} \right\}, \quad (1)$$

where L_{ct}^M is average manufacturing wage earners and L_{ct}^C is total classified employment. I let P_c denote the census year attaining this maximum. The peak window begins in 1880, before the national expansion of manufacturing and at the first major improvement in the quality of U.S. census enumeration.⁶ The years 1890 and 1910 do not feature in the data because most of the 1890 population census was lost and the required county manufacturing tabulations are unavailable in 1910. Manufacturing scale is log manufacturing workers per establishment in the same peak year P_c , also measured in the Census of Manufactures.

To measure the organizational capabilities employed inside manufacturing, I identify workers performing the functions required to organize production at scale and construct

$$MHSPS_{ct} = 100 \times \frac{\text{HSPS-occupation workers in manufacturing}_{ct}}{L_{ct}^C}. \quad (2)$$

The manufacturing HSPS classification contains four functional groups: accounting and record keeping; purchasing, logistics, and coordination; personnel and labor relations; and technical and scientific work, including engineers, draftsmen, statisticians, surveyors, technicians, and inspectors. These occupations capture the information-processing, coordination, and technical functions that are necessary to operate large manufacturing enterprises (Chandler, 1977, 1990).⁷ The service-side measure asks whether employment subsequently shifts toward industries that supply these capabilities on the market. I define

$$SHSPS_{ct} = 100 \times \frac{\text{employment in HSPS service industries}_{ct}}{L_{ct}^C}. \quad (3)$$

In the modern data, HSPS services comprise Information, Professional, Scientific and Technical Services, and Management of Companies—NAICS sectors 51, 54, and 55. These industries provide the closest service-sector counterpart to the organizational capabilities emphasized by Chandler (1977, 1990). NAICS 54 includes accounting, legal, engineering, scientific, technical, and consulting services that supply externally many functions historically organized within large manufacturers; NAICS 55 captures the management and coordination of enterprises; and NAICS 51 captures the production and transmission of

⁵ More specifically, I define total classified employment as employment in primary industries (IND1950 codes 105, 116, and 126), manufacturing (IND1950 codes 306–499), and services (IND1950 codes 506–946). Mining, construction, unclassified employment, and industries outside these categories are excluded.

⁶ Before 1880, federal censuses were conducted by U.S. marshals; the 1880 Census was the first administered by a professional census office and enumerator force.

⁷ The core manufacturing HSPS set contains OCC1950 codes 0, 35, 41–49, 72, 83, 92, 95, 96, 200, 201, 280, 310, 321, 322, 341, 342, and 533. Appendix A reports the full labels and construction.

information. I exclude NAICS 56 because it combines administrative support with many routine activities not closely tied to the mechanism, and I exclude the broad NAICS 52–53 aggregates because they contain substantial financial intermediation, insurance, property operation which are primarily consumer facing.⁸ Appendix A reports the historical IND1950 and SIC concordances used to carry this functional definition backward in time. The county-industry panel within services uses the Population Census through 1950, County Business Patterns in 1960 and 1970, and the harmonized CBP panel of [Eckert et al. \(2020\)](#) thereafter; the manufacturing-HSPS measure uses the Population Census throughout because it requires occupation-by-industry information.⁹ For the 2000 measure of service composition, I use the same HSPS industry definition but divide employment in NAICS 51, 54, and 55 by total private-service employment from the BLS Quarterly Census of Employment and Wages (QCEW), which directly reports the county-level private service-providing aggregate ([U.S. Bureau of Labor Statistics, 2004](#)).¹⁰

3.2 Manufacturing HSPS and the later rise of specialized services

I first trace how managerial and technical employment within manufacturing and employment in specialized HSPS service industries evolve around county-specific manufacturing peaks. Let M_c^{peak} denote the manufacturing employment share at county c 's peak. I estimate by OLS the following specification separately for $MHSPS_{ct}$ and $SHSPS_{ct}$:

$$Y_{ct} = \alpha_c + \gamma_{s(c),t} + \lambda_{t-P_c} + \sum_{k \neq -40} \beta_k \left[\mathbf{1}\{t - P_c = k\} \times \frac{M_c^{\text{peak}}}{0.10} \right] + X'_c \delta Trend_t + \varepsilon_{ct}. \quad (4)$$

The specification absorbs county, state-by-calendar-year, and event-time fixed effects. The vector X_c contains 1850 log population, urbanization, colleges, school enrollment, and log enslaved population, together with 1870 shares of process, machinery, and traditional manufacturing. Each covariate is interacted with a linear calendar-time trend. Event time -40 is omitted. Because repeated county observations may be spatially correlated across neighboring local economies, the confidence intervals use a spatial covariance estimator allowing arbitrary covariance across all observations from counties whose centroids are within 100 kilometers.¹¹ This is a continuous-treatment event-study specification: each β_k measures the differential evolution of the outcome associated with a 10-percentage-point larger manufacturing employment share at the peak. The sample contains 1,116 fixed-geography counties peaking in 1900, 1920, or 1940 for which the peak treatment and a common forty-year pre-period are observed.

Figure 3 establishes a clear sequence. Counties with higher manufacturing peaks accumulate more HSPS employment within manufacturing before the peak, and this differential continues to widen for several decades afterward. The manufacturing coefficient averages 0.063 percentage points over event

⁸ The cross-country ETD analysis necessarily uses a broader sectoral classification. Its business-services category includes information and communication and professional and technical activities, while finance and real estate are observed only as broad sectors.

⁹ Throughout, L_{ct}^C denotes the same broad classified-employment concept. For each outcome and data vintage, it is reconstructed from the same source as the corresponding numerator, avoiding comparisons between workplace-based postwar employment and residence-based Census employment.

¹⁰ Appendix A.5 describes the treatment of suppressed QCEW cells and reports corresponding results using alternative denominators, the harmonized CBP data, and a 2010 endpoint.

¹¹ Clustering standard errors at the state level leaves results unchanged. I cluster spatially because this is the appropriate choice given industrialization historically exploited geographic features, such as the fall-line, coal mines, lakes and rivers, which transcend state boundaries.




Figure 3: Manufacturing exposure at the peak and the capability sequence. The sample contains 1,116 fixed-geography counties peaking in 1900, 1920, or 1940. Each point is the coefficient associated with a 10-percentage-point larger manufacturing employment share at the county’s peak; event time -40 is omitted. The upper outcome is manufacturing HSPS occupations divided by classified county employment, as defined in equation (2). The lower outcome is employment in HSPS service industries—the historical analogues of NAICS 51, 54, and 55—divided by the broad classified-employment base defined in equation (3). The lower-panel numerator and denominator are constructed from the same employment universe in each vintage: harmonized Census records through 1950, imputed early CBP in 1960 and 1970, and harmonized EFSY CBP thereafter. In the postwar years, both are therefore workplace-based employment counts rather than a workplace numerator divided by employed county residents. The regressions include county, state-by-calendar-year, and event-time fixed effects and baseline-covariate linear trends. Bars are 95-percent confidence intervals using a 100-kilometer spatial covariance estimator. Because the panel ends in 2000, the longest post-peak horizons are supplied by the earlier peak cohorts.

times -30 through zero and 0.254 points over $+10$ through $+50$ ($p < 0.001$ in both periods), before declining thereafter. HSPS services increase later. Their average differential is -0.0004 percentage points over event times -30 through zero ($p = 0.978$) and -0.026 points over $+10$ through $+50$ ($p = 0.439$). It then rises to 0.394 points over $+70$ through $+100$ ($p < 0.001$) in the same period when HSPS employment within manufacturing declines. Appendix A.5.2 shows that this pattern is not a generic consequence of population growth or urbanization. I distinguish high-throughput manufacturing industries, which the historical literature identifies as operating at large scale and requiring intensive managerial and technical coordination, from the remaining manufacturing industries. Peaks in both groups predict subsequent population growth and urbanization, but only high-throughput peaks predict a buildup of HSPS intensity within manufacturing and, later, greater HSPS service employment.

3.3 Manufacturing scale, HSPS employment, and persistent service composition

I next ask whether the long-run association between manufacturing peaks and service composition is accounted for by the two margins implied by the mechanism: establishment scale and HSPS employment within manufacturing. Larger establishments create greater returns to formalizing organizational tasks, while manufacturing HSPS employment measures the managerial and technical functions developed inside manufacturing. I contrast these margins with a broader local-development account in which manufacturing primarily operates through higher wages, urbanization, or agglomeration.

3.3.1 Channel accounting

Let $Y_{c,2000}$ denote the QCEW HSPS share of private-service employment described above. I estimate

$$Y_{c,2000} = \beta M_c^{\text{peak}} + \theta S_c^{\text{peak}} + \phi H_c^{\text{peak}} + A_c^{\text{peak}'} \kappa + X'_{c,1850} \Gamma + \delta_{s(c)} + \lambda_{P_c} + \varepsilon_c, \quad (5)$$

where S_c^{peak} is log manufacturing workers per establishment and H_c^{peak} is manufacturing HSPS employment divided by total classified county employment, the same underlying fraction as in the event study. The vector $A_c^{\text{peak}'}$ contains average manufacturing wages, the urban population share, and log county population, all observed at the manufacturing peak. Every specification includes state and peak-cohort fixed effects and five 1850 controls: log population, urbanization, school enrollment, colleges, and log enslaved population.

Column 1 establishes the baseline long-run relationship on the common sample: a manufacturing peak larger by 10 percentage points predicts a 0.48-percentage-point higher HSPS share of private-service employment in 2000, approximately 8 percent of the 6.36-percent outcome mean. Adding establishment scale reduces the peak coefficient by 50 percent, while adding manufacturing HSPS reduces it by 37 percent. Including both reduces the coefficient from 0.048 to 0.015, an attenuation of 68 percent, and renders it statistically indistinguishable from zero. Manufacturing HSPS remains quantitatively large and precisely estimated in the joint specification: a one-percentage-point increase in manufacturing HSPS relative to classified employment predicts a 1.93-percentage-point increase in the HSPS share of private services. Establishment scale also remains positive and is significant at the 10-percent level.

Column 5 asks whether similar attenuation arises from broader correlates of local development. Urbanization and population at the peak are positively associated with subsequent HSPS specialization, while manufacturing wages are not. Jointly including the three variables attenuates the peak coefficient by 64 percent, but the remaining coefficient stays positive and significant at the 10-percent level. These controls therefore absorb part of the long-run relationship, but they do not eliminate it. Because all of these variables may themselves respond to industrialization, I interpret Table 1 as descriptive channel accounting rather than causal mediation.

3.3.2 National industry shifts

The preceding results suggest that the service-sector legacy of manufacturing is tied to the scale and organizational demands of industrialization. However, industrialization, establishment scale, and HSPS employment are jointly determined, and counties with a pre-existing comparative advantage in managerial or technical activity could exhibit similar development patterns encompassing both the type of manufacturing they develop and the later composition of services. To assuage this concern, I exploit

Table 1: Manufacturing Exposure, Scale, and Persistent HSPS Services

	Baseline	+ Scale	+ Mfg. HSPS	+ Both	+ Alternative channels
Peak manufacturing employment share	0.048*** (0.010)	0.024** (0.012)	0.030*** (0.011)	0.015 (0.012)	0.017* (0.010)
Log manufacturing establishment scale at peak		0.0060*** (0.0022)		0.0041* (0.0023)	
Manufacturing HSPS / classified employment at peak			2.112*** (0.529)	1.927*** (0.535)	
Average manufacturing wage at peak, per \$100					0.0008 (0.0006)
Urban population share at peak					0.0183** (0.0072)
Log county population at peak					0.0188*** (0.0023)
Attenuation of peak coefficient (%)	0.0	50.3	37.4	68.3	63.9
State fixed effects	Yes	Yes	Yes	Yes	Yes
Peak-cohort fixed effects	Yes	Yes	Yes	Yes	Yes
1850 historical controls	Yes	Yes	Yes	Yes	Yes
Observations	1,487	1,487	1,487	1,487	1,487
R^2	0.235	0.239	0.246	0.247	0.291

Notes: The dependent variable is public numeric QCEW private employment in NAICS 51, 54, and 55 divided by the disclosed private service-providing aggregate less published private-utility employment in 2000. The aggregate denominator retains confidential non-HSPS service employment. Peak manufacturing employment is Census of Manufactures average wage earners divided by total classified county employment and is defined as the maximum of this ratio over 1880, 1900, 1920, 1930, and 1940. Manufacturing HSPS is employment in the core accounting, coordination, personnel, technical, and scientific occupations listed in the text, conditional on employment in manufacturing, divided by total classified county employment. Establishment scale, manufacturing HSPS, and the alternative-channel variables are measured in the county's manufacturing-peak year. Missing peak-year manufacturing wages are replaced with the corresponding peak-cohort median and the specification includes a missing-wage indicator. Every column uses the same 1,487 counties and includes state and peak-cohort fixed effects and 1850 controls for log population, urbanization, school enrollment, colleges, and log enslaved population. Conley standard errors with a 100-kilometer cutoff are in parentheses. The attenuation exercise is descriptive channel accounting. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

counties' differential exposure to national manufacturing-industry growth over 1880–1900 by interacting their 1880 industry composition with subsequent leave-county-out national growth. Because 1880 specialization may itself reflect persistent local comparative advantage, the design does not require unconditional balance in the initial shares. The exclusion restriction is instead conditional: after accounting for observed initial conditions and complete 1850–1870 histories of manufacturing employment, manufacturing HSPS, and HSPS services, subsequent national industry growth must affect service composition through the local industrialization episode it induces rather than through an independent continuation of pre-existing county trends. Identification therefore compares counties in the same state and peak cohort that followed similar observed development paths before 1880 but were differently exposed to the subsequent national growth of the industries in which they were specialized.

This industry-level design also permits a sharper test of the mechanism. National growth across manufacturing as a whole may deepen local industrialization regardless of how production is organized and which manufacturing industries expand. The mechanism, however, predicts a persistent change in services specifically when expansion occurs in industries in which operating at scale creates intensive

managerial, informational, and technical demands. Following [Chandler \(1990\)](#), I classify continuous-process, large-batch, and complex-assembly industries as high throughput and distinguish them from traditional low-throughput industries that could operate at lower scale. Examples of high-throughput industries include iron and steel, chemicals, machinery, and processed foods, while low-throughput industries include textiles and apparel, wood and furniture, printing, and leather goods.¹² The classification follows Chandler’s technological taxonomy rather than being constructed from the county outcomes. High-throughput industries are those in which large fixed investments and continuous or high-volume production were necessary to achieve low unit costs, making large establishments and intensive coordination integral to the production technology; low-throughput industries could operate efficiently at substantially smaller scale. For industry group $g \in \{A, H, L\}$, where A denotes all manufacturing, H high-throughput manufacturing, and L traditional low-throughput manufacturing, I construct

$$Z_c^g = \sum_{j \in g} \omega_{cj,1880}^g G_{j,-c,1880-1900}, \quad (6)$$

where $\omega_{cj,1880}^g$ is industry j ’s share of county c ’s 1880 manufacturing employment within group g , and $G_{j,-c,1880-1900}$ is national employment growth in industry j between 1880 and 1900 after removing county c . Each shifter is standardized within the common estimation sample. The within-group normalization makes the exposure shares sum to one and isolates which industries within a production regime a county initially specialized in rather than the overall size of the group in the local economy. Appendix A.5 reports economy-scaled alternatives and specifications that control directly for the initial high-throughput share.

For each shifter Z_c^g , the first and second stages are

$$M_c^{\text{peak}} = \pi_g Z_c^g + X'_{c,1850} \Pi + \mathcal{H}'_{c,<1880} \Omega + \delta_{s(c)} + \lambda_{P_c} + u_c, \quad (7)$$

$$Y_{c,2000} = \beta_g \widehat{M}_c^{\text{peak}} + X'_{c,1850} \Gamma + \mathcal{H}'_{c,<1880} \Psi + \delta_{s(c)} + \lambda_{P_c} + \varepsilon_c, \quad (8)$$

Where the vector $X_{c,1850}$ contains the five baseline controls measured in 1850, while $\mathcal{H}_{c,<1880}$ collects the complete 1850–1870 histories of manufacturing employment, manufacturing HSPS employment, and HSPS-service specialization. The terms $\delta_{s(c)}$ and λ_{P_c} denote state and manufacturing-peak-cohort fixed effects. Table 2 reports the corresponding OLS benchmark and three just-identified 2SLS specifications on a common sample of 1,487 counties.

The all-industry and high-throughput shifters both strongly predict county manufacturing peaks, with Conley first-stage F statistics of 14.67 and 12.87 and Kleibergen–Paap statistics above 20. However, their long-run implications for the composition of services are different. Using national growth across manufacturing as a whole yields a 2SLS coefficient of 0.133, but neither the conventional estimate nor the weak-instrument-robust reduced form is statistically distinguishable from zero ($p_{AR} = 0.203$). By contrast, the high-throughput IV estimate is 0.283 with a Conley standard error of 0.118. A 10-percentage-point increase in the manufacturing peak induced by national high-throughput growth therefore raises

¹² Appendix A reports the complete IND1950 concordance, the corresponding Chandler references, and the rationale for every included industry.

Table 2: National Manufacturing Industry Shifts and Persistent HSPS Services

	OLS	All industries IV	High-throughput IV	Low-throughput IV
Peak manufacturing employment share	0.029** (0.012)	0.133 (0.105)	0.283** (0.118)	−1.095 (8.010)
First-stage coefficient		0.0150*** (0.0039)	0.0193*** (0.0054)	0.0006 (0.0042)
First-stage Conley F statistic		14.67	12.87	0.02
Kleibergen–Paap F statistic		20.53	21.94	0.03
Anderson–Rubin p -value		0.203	0.009	0.596
State fixed effects	Yes	Yes	Yes	Yes
Peak-cohort fixed effects	Yes	Yes	Yes	Yes
1850 historical controls	Yes	Yes	Yes	Yes
Complete 1850–1870 histories	Yes	Yes	Yes	Yes
Observations	1,487	1,487	1,487	1,487

Notes: The dependent variable is the QCEW HSPS share of private-service employment described in Section 3.1. Peak manufacturing employment is Census of Manufactures average wage earners divided by total classified county employment and is defined over 1880, 1900, 1920, 1930, and 1940. Column 1 reports OLS. Columns 2–4 instrument the peak manufacturing employment share with standardized shift-share instruments based on, respectively, all manufacturing industries, Chandler high-throughput industries, and traditional low-throughput industries. The instruments combine 1880 county industry-employment shares within the relevant group with leave-county-out national industry growth from 1880 to 1900. All columns use the same 1,487 counties and include state and peak-cohort fixed effects, the five 1850 controls, and complete 1850–1870 histories of manufacturing employment and core manufacturing HSPS relative to total classified employment and of HSPS services relative to private-service employment. Coefficient and first-stage standard errors are Conley with a 100-kilometer cutoff. The first-stage Conley F statistic is the square of the Conley first-stage t statistic. The Kleibergen–Paap statistic is the heteroskedasticity-robust excluded-instrument Wald statistic. The Anderson–Rubin p -value is the corresponding Conley reduced-form test. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

the HSPS share of private-service employment in 2000 by 2.83 percentage points—approximately 45 percent of the 6.36-percent cross-county mean. The Anderson–Rubin test rejects zero at the 1-percent level ($p = 0.009$).

The traditional low-throughput shifter has essentially no first stage, so its 2SLS coefficient is not informative. The economically relevant contrast is between broad manufacturing growth, which strongly moves the manufacturing peak without producing a detectable shift in long-run service composition, and high-throughput growth, which moves both.

The contrast points to a scale-dependent organizational legacy of industrialization: persistent changes in service composition arise when manufacturing expansion occurs in production regimes where operating at scale creates intensive managerial, informational, and technical capabilities. Appendix A.5 subjects this result to the principal measurement, classification, exposure, and inference diagnostics.¹³ The result is not driven by any single high-throughput industry, remains positive across alternative classification boundaries and endpoints, and survives substantially wider spatial-correlation cutoffs. Combined with

¹³ The appendix reports QCEW denominator and disclosure bridges, a 2010 endpoint, interior peak cohorts, EFSY-imputed commuting-zone outcomes, alternative Chandler classifications and exposure normalizations, Rotemberg-weight and leave-one-industry-out diagnostics, direct high-versus-comparison reduced forms, shock-level inference, additional geography controls, and Conley cutoffs through 500 kilometers. With exact leave-county-out shock-level inference, the high-throughput estimate remains 0.283, with a t -based p -value of 0.088 and a weak-instrument-robust Anderson–Rubin p -value of 0.062; the common-shock representation yields $p = 0.041$ and $p_{AR} = 0.047$.

the event-time and channel-accounting evidence, these findings motivate the quantitative model in Section 4, which asks how this organizational legacy shapes structural transformation, skill accumulation, and long-run growth.

4 Model

4.1 Overview

Time is discrete and one period is one year. The economy contains three final-good sectors—agriculture, manufacturing, and consumer services—indexed by $j \in \mathcal{J} \equiv \{a, m, c\}$. Firms enter endogenously, draw permanent productivity upon entry, and face exogenous exit risk. Final goods are used for consumption and as intermediate inputs, while producer services help firms organize and coordinate production. Nonhomothetic demand reallocates expenditure from agriculture toward manufacturing and services as income rises, generating the economy’s structural transformation path. Households choose consumption, fertility, and education, so population and the supply of skilled labor evolve endogenously.

The central mechanism is that firm scale raises the return to formalizing recurring organizational tasks. Every firm faces a unit measure of organizational tasks required to operate. Each task generates recurring organizational problems, with their frequency increasing linearly in firm size. These tasks are common across sectors: firms must coordinate workers, process information, manage inputs, and organize decision-making regardless of what they produce. Tasks can be handled informally, in which case the underlying problems must be resolved repeatedly, or formalized by installing a reusable procedure. The stock of formalized tasks constitutes the firm’s *organizational capital*. Because a given task generates more problems in a larger firm, formalization avoids more repeated problem solving and is therefore more valuable at larger scale. Larger firms consequently formalize more tasks and demand more high-skill producer-service workers to do so.

Formalization can also create knowledge that extends beyond the firm. When a firm formalizes a task for which no established procedure exists, it must use skilled labor to develop a solution from scratch. The resulting method is initially proprietary, but may subsequently diffuse into a public stock of organizational knowledge, K_t , allowing other firms to formalize the same task at lower cost. The originating firm therefore does not capture the full value of the knowledge it creates. Because organizational tasks such as accounting, personnel management, quality control, and workflow design arise across sectors, this externality is naturally cross-sectoral: methods developed inside one sector can later lower the cost of organizing production in another.¹⁴

Manufacturing matters in the model not because it uses intrinsically different organizational tasks, but because historically it operated at the scale at which formalization became privately worthwhile. A manufacturing-intensive development path therefore generates more organizational knowledge than an economy which is composed of small firms throughout its development path. The service economy that follows manufacturing inherits organizational knowledge, formalizes more activities, and demands

¹⁴ Historical work documents both sides of this mechanism. Large industrial firms developed new systems of cost accounting, return-on-investment control, statistical quality control, budgeting, workflow design, and personnel management as they confronted the organizational demands of operating at scale (Johnson, 1975; Kaplan, 1984; Chandler, 1977; Shewhart, 1931). Many of these methods were subsequently codified, professionalized, and applied outside manufacturing (McKinsey, 1922, 1924; Yates, 1993; McKenna, 2006; Mol and Birkinshaw, 2014). Appendix C discusses several examples in detail.

more HSPS labor. By raising the return to skill, this organizational legacy also affects education and, ultimately, the economy's long-run skill composition and productivity growth.

The productivity-growth block builds on the tradition of [Jones \(1995\)](#). Unlike standard models in this tradition, however, the long-run growth rate is independent of both population scale and population growth and is instead determined by the economy's *endogenous* share of skilled workers. I circumvent the reliance of this class of models on positive asymptotic population growth to generate income growth by adopting the endogenous fertility mechanism present in [Galor and Weil \(2000\)](#) which creates a feedback loop between population and income growth which drives fertility to replacement and stabilizes long-run income growth. This also embeds the transition from agriculture to services within a broader transition from stagnation to sustained growth, unifying two central dimensions of development: structural transformation and the emergence of sustained income-per-capita growth. Proposition 2 contains the details.¹⁵

The remainder of the section proceeds in four steps. I first characterize firms' organizational choices and show how scale determines formalization and HSPS demand. I then describe the creation, adaptation, and diffusion of organizational procedures and the entry margin that determines equilibrium firm scale. Next, I embed these decisions in the structural-transformation equilibrium and introduce the household choices that determine education and population growth. Finally, I characterize the limiting equilibrium, including the long-run composition of services, skill supply, and productivity growth. The main text emphasizes the economic mechanisms, while derivations, equilibrium definitions, and proofs are collected in Appendix D.

4.2 Production and organizational choice

Final-good production. A firm f in sector j has permanent ability $a_f \sim G$ and an inherited formalization frontier s_{fjt} . For a generic firm I suppress the f subscript and write its state as (a, s) . Skilled and unskilled labor form a CES composite, and value added exhibits decreasing returns to firm scale:

$$\begin{aligned}\ell_{fjt} &= \left[\alpha_j (h_s L_{fjt}^s)^{\frac{\sigma_L-1}{\sigma_L}} + (1 - \alpha_j) (h_u L_{fjt}^u)^{\frac{\sigma_L-1}{\sigma_L}} \right]^{\frac{\sigma_L}{\sigma_L-1}}, \\ v_{fjt} &= A_{jt} a_f \ell_{fjt}^\eta, \quad 0 < \eta < 1.\end{aligned}\tag{9}$$

Gross output combines value added, agricultural and manufacturing materials, and producer services used directly in production:

$$\begin{aligned}y_{fjt} &= Z_{jt} v_{fjt}^{\theta_{jt}} (m_{fjt}^a)^{\gamma_{jt}^a} (m_{fjt}^m)^{\gamma_{jt}^m} (P_{fjt}^p)^{\gamma_{jt}^p}, \\ \theta_{jt} + \gamma_{jt}^a + \gamma_{jt}^m + \gamma_{jt}^p &= 1.\end{aligned}\tag{10}$$

Let w_{jt}^v denote the unit cost of the labor composite and let $n = \nu_{jt} \ell$ denote the associated employment headcount. Both follow from cost minimization in the CES labor nest. The distinction between ℓ and n matters because production depends on effective labor, whereas organizational problems arise from

¹⁵ See [Jones \(2022\)](#). In standard semi-endogenous growth models such as [Jones \(1995\)](#), sustained long-run productivity growth ultimately requires sustained population growth, so growth vanishes when population becomes stationary, an increasingly problematic prediction of these models in a world where fertility rates are below replacement in many advanced economies. Models outside the semi-endogenous tradition can combine replacement fertility with positive growth ([Peretto and Valente, 2015](#)), but through different mechanisms.

coordinating workers and therefore depend on employment headcount. The input-output coefficients are initialized to imply value-added production in 1850 and then evolve gradually toward their 2015 BEA input-output shares.¹⁶

Organizational tasks. Every firm faces the same unit mass of organizational tasks. Tasks differ in complexity $z \geq 0$, which determines how difficult they are to formalize and is exponentially distributed:

$$\mathcal{F}(z) = 1 - e^{-z}.$$

A firm's organizational state is summarized by its formalization frontier s : all tasks with $z \leq s$ have been formalized, so its organizational-capital stock can equivalently be represented by the formalized task share

$$q(s) \equiv \mathcal{F}(s).$$

Each task generates $b(n) = n/\bar{N}$ organizational problems, where n is firm employment and $\bar{N}$ normalizes firm scale. Thus, the same organizational task generates more instances requiring resolution in a larger firm. If a task is handled informally, these problems must be repeatedly diagnosed and resolved; once the task is formalized, the firm can instead apply a reusable procedure. Let $r > 1$ denote the current-period-equivalent amount of informal problem solving required relative to operating under such a procedure.¹⁷ Conditional on (n, s) , recurring demand for high-skill formal and low-skill informal producer services is therefore

$$H^o(n, s) = \frac{n}{\bar{N}} q(s), \quad L^o(n, s) = \frac{n}{\bar{N}} r[1 - q(s)]. \quad (11)$$

At competitive prices p_t^H and p_t^L , recurring coordination costs are

$$C_t^o(n, s) = \frac{n}{\bar{N}} [q(s)p_t^H + r[1 - q(s)]p_t^L]. \quad (12)$$

Formalization lowers these recurring costs whenever $rp_t^L > p_t^H$, because it replaces repeated informal problem solving with the application of an established procedure. Since the number of problems rises with employment, the associated saving is larger in bigger firms.

The same organizational state also determines the composition of producer services used directly in production:

$$P^P(s) = \left[q(s)^{1/\sigma_P} (H^P)^{\frac{\sigma_P-1}{\sigma_P}} + [1 - q(s)]^{1/\sigma_P} (L^P)^{\frac{\sigma_P-1}{\sigma_P}} \right]^{\frac{\sigma_P}{\sigma_P-1}}. \quad (13)$$

A higher formalization share therefore shifts producer-service use toward high-skill inputs. Because high-skill producer services are more expensive, formalization can raise the unit cost of production even as it reduces repeated coordination costs. The firm's organizational choice is therefore governed by a single tradeoff: formalization economizes on recurring problem solving but requires a more skill-intensive organization. The next subsection shows how the net return to this tradeoff varies with firm scale.

¹⁶ Allowing the input-output structure to evolve avoids imposing the modern production network on the 1850 economy. With fixed modern IO coefficients, historical changes in employment generated by the expansion of intermediate-input use would instead be absorbed by the estimated preference and productivity parameters.

¹⁷ More precisely, if R denotes the discounted number of informal resolutions associated with the problems generated by a task over the expected life of the firm, then $r = [1 - \beta(1 - \delta_f)]R$. The parameter R is disciplined primarily by sectoral high-skill producer-service employment shares.

Static firm problem. After optimizing over labor composition, materials, and productive producer services, the firm's static operating profit can be written as

$$\pi_{jt}^*(a, s) = \max_{\ell \geq 0} \{p_{jt}^v(q(s))A_{jt}a\ell^\eta - [w_{jt}^v + c_{jt}^o(q(s))]\ell\}, \quad (14)$$

where

$$c_{jt}^o(q) = \frac{\nu_{jt}}{\bar{N}} [qp_t^H + r(1 - q)p_t^L]$$

is recurring organizational cost per unit of effective labor, and $p_{jt}^v(q)$ is the shadow value of firm value added after materials and productive producer services have been chosen optimally. Its dependence on q captures the fact that using different organizational inputs changes the firm's unit cost in production. Appendix D derives both objects from the gross-output technology.

The firm's optimal effective labor and employment are

$$\ell_{jt}^*(a, s) = \left[\frac{\eta p_{jt}^v(q(s))A_{jt}a}{w_{jt}^v + c_{jt}^o(q(s))} \right]^{\frac{1}{1-\eta}}, \quad n_{jt}^*(a, s) = \nu_{jt}\ell_{jt}^*(a, s). \quad (15)$$

Thus more productive firms are larger: employment rises with ability with elasticity $1/(1 - \eta)$.

The key comparative static concerns the return to formalization. Combining the envelope condition with the firm's employment choice gives

$$\frac{\partial \pi_{jt}^*(a, q)}{\partial q} = n_{jt}^*(a, q)\Psi_{jt}(q), \quad (16)$$

where

$$\Psi_{jt}(q) \equiv \frac{w_{jt}^v + c_{jt}^o(q)}{\eta \nu_{jt}} \frac{\partial \log p_{jt}^v(q)}{\partial q} + \frac{rp_t^L - p_t^H}{\bar{N}}. \quad (17)$$

The first term is the marginal effect of formalization through the productive producer-service input. When high-skill producer services are more expensive than low-skill producer services, $p_t^H > p_t^L$, this term is negative: hiring skilled producer service workers raises unit cost and, in isolation, reduces desired firm scale. The second term is the recurring coordination saving: formalization replaces informal problem solving costing rp_t^L per problem with application of a formal procedure costing p_t^H . Thus $\Psi_{jt}(q)$ is the net marginal operating gain per worker from formalization, and the total operating gain is proportional to firm employment.

Proposition 1 (Firm scale and organizational formalization). *Fix sector j , date t , and organizational state q . Conditional on equilibrium prices, if $\Psi_{jt}(q) > 0$, the marginal operating benefit of formalizing an additional task is strictly increasing in firm employment:*

$$\frac{\partial \pi_{jt}^*(a, q)}{\partial q} = n_{jt}^*(a, q)\Psi_{jt}(q).$$

Hence larger and more productive firms obtain a greater operating return from the same organizational procedure. Moreover, if $rp_t^L > p_t^H$, then $\Psi_{jt}(q) > 0$ implies that greater formalization also raises the firm's desired employment.

Proof. See Appendix D.

The scale dependence in Proposition 1 becomes consequential once formalization requires a one-time installation cost. Because this cost does not scale proportionally with firm employment, larger firms find it profitable to formalize a broader set of tasks. The next subsection introduces the one-time cost of formalization, characterizes firms' dynamic organizational investment decisions, and then shows how newly created procedures can diffuse into public organizational knowledge.

4.3 Organizational investment, diffusion, and entry

This subsection turns the static scale mechanism into a dynamic one. Firms pay a one-time cost to formalize organizational tasks, and the procedures they create can subsequently diffuse into a public stock of organizational knowledge, lowering the cost of formalization for other firms and sectors. Entry determines equilibrium firm size and hence whether these organizational investments become profitable.

Creation and adaptation. Let $K_t \in [0, 1]$ denote the public stock of organizational knowledge, measured as the share of tasks for which a reusable procedure is publicly available. Because tasks are ordered by complexity, public knowledge takes the form of a frontier: a procedure for task z is publicly available whenever $\mathcal{F}(z) \leq K_t$. Formalizing a task requires a one-time organizational investment. If no public procedure exists, the firm must create one from scratch using high-skill producer-service labor; if a procedure is already contained in K_t , the firm can instead adapt it at a fraction of the creation cost:

$$C_t^C(z) = \frac{p_t^H}{\bar{e}_t^o} e^{z-1}, \quad (18)$$

$$C_{jt}^A(z) = \frac{\rho}{\xi_j} C_t^C(z), \quad \xi_m = 1, \quad \rho \in (0, 1). \quad (19)$$

Here $\bar{e}_t^o$ measures the average education embodied in the current workforce, inherited from the education choices made by households one generation earlier. For a given task, a more educated workforce requires fewer skilled-labor resources to produce a reusable procedure. Because education rises gradually over development, this technology also links the historical expansion of education to the timing at which increasingly complex organizational methods become economical to develop. The parameter ρ governs the common cost advantage of adapting an existing procedure, while ξ_j determines how adaptation costs vary across sectors.

A firm moving from inherited formalization frontier s to a new frontier $u \geq s$ therefore incurs the one-time investment cost

$$\begin{aligned} \mathcal{C}_{jt}^O(s, u; K_t) &= \int_s^u \tilde{C}_{jt}(z; K_t) d\mathcal{F}(z), \\ \tilde{C}_{jt}(z; K_t) &= \begin{cases} \min\{C_t^C(z), C_{jt}^A(z)\}, & \mathcal{F}(z) \leq K_t, \\ C_t^C(z), & \mathcal{F}(z) > K_t. \end{cases} \end{aligned} \quad (20)$$

Thus a task beyond the public frontier must be formalized by creating a new procedure, whereas a task inside the frontier can exploit the accumulated organizational knowledge of previous firms.

A firm produces in period t using its inherited frontier s and chooses the frontier $u \geq s$ carried into period $t + 1$. Because the economy is generally away from its limiting equilibrium, this is a nonstationary dynamic problem: the firm's choice depends on the entire perfect-foresight continuation path of aggregate states and prices. I suppress this dependence in the notation. Taking that continuation path

as given, the firm's value satisfies

$$W_{jt}(a, s) = \pi_{jt}^*(a, s) + \max_{u \geq s} \{ -\mathcal{C}_{jt}^O(s, u; K_t) + \beta(1 - \delta_f)W_{j,t+1}(a, u) \}. \quad (21)$$

Because the organizational investment cost is fixed and Proposition 1 shows that the benefit to formalization rises with firm scale, larger firms therefore find it profitable to formalize a broader range of tasks and incidentally create more organizational knowledge.

Diffusion. Newly created procedures are initially proprietary. Let $\mu_{jt}(a, s)$ denote the distribution of active firms in sector j , let

$$N_{jt} = \int n_{jt}^*(a, s) d\mu_{jt}(a, s), \quad N_t = \sum_{j \in \mathcal{J}} N_{jt},$$

and let $c_{jt}(a, s; z) \in \{0, 1\}$ indicate that a firm creates at date t a procedure for task z that lies beyond the current public frontier. Economy-wide employment-weighted exposure to newly created proprietary procedures is

$$\mathcal{E}_t = \frac{1}{N_t} \sum_{j \in \mathcal{J}} \int n_{jt}^*(a, s) \left[\int_{\mathcal{F}(z) > K_t} c_{jt}(a, s; z) d\mathcal{F}(z) \right] d\mu_{jt}(a, s).$$

Thus $\mathcal{E}_t$ is larger when more procedures are created and when those procedures are created in firms employing more workers. Public organizational knowledge evolves according to

$$K_{t+1} = \min\{1, K_t + \delta_K \mathcal{E}_t\}, \quad (22)$$

where δ_K governs the rate at which proprietary procedures diffuse into the public stock. Because K_t is irreversible, a temporary period of large-scale production can therefore leave a persistent change in equilibrium skill-demand.

Entry. A potential entrant pays its setup cost before drawing permanent ability $a \sim G$ and enters with frontier $s = 0$. Let I_{jt} denote entry. Convex setup-input supply gives the marginal entry cost $F_j I_{jt}^{\phi_E}$. Free entry and the law of motion for establishment counts are summarized by

$$\begin{aligned} V_{jt}^E &= \int W_{jt}(a, 0) dG(a), & M_{jt} &= (1 - \delta_f)M_{j,t-1} + I_{jt}, \\ I_{jt} &\geq 0, & F_j I_{jt}^{\phi_E} - V_{jt}^E &\geq 0, \\ & & I_{jt} [F_j I_{jt}^{\phi_E} - V_{jt}^E] &= 0. \end{aligned} \quad (23)$$

The shifter F_j governs equilibrium entry and therefore average firm size at the sectoral level. Through Proposition 1, it consequently determines the strength of firms' incentives to invest in formal organization.

4.4 Households and aggregate dynamics

The household block connects firms' organizational choices to the economy's structural transformation and long-run growth path. Nonhomothetic demand reallocates expenditure from agriculture through

manufacturing and into services. Endogenous skill acquisition allows the resulting changes in firms' demand for high-skill labor to alter the skill composition of the workforce, while fertility and education determine population dynamics. Together, these choices close the feedback between structural transformation, skill supply, population, and productivity growth. The household side builds on the quantity–quality mechanism of [Galor and Weil \(2000\)](#) and [Lagerlöf \(2006\)](#) and the nonhomothetic demand system of [Comin et al. \(2021\)](#); the productivity-growth technology introduced below follows the tradition of [Jones \(1995\)](#).

Fertility and education. Following [Galor and Weil \(2000\)](#) and [Lagerlöf \(2006\)](#), households choose fertility and educational investment in their children. Adults differ by employment origin $i \in \{a, m, s\}$, with s pooling all service activities and $\lambda_{i,t}$ denoting the corresponding population shares. Origin affects only the fixed time cost of raising a child, $\tau_i \equiv \tau \tau_i^n$; labor income and education costs are otherwise pooled across households. Preferences and child quality are

$$U_{i,t} = (1 - \gamma) \log C_{i,t} + \gamma \log[n_{i,t} h_{i,t}], \quad (24)$$

$$h_{i,t} = \frac{e_{i,t} + \tau_h}{e_{i,t} + \tau_h + g_t} \quad 0 < \tau_h \ll 1 \quad (25)$$

Where $n_{i,t}$ and $h_{i,t}$ denote fertility and child human capital, respectively. Here $g_t \equiv \sum_{i \in \{a, m, s\}} \lambda_{i,t} g_{i,t}^A$ is aggregate productivity growth. The small constant $\tau_h > 0$ keeps child quality strictly positive at the zero-education corner. As in [Galor and Weil \(2000\)](#), faster technological change raises the return to education, inducing households to invest more in their children and reduce fertility. This quantity–quality response provides the demographic feedback that stabilizes population in the long run: population becomes an endogenous adjustment margin rather than an exogenous source of sustained growth. Raising and educating a child requires $\tau_i + \tau_e(S_t)e_{i,t}$ units of parental time, where

$$\tau_e(S) = 1 - S(1 - \tau_e^s), \quad \tau_e^s \in (0, 1).$$

Thus, a more skilled economy makes educational investment less costly in parental time. This reduced-form complementarity captures the greater availability of educated parents, teachers, peers, and educational institutions in a more skilled environment and is consistent with education technologies featuring dynamic complementarity ([Cunha and Heckman, 2007](#); [Cunha et al., 2010](#)). It also links the terminal skilled share to long-run growth. At a given rate of technological change, a higher S_t lowers the time cost of educating children and raises fertility; restoring replacement fertility therefore requires faster technological change, which induces sufficiently greater educational investment. Proposition 2 shows that this makes the limiting growth rate increasing in the terminal skilled share. Let $\bar{w}_t$ denote pooled potential labor earnings per adult and d_t net dividends per adult. The household budget constraint is

$$\sum_{j \in \mathcal{J}} p_{j,t}^c C_{ij,t} = [1 - n_{i,t} \{\tau_i + \tau_e(S_t)e_{i,t}\}] \bar{w}_t + d_t.$$

Final demand and skill acquisition. Final consumption follows the nonhomothetic CES system of [Comin et al. \(2021\)](#):

$$\sum_{j \in \mathcal{J}} \omega_j^{1/\sigma} \left(\frac{C_{ij,t}}{C_{i,t}^{\varepsilon_j}} \right)^{\frac{\sigma-1}{\sigma}} = 1, \quad 0 < \sigma < 1. \quad (26)$$

With the parameter ordering used below, income and relative-price effects reallocate expenditure away from agriculture, generate a hump-shaped manufacturing share, and eventually shift activity toward services. This provides the structural-transformation path along which firm scale, organizational investment, and knowledge diffusion evolve.

Members of the entering cohort also choose whether to acquire skill. I summarize this decision by

$$p_t^S = \mathcal{S}(\Pi_{t+1}), \quad \mathcal{S}'(\Pi) > 0, \quad (27)$$

where Π_{t+1} is the expected equilibrium skilled-wage premium next period and expectations are rational. In the quantitative model, $\mathcal{S}$ is parameterized as a smooth logistic choice probability.¹⁸ Greater formalization raises demand for high-skill producer-service labor and therefore affects the skilled composition of subsequent cohorts through the equilibrium skill premium.

Laws of motion. Let Δ denote generation length. Following the annualization in Lagerlöf (2006), a fraction $1/\Delta$ of the workforce is replaced each year.¹⁹ Let $\bar{n}_t$ denote average fertility. Population and the skilled share evolve according to

$$L_{t+1} = L_t \bar{n}_t^{1/\Delta}, \quad S_{t+1} = \left(1 - \frac{1}{\Delta}\right) S_t + \frac{1}{\Delta} p_t^S. \quad (28)$$

Sectoral value-added productivity growth increases with the economy's skilled-worker mass:

$$g_{j,t}^A \equiv \frac{A_{j,t+1}}{A_{j,t}} - 1 = \left[1 + \lambda_A (S_t L_t)^{\phi_j}\right]^{1/\Delta} - 1, \quad \phi_j \in (0, 1). \quad (29)$$

Consumer services and both producer-service industries share the service exponent ϕ_c , while gross-output productivity grows exogenously at rate g^{GO} .

Equation (29) builds on the central insight of Jones (1995) that a larger supply of workers raises productivity growth. In the standard semi-endogenous formulation, diminishing returns to the accumulated stock of knowledge—the idea that new ideas become progressively harder to find—eliminate the population-scale effect, but in doing so make the long-run growth rate depend on exogenous population growth. I take a different route. I omit the scale effect through the demographic equilibrium: technological change affects education and fertility, fertility converges to replacement, and population converges to a finite endogenous level. This formulation is natural in the present setting because over an extended time window population and the supply of skilled workers are themselves outcomes of the development process rather than exogenous inputs into growth. The resulting limiting growth rate is therefore fully endogenous and remains positive even when population growth is zero. Proposition 2 characterizes this equilibrium formally.

Equilibrium. Manufacturing is the numeraire. Given the initial state of the economy, a perfect-foresight competitive equilibrium is a sequence of prices, household choices, firm policies, entry, firm

¹⁸ Specifically, $\mathcal{S}(\Pi) = \left[1 + \exp\left(-\frac{\beta_\pi \log \Pi - c^{\text{skill}}}{\sigma_{\text{skill}}}\right)\right]^{-1}$.

¹⁹ Time is measured annually throughout the model, while Δ denotes generation length. Household fertility and cohort skill choices are generational objects, which I map into annual population and workforce dynamics following the annualization in Lagerlöf (2006).

distributions, and aggregate states such that households optimize; firms solve (14) and (21); entry satisfies (23); final-good, producer-service, and labor markets clear; and the firm distribution, public organizational knowledge, population, skills, and sectoral productivities follow their respective laws of motion. Aggregate uncertainty is absent, although firms face idiosyncratic exit risk and entrants draw heterogeneous ability. A formal definition is provided in Appendix D.²⁰

4.5 Limiting growth and the manufacturing legacy

I characterize the model's long-run behavior in two steps. First, I show that the terminal skilled share pins down the economy's limiting productivity growth rate, while population scale adjusts to the finite level required to sustain it. Second, I show how a temporary manufacturing phase can permanently raise the terminal skilled share, and thus long-run growth, by expanding the stock of public organizational knowledge inherited by the service economy. For the analytical results, I assume that broad services, including consumer and producer services, absorb the workforce in the limit; in the calibrated economy they account for 92 percent of terminal employment.²¹

Proposition 2 (Endogenous population scale and limiting growth). *Fix a terminal skilled share $S \in (0, 1)$ and let*

$$\zeta \equiv 1 + \frac{d^\infty}{\bar{w}^\infty}$$

denote the full-income ratio in the terminal service equilibrium. Suppose $\gamma\zeta > \tau_s$ and, for this analytical result, $\tau_h = 0$. Population converges from every $L_0 > 0$ to a unique steady state $L^\infty(S; \zeta)$, and fertility converges to replacement. At that steady state,

$$g^\infty = \mathcal{G}(S; \zeta) \equiv \frac{(\gamma\zeta - \tau_s)^2}{\tau_s} \times \frac{1}{\tau_e(S)}, \quad (30)$$

$$L^\infty(S; \zeta) = \frac{1}{S} \left[\frac{[1 + \mathcal{G}(S; \zeta)]^\Delta - 1}{\lambda_A} \right]^{1/\phi_c}. \quad (31)$$

Holding ζ fixed,

$$\frac{\partial \mathcal{G}(S; \zeta)}{\partial S} > 0.$$

Proof. See Appendix D.

Proposition 2 follows from a stabilizing feedback between population, technological change, and fertility. Conditional on the terminal skilled share, a larger population raises the mass of skilled workers SL and hence productivity growth. Faster technological change increases the return to educating children, raising educational investment and reducing fertility. Consequently, when population is below L^∞ , productivity growth is too low to induce the education required for replacement fertility and population expands; when population is above L^∞ , the reverse occurs. This feedback selects a unique finite population at which fertility equals replacement. The dependence of limiting growth on the skilled share follows from $\tau'_e(S) < 0$: a higher S makes education more effective, so at the previous growth rate children require less parental time and fertility rises above replacement. Restoring replacement fertility

²⁰ Computationally, I solve the perfect-foresight transition using the extended-path logic of [Fair and Taylor \(1983\)](#). I conjecture the aggregate path, solve firms' dynamic decisions backward along that path, propagate the firm distribution and aggregate states forward, and iterate until the conjectured and implied aggregate sequences coincide.

²¹ The remaining nonservice employment in the calibrated economy reflects input-output linkages: even at high income levels, service production continues to require agricultural and manufacturing intermediates.

requires faster technological change, which induces sufficiently greater educational investment. Thus the terminal skilled share determines the limiting growth rate, while population adjusts endogenously to the level required to sustain it.

I next ask when passing through manufacturing places the economy on a permanently higher growth path. Proposition 2 establishes that limiting productivity growth is determined by the terminal skilled share. Manufacturing therefore raises long-run growth only if it leaves the subsequent service economy sufficiently more skill intensive to raise S^∞ . In the model, this occurs through organizational knowledge: large firms create reusable procedures, those procedures diffuse, and service firms later formalize more tasks and demand more high-skill producer-service labor than they would have without the manufacturing phase. To characterize when this chain of events occurs in the model, consider two economies beginning from the same state (S_1, L_1, K_1) . Economy A passes through manufacturing before entering services, whereas economy B moves directly from agriculture into services. I isolate the mechanism using a three-stage comparison with symmetric static firms, common prices and operating-surplus terms, and a common terminal full-income ratio.²² I additionally require that the service economy is not organizationally saturated, so that access to an additional public procedure induces some service firms to formalize an additional task. Appendix D states this condition formally.

Proposition 3 (Manufacturing and long-run growth). *Under the preceding restrictions,*

$$g_A^\infty > g_B^\infty \iff S_A^\infty > S_B^\infty \iff K_{A,3} > K_{B,3} \iff \underbrace{\frac{\tau_m}{\tau_s}}_{\text{education}} \underbrace{(S_1 L_1)^{\phi_m - \phi_c}}_{\text{productivity growth}} \underbrace{\left(\frac{F_m}{F_c}\right)^{\frac{2}{1+\phi_E}}}_{\text{firm scale}} > 1. \quad (32)$$

Proof. See Appendix D. Proposition 3 makes the model's central mechanism explicit. Manufacturing raises long-run growth when it leaves the subsequent service economy with a larger stock of public organizational knowledge: cheaper adaptation induces service firms to formalize more tasks, raising HSPS demand, the return to skill, and ultimately S^∞ and g^∞ . Condition (32) shows when this occurs. The terms τ_m/τ_s and $(S_1 L_1)^{\phi_m - \phi_c}$ capture the education force along the transition, through differences in household education choices and sectoral productivity growth, respectively; both affect the availability of educated labor able to develop new procedures. The term $(F_m/F_c)^{2/(1+\phi_E)}$ captures the complementary scale force: through entry, it determines relative firm size and hence the private return to formalization. The condition therefore summarizes the interaction between the capacity to create organizational knowledge and the scale at which creating it becomes worthwhile. The quantitative model relaxes the simplifying restrictions behind the proposition and determines the strength of these forces in full general equilibrium.

5 Calibration

I calibrate the model at an annual frequency to the U.S. development path from 1850 to 2019. The economy is initialized in 1790, allowing the endogenous state variables to evolve before the first targeted observations, and is solved through a terminal horizon well beyond 2019. Parameters fixed by normalization, timing conventions, direct empirical evidence, or prior literature are assigned externally, while

²² For this analytical comparison I also set $\tau_h = 0$ and suppress the annualization.

the remaining 19 parameters are estimated jointly by simulated method of moments (SMM). Appendix E describes the data, external parameter choices, numerical solution, parameter estimates, and local identification analysis.

Let $m_{bk}(\Theta)$ and $\hat{m}_{bk}$ denote the model and data values of moment k in block b . I estimate Θ by minimizing the block-weighted SMM objective

$$r_{bk}(\Theta) = \frac{m_{bk}(\Theta) - \hat{m}_{bk}}{\max\{|\hat{m}_{bk}|, 0.01\}}, \quad \mathcal{L}(\Theta) = \frac{\sum_b \omega_b K_b^{-1} \sum_{k=1}^{K_b} r_{bk}(\Theta)^2}{\sum_b \omega_b}. \quad (33)$$

The normalization expresses discrepancies relative to the magnitude of each target, with the 0.01 floor preventing moments close to zero from receiving disproportionate weight. Averaging within each block prevents long time series from mechanically dominating scalar targets.²³

The set of moment blocks disciplines the main quantitative margins of the model. Sectoral employment paths, normalized GDP per capita, and the height and timing of the manufacturing employment peak discipline structural transformation and aggregate development. Fertility paths and the agriculture-to-services fertility ratio discipline the demographic transition. The aggregate skilled share, skilled earnings premium, and sectoral value-added skill shares discipline skill supply and demand. Sectoral and aggregate high-skill producer-service (HSPS) shares, together with sectoral firm size, discipline the scale-formalization mechanism. Finally, sectoral value-added productivity growth disciplines the productivity-growth technology. Scalar targets are measured in 2019 in both the model and the data unless otherwise noted; the agriculture-to-services fertility ratio is measured over its historical window, while manufacturing peak timing is treated as an interval target, with dates between 1930 and 1960 receiving no penalty. Appendix E reports the data sources, measurement windows, block weights, and construction of each target.

5.1 Measurement

Producer and consumer services. The model distinguishes services by their use rather than by the industry of the supplier. This distinction matters because transportation, wholesale trade, finance, insurance, real estate, legal services, and other business services sell to both firms and final consumers. I combine 2019 IPUMS USA microdata with the 2019 BEA Summary Use Table. Within agriculture, manufacturing, and consumer-service industries, occupations are divided between sectoral value-added work and organizational or technical producer-service tasks. Within each group workers are then divided by education, with high skill defined as at least a bachelor’s degree.

Employment in industries that primarily supply intermediate services is allocated across destination sectors using BEA intermediate-use expenditure shares; the portion sold to final consumers is assigned to consumer-service value added. This construction yields sector-specific HSPS and skilled value-added employment shares without imposing a mutually exclusive producer-service industry classification. The IPUMS–BEA concordance is described in Appendix E. The allocation applies a common skill composition and labor requirement to an industry’s sales across buyers, the identifying restriction required because worker records cannot be linked to the purchaser of the industry’s output.

²³ I assign a larger weight to the manufacturing peak share, manufacturing peak timing, sectoral productivity growth, and sectoral HSPS shares. The additional weight is placed on moments that directly discipline the paper’s central structural-transformation and organizational mechanisms. Appendix E reports the complete weighting scheme.

Firm scale. Modern firm-size targets come from the 2019 Census Statistics of U.S. Businesses and are measured as employees per employer firm. Average firm sizes are 7.435 workers in agriculture, 49.694 in manufacturing, and 20.454 in broad services. I calibrate the model’s scale dimension using firm size rather than establishment size as an establishment size measure would disproportionately lower scale in services, where multi-establishment firms are especially common (Hsieh and Rossi-Hansberg, 2023). A common initial scale of 10.7 workers is anchored in both manufacturing and services to the 1880 Census of Manufactures; the subsequent historical manufacturing observations, excepting the 2019 targets, are used to assess the model’s performance in generating the timing of firm size growth and decline over the development path.

5.2 Parameters

Parameters governing standard preferences, firm dynamics, substitution elasticities, and units are fixed before estimation using conventional values or normalizations; Appendix E reports the complete list and their sources. I discuss here the assignments that are most specific to the mechanism.

The dispersion of idiosyncratic skill-acquisition costs, σ_{skill} , determines the elasticity of skill supply to the expected skilled earnings premium. With the normalization $\beta_{\pi} = 1$, the logit rule in equation (27) implies

$$\sigma_{\text{skill}} = \frac{\Delta \log \Pi}{\Delta \log \{p^S / (1 - p^S)\}}.$$

I discipline this parameter using the long-run U.S. joint evolution of college attainment and the college earnings premium. Using bachelor-or-higher attainment and the corresponding skilled–unskilled earnings premium gives an implied value close to 0.10, which I use as the baseline. This follows the quantitative education-choice approach of Heathcote et al. (2010), who discipline heterogeneity in college costs using aggregate changes in educational attainment and the college wage premium; the logit random-utility structure follows McFadden (1974). The location parameter c_{col} is instead estimated jointly and primarily disciplines the level of the skilled share.

The knowledge-diffusion parameter δ_K governs how quickly exposure to newly created proprietary procedures expands the public stock of organizational knowledge. I set $\delta_K = 0.09$. Mansfield (1961) reports a mean time of 7.8 years from first commercial use to 50 percent adoption across twelve industrial innovations, corresponding to an annual closing rate of $1 - 2^{-1/7.8} = 0.085$. Appendix E reports the full mapping and shows that the headline results are robust to significant variations in δ_K .

Two additional organizational parameters are assigned rather than estimated directly. The scale normalization $\bar{N} = 25$ implies that recurring coordination absorbs approximately five percent of employment in the calibrated economy, an order of magnitude consistent with the organization-capital evidence in Atkeson and Kehoe (2005). The adaptation-cost ratio is set to $\rho = 0.40$, implying a substantial but incomplete cost advantage to adapting an existing procedure rather than developing it from scratch. This is consistent with evidence that established management practices can be transferred across firms through consulting, training, and managerial knowledge flows, and that informational barriers are an important obstacle to their adoption (Bloom et al., 2013; Giorcelli, 2019). Because these studies do not identify the adaptation-to-creation cost ratio directly, I treat ρ as a maintained benchmark and report sensitivity to it in Appendix E.

The remaining 19 parameters are chosen jointly by simulated method of moments. Table ?? reports their estimates and principal empirical disciplines. Fertility moments primarily discipline (τ, τ_a^n) ; sectoral

Table 3: Parameter-Level Local Identification Diagnostics

Parameter	Estimate	Unique share (%)	HC1 relative SE (%)
τ	0.03236	17.92	4.86
τ_a^n	0.44017	30.66	5.92
ϕ_a	0.99840	15.49	10.15
ϕ_m	0.91919	14.04	10.38
ϕ_c	0.69046	15.97	10.86
g_{GO}	0.00318	61.33	24.31
c_{col}	0.56821	35.90	5.36
α_a	0.24338	81.42	2.81
α_m	0.33001	84.90	2.96
α_c	0.40995	82.84	2.62
ϵ_c	1.01048	24.87	30.15
w_a	0.65105	34.49	12.44
w_c	3.82326	25.57	39.53
R	22.78696	33.06	5.33
ξ_a	3.56402	64.17	88.56
ξ_c	6.87739	70.22	491.31
F_a	25.872	80.00	48.09
F_m	187.625	69.56	60.27
F_c	85.117	91.00	22.85


Notes: The unique shares use the exact, smooth, nonredundant 82-by-19 Jacobian. HC1 standard errors are computed in log-parameter space and translated to levels by the delta method.

productivity growth and the GDP-per-capita path discipline ($\phi_a, \phi_m, \phi_c, g_{GO}$); aggregate and sectoral skill shares discipline ($c_{col}, \alpha_a, \alpha_m, \alpha_c$); structural-change paths discipline (ϵ_c, w_a, w_c);²⁴ sectoral HSPS shares discipline (R, ξ_a, ξ_c); and sectoral firm-size moments discipline (F_a, F_m, F_c). These mappings describe the principal sources of variation rather than one-parameter-one-moment identification: all 19 parameters are estimated jointly.


The sectoral adaptability parameters are the least precisely identified objects in the calibration. In particular, $\xi_c = 6.8774$ places consumer services near the saturation region of the adaptation technology, where further increases have little effect on model moments. The HSPS targets strongly reject low adaptability but are less informative about its precise value once service firms are able to exploit most procedures developed elsewhere.²⁵ Appendix Figure 13 reports the corresponding parameter profiles and singular-value diagnostics. Two additional parameters lie close to economically meaningful bounds. The estimated service nonhomotheticity parameter, $\epsilon = 1.0105$, implies only a modest direct luxury-good effect for services. The estimate is lower than the service nonhomotheticity typically obtained in models using CLM-type preferences, but the comparison is not one-to-one: here ϵ is estimated jointly with sector-specific productivity growth, time varying input-output linkages, and endogenous firm organization. As emphasized by Herrendorf et al. (2013), the apparent importance of income effects in the structural change literature depends importantly on production structure of the model and how it is mapped to the data. In this model, the estimate therefore implies that the U.S. service transition is not being generated primarily by a large income effect. The agricultural growth elasticity, $\phi_a = 0.9984$, is

²⁴ I set ϵ_a to 0.01 consistent with significant subsistence constraints, Kongsamut et al. (2001); Herrendorf et al. (2014).


²⁵ Given observed service-sector firm size, matching the consumer-service HSPS share requires firms to make extensive use of procedures created outside the sector. Once this margin is nearly exhausted, higher values of ξ_c have little additional effect.




(a) Decadal GDP-per-capita growth



(b) Basic education and school enrollment



(c) Sectoral firm size



(d) High-skill producer-service employment

Figure 4: Historical-path validation. Panel (a) reports annualized GDP-per-capita growth by decade; the model line is the seven-decade quadratic low-frequency trend, with the unsmoothed series available in the replication files. Panel (b) compares the model's basic-education index with the primary-plus-secondary enrollment ratio. Because the two education measures have different units, panel (b) evaluates timing and shape. Panel (c) compares the model's sectoral firm-size paths with historical manufacturing establishment sizes, later manufacturing employer-firm sizes, and the 2019 sectoral calibration targets. Panel (d) compares the model-implied employment share of high-skill producer services with its empirical counterpart, constructed using harmonized IPUMS occupations and industries together with BEA input-use weights.

similarly close to its upper bound and implies nearly linear sensitivity of agricultural productivity growth to skilled-worker mass. This reflects the high historical rate of agricultural productivity growth relative to manufacturing and services, which the common growth shifter must reconcile through sectoral elasticities.

5.3 Fit and untargeted validation

Figure 5 compares the model with the U.S. employment transition. The model reproduces the decline of agriculture, the rise of services, and the hump-shaped manufacturing path. Manufacturing peaks in 1928 at 24.93 percent of employment, compared with 28.28 percent in 1950 in the data. Table ?? reports all the calibrated parameters next to the fit associated with their principal empirical discipline. The moments most directly tied to the paper's mechanism—sectoral firm size and HSPS shares—are matched almost exactly, with the largest deviation below 1 percent.




Figure 5: Fit to U.S. structural change. Lines are model outcomes and circles are historical employment shares. The sectoral employment paths enter the calibration objective.

Four untargeted implications evaluate the calibrated transition. First, decadal GDP-per-capita growth is computed from the targeted GDP-level series but does not itself enter the objective. Across the 22 decades in Figure 4, the model matches the hump-shaped pattern from low growth in the early nineteenth century, to faster growth during the industrial expansion, and then to a gradual slowdown as the economy transitions toward services. Second, I compare the model’s household education choice with the average of primary and secondary enrollment rates over 1825–2020. The correlation between the model’s education choice and enrollment is 0.94, so the quantity–quality household block captures the secular expansion of education remarkably well despite the historical education path being untargeted. Third, the model generates substantial cross-firm dispersion in formalization and a secular increase in the dispersion of firm size, both untargeted and well-known empirical features of firm dynamics.²⁶ Finally, the model captures the rise and subsequent decline of average manufacturing firm size and the growth of the HSPS employment share even though only the terminal observations of these series are targeted.²⁷ The next subsection provides the strongest untargeted validation of the model’s novel mechanism by showing that the accumulation of public organizational knowledge in the model closely tracks the historical production of books and manuals codifying new methods of business organization and management in the data.

²⁶ See, for instance, [Poschke \(2018\)](#) for the secular rise in firm-size dispersion and [Bloom et al. \(2012\)](#) for dispersion in management practices across firms. The employment-weighted standard deviation of firm-level formalization in the model is 0.176 in agriculture, 0.187 in manufacturing, and 0.049 in services. I construct data-based proxies of 0.229, 0.205, and 0.211 using the DEFRA Farm Business Survey 2024/25 for agriculture and the ONS Management and Expectations Survey 2019 for manufacturing and services. None of these moments enters the calibration objective.

²⁷ For firm size, the initial level is normalized to its value in the data, so the untargeted validation comes from the subsequent path, in particular the decline in manufacturing firm size after the manufacturing peak.




Figure 6: Public organizational knowledge and copyright publications. Panel A compares the annual increase in the model’s public knowledge stock with business-organization publication intensity. The displayed series are centered five-year averages normalized by their own maxima; the reported correlation uses unsmoothed annual observations. Panel B reports the corresponding annual correlation for the business-organization category and eleven placebo publication categories. All publication series use identical processing and the common 1906–1977 support.

5.4 Organizational knowledge in the historical record

I assess the timing of public organizational knowledge using an independent historical series. The model’s annual increment,

$$\Delta K_t \equiv K_{t+1} - K_t, \quad (34)$$

measures additions to the stock of public knowledge concerning business organization which firms can draw on to formalize production. This object has a natural empirical counterpart: the publication of books and manuals through which methods of business organization and management became codified and publicly available. I construct this series from digitized volumes of the Library of Congress *Catalog of Copyright Entries*. Individual titles are parsed and classified using terms associated with accounting and record keeping, office and business systems, management methods, and the organization of sales, personnel, and multi-unit firms. The classification includes general Taylorist concepts such as “scientific management,” “time study,” and “motion study,” and excludes titles containing explicit manufacturing or industrial language.²⁸ This produces a measure of organizational knowledge that is not mechanically driven by references to manufacturing. Using the same underlying data source and classification method, I also construct analogous series for 11 different placebo topics ranging from children’s literature to fields which are more plausibly connected to manufacturing such as chemistry and applied physics.

Because the number and size of digitized catalog volumes vary across years, I measure publication intensity rather than raw counts. For publication category g , define

$$P_{gt} \equiv \frac{\text{matched copyright records}_{gt}}{\text{OCR words}_t}. \quad (35)$$

²⁸ For instance titles including the words “industrial, manufacturing, factory, etc.” are excluded.

All categories use the same annual support and OCR denominator. The common sample contains 72 annual observations from 1906 through 1977.

Figure 6 presents the comparison. Panel A compares ΔK_t with business-organization publication intensity. The displayed lines are centered five-year averages normalized by their respective maxima; the reported correlation uses the unsmoothed annual observations. The correlation is 0.752. Both series rise during the first half of the twentieth century, peak around mid-century, and subsequently decline.

Panel B evaluates the specificity of this relationship using eleven placebo categories spanning scientific, technical, household, and literary subjects. Business-organization publications rank first, with an annual correlation of 0.752. The closest placebo is household manuals, with a correlation of 0.486; all remaining correlations are 0.416 or lower. In one-sided comparisons using Newey–West standard errors with four annual lags, the business correlation exceeds each placebo correlation. The largest comparison p -value is 0.069. The ranking does not depend solely on the common hump: using five-year log changes, business organization remains first ($r = 0.543$, compared with 0.355 for fiction and novels), and the difference is significant using Newey–West standard errors with five lags (one-sided $p = 0.033$). During the period in which the model predicts rapid additions to public organizational knowledge, the historical record shifts disproportionately toward works about business organization. This pattern reflects the co-movement of manufacturing employment and organizational-knowledge production that appears in both the model and the data. To the best of my knowledge, this is the first study to document that the production of codified business-organization knowledge moved systematically with the rise and decline of manufacturing employment and to relate that historical comovement to the knowledge flow generated by a structural model.

6 Results

6.1 Manufacturing exposure and development outcomes

I use the calibrated model to quantify how the historical importance of manufacturing shapes subsequent development. Starting from the U.S. calibration, I vary only the manufacturing weight in final demand, w_m . This traces an equilibrium locus of otherwise identical economies with different endogenous peak manufacturing employment shares.²⁹ The eight model economies in Figure 7 span manufacturing peaks from 6.35 to 38.80 percent of employment. Larger manufacturing peaks are associated with greater tertiary education, a composition of services tilted towards HSPS employment, higher 2019 GDP per capita, and faster GDP-per-capita growth over 2010–2019. They are also associated with lower population growth over 1850–2019, as in the data. In the model, deeper industrialization raises productivity growth, which increases the return to education and lowers fertility. The model performs best on the outcomes closest to the organizational mechanism: it accounts for 98 percent of the empirical HSPS gradient and 32 percent of the tertiary-education gradient. It also accounts for 40 percent of the 2010–2019

²⁹ The experiment also has a price-wedge interpretation. Suppose the manufacturing demand weight changes from w_m to $f_m w_m$, and let $k_j \equiv \varepsilon_j / \varepsilon_m$. Holding w_m fixed, the same relative final-demand schedules are obtained by applying wedges to the prices of agriculture and consumer services such that

$$1 + \tau_a = f_m^{-k_a/(1-\sigma)}, \quad 1 + \tau_c = f_m^{-k_c/(1-\sigma)}.$$

Varying w_m is therefore isomorphic, for household final demand, to varying relative price wedges on agricultural and service consumption. The exercise can be read as a parsimonious change in the relative pull of manufacturing, rather than as a structural change in preferences. Appendix E derives this equivalence.




Figure 7: Manufacturing peaks and development outcomes. Red circles show means within 30 approximately equal-sized country bins; blue diamonds show model economies. Dashed lines are OLS fits to the underlying observations; β_M/β_D denotes the ratio of model to empirical slopes. The large red circle identifies the United States. Manufacturing peaks use the harmonized historical-employment series described in Appendix A.1. GDP levels (2019, in logs relative to the United States) and annual log growth (2010–2019) use Gapminder’s PPP GDP-per-capita series. Completed tertiary attainment averages Barro–Lee observations for 2010 and 2015. HSPS is employment in ISCO-08 groups 1–3 within ISIC Rev. 4 divisions 62–66 and 69–75, divided by employment in market services (sections G–N and R–S), using ILOSTAT data averaged over 2010–2019. Samples comprise 172 countries for GDP, 137 for education, and 48 for HSPS. U.S. manufacturing and education observations use the empirical calibration targets.

GDP-per-capita growth gradient, 17 percent of the 2019 GDP-per-capita gradient, and 17.6 percent of the population-growth gradient. Variation in manufacturing history therefore generates a substantial share of the observed cross-country differences in service composition, higher education, and growth in income per capita, while accounting for a smaller, but still substantial, share of differences in income levels and demographic change.

The cross-country accounting exercise describes differences observed around the end of the calibration period, but the model implies that the consequences of industrial history persist beyond that point. Because a deeper manufacturing phase leaves the service economy with a larger skilled share, Proposition 2 implies a permanently higher limiting productivity-growth rate. Table 4 illustrates this mechanism by comparing two otherwise identical economies with different manufacturing histories. The low-peak economy reaches a manufacturing employment share of 11.19 percent, close to the average peak among low- and middle-income economies, while the high-peak economy reaches 29.88 percent. The high-peak

economy enters the service era with a persistently larger skilled-employment share and therefore grows faster thereafter. The annual growth gap is modest—0.25 percentage points in 2030 and 0.17 points in 2070—but compounds over time: GDP per capita in the high-peak economy rises from 39 percent above the low-peak economy in 2030 to 51 percent above it by 2070. In the model, industrial history therefore affects not only the level of income reached by the end of industrialization, but also the subsequent growth path of the service economy.

Table 4: Manufacturing history and subsequent growth paths

Economy	Outcome	2030	2050	2070
High-peak economy	GDP per capita	50.36	78.43	121.80
	Skilled-employment share (%)	34.18	34.64	35.02
	GDP-per-capita growth (% p.a.)	2.24	2.22	2.20
Low-peak economy	GDP per capita	36.12	53.88	80.72
	Skilled-employment share (%)	28.74	29.37	29.83
	GDP-per-capita growth (% p.a.)	1.99	2.01	2.03
High-low comparison	GDP per capita (low peak = 100)	139.44	145.57	150.89
	Skilled-employment-share gap (pp)	5.44	5.27	5.19
	GDP-per-capita growth gap (pp p.a.)	0.25	0.21	0.17

Notes: The low-peak economy sets $w_m = 0.60$ and reaches a manufacturing-employment peak of 11.19 percent in 1941. The high-peak economy sets $w_m = 3.49$ and reaches a peak of 29.88 percent in 1927. GDP-per-capita levels equal nominal GDP divided by the consumption price index and population. The skilled-employment share is the model stock S_t . GDP-per-capita growth is the annualized log change in chain-weighted real GDP per capita over the preceding ten years. In the GDP-per-capita comparison, the low-peak economy is normalized to 100. The skill and growth comparisons report high-minus-low differences in percentage points.

6.2 Which mechanisms matter?

To understand which source of sector-heterogeneity is driving the results, I repeat the manufacturing-demand experiment after removing one source of heterogeneity at a time. Each experiment equalizes one of the following: entry barriers F_j , sectoral productivity growth ϕ_j , adaptability ξ_j , or value-added skill intensity α_j . The final experiment sets $\delta_K = 0$, eliminating the diffusion of newly formalized routines into public organizational knowledge. All other parameters remain at their calibrated U.S. values. Because each intervention changes the mapping from w_m into structural transformation, every coefficient is estimated against the counterfactual economy’s realized manufacturing peak. Table 5 shows the results.

The results in Table 5 isolate one source of cross-sector heterogeneity as the main force behind the paper’s quantitative cross-country accounting results. Equalizing F_j leaves 49.0 percent of the GDP-level slope and 40.7 percent of the growth slope, *reverses* the tertiary-attainment slope, and reduces the HSPS slope to 12.7 percent of baseline. Thus, manufacturing exposure raises the skill-share and producer-service employment primarily when sectoral entry conditions allow manufacturing firms to operate at scale. Knowledge diffusion propagates this effect. Setting $\delta_K = 0$ preserves 71.3 percent of the GDP-level slope and 67.6 percent of the growth slope, but only 26.5 percent of the tertiary slope and 47.2 percent of the HSPS slope. Private formalization still depends on manufacturing exposure through the model’s other channels, but the effect of the manufacturing legacy is substantially reduced. The population-growth relationship is driven more directly by sectoral productivity growth: equalizing ϕ_j reduces the

Table 5: Mechanism accounting for the model’s cross-country slopes

Outcome	Baseline	Equal entry barriers	Equal productivity growth	Equal adaptability	Equal VA skill intensity	No knowledge spillover
GDP-per-capita level, 2019	1.720 [100.0]	0.843 [49.0]	1.314 [76.4]	1.594 [92.7]	1.871 [108.8]	1.225 [71.3]
GDP-per-capita growth, 2010–2019	0.015 [100.0]	0.006 [40.7]	0.018 [118.7]	0.013 [89.7]	0.015 [99.8]	0.010 [67.6]
Population growth, 1850–2019	−0.678 [100.0]	−0.244 [36.0]	−0.010 [1.5]	−0.630 [93.0]	−0.832 [122.8]	−0.460 [67.9]
Tertiary attainment, 2010–2019	0.301 [100.0]	−0.116 [−38.4]	0.246 [81.8]	0.235 [77.9]	0.340 [113.0]	0.080 [26.5]
HSPS share of services, 2010–2019	0.580 [100.0]	0.074 [12.7]	0.691 [119.1]	0.485 [83.6]	0.579 [99.7]	0.274 [47.2]

Notes: Each cell reports a model slope; brackets report its percentage of the baseline slope. Outcomes match Figure 7. All columns use the converged common support $w_m \in \{1.16, 1.8, 2.56, 3.49, 4.69\}$ and realized manufacturing peaks. Counterfactuals equalize F_j , ϕ_j , or α_j across sectors; set $\xi_a = \xi_c = \xi_m = 1$; or set $\delta_K = 0$, as indicated.

population-growth slope to 1.5 percent of baseline. This is consistent with the demographic mechanism in the model, through which faster productivity growth raises the return to education and lowers fertility.

Equalizing value-added skill shares leaves all four slopes essentially unchanged or slightly larger on account of the large value-added skill-share of consumer services. Equal adaptability retains between 77.9 and 92.7 percent of the baseline slopes. Equal productivity growth reduces the GDP-level and tertiary slopes while increasing the growth and HSPS slopes. Sectoral entry conditions therefore generate the scale margin which makes manufacturing especially important and knowledge diffusion transmits its effects to the broader economy.

6.3 Robustness

The results are stable across a broad sensitivity battery. Across 18 economic variants all 90 cross-country slopes retain their benchmark sign. Median absolute changes are 2.9 percent for GDP levels, 8.5 percent for GDP growth, 10.2 percent for tertiary education, 2.4 percent for HSPS, and 4.3 percent for population growth. Excluding the two deliberately radical exercises—eliminating knowledge spillovers and replacing employment weights with equal sector weights—the largest absolute changes are 15.1, 17.5, 27.3, 16.4, and 13.7 percent, respectively. A separate exact grid over $\delta_K \in [0.01, 0.15]$ also preserves every sign. Appendix E reports the complete sensitivity results.

7 Industrial Policy: Skill Supply and Firm Scale

The policy experiments separate two margins: the supply of skills and firms’ demand for organizational capability. I consider temporary interventions that operate on one or both margins and ask a simple question:

Can temporary development policies leave a persistent effect after they are withdrawn?

By Proposition 2, a temporary policy can change the economy’s asymptotic growth rate only if it changes the skilled share to which the economy ultimately converges. In the model, the natural state variable through which such a persistent effect can arise is public organizational knowledge. A temporary intervention that induces firms to create additional organizational procedures can raise K_t permanently; once these procedures diffuse, they lower the cost of formalization for future firms, increase demand for high-skill producer-service labor, and thereby alter the equilibrium supply of skill. The effectiveness of such interventions depends critically on firm scale because the private return to installing

a reusable organizational procedure rises with the scale over which it is used. I therefore first introduce a developing-country firm-scale distortion and then study how temporary skill and organizational policies interact with its removal.

7.1 A LMIC policy laboratory

To introduce a disciplined source of small firms, I extend the entry block in the spirit of [Bento and Restuccia \(2017\)](#) and [Bento and Restuccia \(2021\)](#). Before drawing permanent ability $a \sim G$, a prospective entrant chooses establishment setup intensity $h \in [h, 1]$, which determines permanent effective productivity $x = ha$. In manufacturing and consumer services, effective productivity is subject to the correlated distortion

$$1 - \tau_{i,t} = \kappa_{j,t} x_{i,t}^{-d_j}, \quad x_{i,t} = h_{i,t} a_i, \quad (36)$$

where $d_j > 0$ implies that more productive firms face systematically larger distortions. I choose $\kappa_{j,t}$ in each sector and year so that the revenue-weighted mean of $\tau_{i,t}$ is zero. Reform therefore removes the correlation between distortions and firm productivity rather than changing the average tax burden.

The setup choice is linked directly to the benchmark convex entry block. An entrant choosing setup intensity h requires

$$e(h, d_j) = \exp\{-\lambda_E d_j (1 - h)\} \quad (37)$$

units of the existing setup input. A larger correlated distortion lowers the return to establishing a high-productivity firm, while a lower h economizes on the setup input required for entry. Entrants therefore respond to stronger distortions by choosing simpler, less productive establishments, allowing the distorted economy to sustain a larger mass of small firms even though aggregate setup-input supply remains convex. When $d_j = 0$, $e(h, 0) = 1$ for every h , so entrants choose $h = 1$, effective productivity becomes $x = a$, and the benchmark entry block is recovered exactly.³⁰

I set the excess correlated distortion to $d_m = d_s = 0.41$, corresponding to the difference between the high-distortion and U.S. benchmarks in the misallocation literature ([Bento and Restuccia, 2017, 2021](#)). The common setup-response parameter λ_E is disciplined externally by the establishment-size response to removing correlated distortions reported by [Bento and Restuccia \(2021\)](#). Conditional on these objects, I re-estimate the manufacturing and service setup-cost shifters F_m and F_c to match the remaining gaps in average firm size. I also re-estimate the productivity-growth parameters ϕ_j , the manufacturing-demand shifter ω_m , and the skill-acquisition cost c^{skill} to match, respectively, sectoral value-added productivity growth, peak manufacturing employment, and the skilled share of the median LMIC economy. Table 6 reports the target values and resulting fit.

7.2 Distortion reform creates scale, but not organizational capability

I first ask what happens when the firm-scale distortion is removed without any temporary skill or organizational policy. Beginning in 2026, I reduce d_j linearly to zero by 2035 while holding all other recalibrated

³⁰ Unlike [Bento and Restuccia \(2017, 2021\)](#), who model h as a separate costly productivity investment, I let h determine the amount of the benchmark setup input required to create the establishment. This preserves their central incentive mechanism—correlated distortions reduce the return to establishing a highly productive firm—without adding a second investment technology alongside organizational formalization. The interaction in $e(h, d)$ captures a distortion-avoidance margin through which simpler establishments can economize on the setup infrastructure required to operate at scale. Setting $d = 0$ eliminates this margin and recovers the original entry block exactly.

Table 6: Calibration of the LMIC policy economy

Block	Target moment	Data	Model
Productivity	Agricultural VA productivity growth, 1990–2018 (%)	2.620	2.620
	Manufacturing VA productivity growth, 1990–2018 (%)	2.013	2.013
	Services VA productivity growth, 1990–2018 (%)	1.111	1.111
Firm scale	Agricultural average firm size, 2019	2.220	2.103
	Manufacturing average firm size, 2019	6.238	6.212
	Services average firm size, 2019	1.883	1.907
Education	Tertiary-skilled share, 2019 (%)	12.000	11.773
Demand	Peak manufacturing employment, 1990–2018 (%)	11.713	11.114

Notes: Productivity growth and peak manufacturing employment use the GGDC/UNU–WIDER Economic Transformation Database. Firm-size targets are LMIC medians constructed from [Bento and Restuccia \(2021\)](#). Because that source does not report agriculture separately, the nonagricultural mean is used as the agricultural proxy. The 12-percent skilled-share target is bracketed by ILO employment and WDI tertiary-enrollment measures.

Table 7: Permanent distortion reform: scale without temporary policy

Outcome in 2089	Distortions remain	Full reform
Average manufacturing firm size	6.04	13.76
Average service firm size	2.13	4.89
GDP per capita relative to unreformed path (%)	0.00	7.60
Formalized task-share change (pp)	0.00	0.00
Skilled employment-share change (pp)	0.00	0.01
Public organizational knowledge, K	0.000	0.000

Notes: All entries are from no-temporary-policy paths. Full reform removes the excess productivity-correlated distortion linearly during 2026–2035. Firm-size entries are levels. GDP, formalization, and skilled employment are measured relative to the unreformed LMIC path. The institutional or administrative cost of the permanent reform is not modeled.

LMIC parameters fixed. This experiment isolates the effect of eliminating productivity-correlated distortions on firm scale and asks whether larger firms alone are sufficient to trigger the model’s organizational mechanism.

Table 7 shows that removing productivity-correlated distortions produces substantial gains in firm scale and income. Average firm size approximately doubles in both manufacturing and services, while GDP per capita is 7.6 percent higher by 2089, somewhat smaller than the gains reported by [Bento and Restuccia \(2021\)](#). Yet the formalized task share, public organizational knowledge, and skilled employment share remain essentially unchanged. The reform therefore produces a substantial level gain by allowing more productive firms to operate at larger scale, but it does not cross the organizational thresholds required to create new procedures. Scale alone is not sufficient: firms must also find formalization sufficiently valuable relative to the skilled-labor cost of developing and operating organizational procedures. The reform raises income without changing the organizational state inherited by the future service economy and therefore leaves the limiting growth rate unchanged.

7.3 Temporary policies

I next compare five temporary interventions operating from 2026 through 2045. Each is evaluated both when LMIC distortions remain and when the permanent distortion reform described above is implemented.

1. A *manufacturing wage subsidy* reimburses part of firms' skilled and unskilled manufacturing wage bill, so that firms pay $(1 - \tau_m^s)w$ while workers receive w .
2. An *HSPS wage subsidy* analogously reimburses part of the wage bill associated with high-skill producer-service employment.
3. A *college scholarship* reimburses the pecuniary component of the calibrated college cost.³¹
4. *Firm-formalization support* reimburses part of the one-time cost of creating a new organizational procedure.
5. A *mixed package* combines college scholarships and firm-formalization support, splitting the available fiscal resources equally between the two instruments.

Policy rates are chosen so that realized present-value expenditure is approximately 0.6 percent of GDP during the intervention. Firm-formalization support is capped at a 100-percent reimbursement of the creation cost. Even under full reimbursement, firms must still incur the recurring cost of employing HSPS workers to operate formalized tasks, so making procedure creation free does not imply that every firm formalizes every task. In economies where a 100-percent reimbursement cannot exhaust the 0.6-percent fiscal envelope, I report the maximum feasible expenditure rather than subsidizing firms beyond the cost of creating the procedure. All temporary programs are financed through contemporaneous lump-sum household taxes, so their resource costs enter household consumption and welfare. Consumption-equivalent variation is evaluated over 2026–2089 using the full household utility function.³²

To isolate the state created by the intervention from anticipation of its termination, I use an unexpected-withdrawal design. Through 2045, households and firms behave under the belief that the policy will continue. At the start of 2046, its removal arrives as an unanticipated policy change and I resolve the continuation equilibrium from the complete inherited state, including the firm distribution, installed organizational capital, and public organizational knowledge. The post-2046 paths therefore measure the legacy of the economic state created while the policy was active rather than behavioral responses to a previously announced termination date.

7.4 Firm scale and policy effectiveness

Table 8 reports the long-run effects of the temporary interventions. Every entry is measured relative to the no-temporary-policy path with the same permanent-distortion and entry-cost environment.

Panel A shows that persistent distortions sharply limit the economy's capacity to convert temporary policy support into lasting organizational capability. Even at a 100-percent reimbursement rate, firm-formalization support spends only 0.017 percent of present-value GDP and leaves a modest organizational legacy. The reason is that making procedure creation free does not eliminate the operating economics of formalization. When firms are small, a reusable procedure avoids relatively few instances of repeated problem solving; when skills are scarce, operating the formalized task also requires relatively expensive HSPS labor. Most firms therefore remain below the threshold at which formalization is worthwhile even when the government pays the entire creation cost. College scholarships and the mixed package

³¹ I assume that one unit of the college utility cost maps into one unit of unskilled wages. The government lowers the cost of acquiring skill through a tuition subsidy financed by a lump-sum household tax.

³² I end the welfare calculation in 2089, 200 years before the terminal horizon used to solve the model's perfect-foresight transition, so the results are not driven by behavior near the numerical terminal condition.

Table 8: Temporary policies and their long-run legacies

Temporary policy	PV cost (% GDP)	CEV (%)	GDPpc (%)	Formalized tasks (pp)	Skilled employment (pp)	ΔK (level)
<i>Panel A: LMIC distortions remain</i>						
College scholarship	0.593	2.59	5.56	0.00	0.01	0.000
Firm-formalization support	0.017*	0.18	2.13	0.27	0.13	0.036
Mixed package	0.613	6.95	8.02	0.00	0.02	0.000
Naive HSPS subsidy	0.101 [†]	−3.31	−1.51	0.10	0.04	0.009
Naive manufacturing subsidy	0.598	−3.48	−2.14	0.00	−0.01	0.000
<i>Panel B: Full permanent distortion reform</i>						
College scholarship	0.581	6.92	8.52	1.64	0.29	0.026
Firm-formalization support	0.123*	1.67	8.43	8.63	1.46	0.256
Mixed package	0.600	7.51	10.78	4.48	0.75	0.078
Naive HSPS subsidy	0.580	0.57	1.79	3.15	0.55	0.052
Naive manufacturing subsidy	0.597	−0.16	−0.42	0.00	0.00	0.000
<i>Panel C: Full reform plus U.S. entry block</i>						
College scholarship	0.579	7.17	10.62	9.40	1.41	0.069
Firm-formalization support	0.130*	2.30	15.43	26.60	4.04	0.381
Mixed package	0.585	7.77	13.40	13.98	2.10	0.140
Naive HSPS subsidy	0.620	0.70	3.82	11.02	1.63	0.086
Naive manufacturing subsidy	0.597	−0.15	−0.41	0.00	0.00	0.000

Notes: Programs operate from 2026 through 2045 and are withdrawn unexpectedly at the start of 2046. PV cost is discounted government spending divided by discounted GDP during the policy period. CEV is evaluated over 2026–2089; all other outcomes are measured in 2089. Every entry is relative to the no-temporary-policy path with the same permanent-distortion and F_j environment. Formalization is the employment-weighted formalized-task share. Panel C sets $F_m = 187.625$ and $F_s = 85.117$, their U.S.-calibration values, from the common inherited 2026 state.

The naive HSPS subsidy applies to the entire HSPS wage bill, while the naive manufacturing subsidy applies to the entire skilled and unskilled manufacturing wage bill.

*Firm-formalization support is capped at a 100-percent reimbursement rate and cannot absorb the 0.6-percent-PV-GDP target.

[†]In Panel A, the HSPS fiscal correspondence jumps from 0.101 to 2.083 percent of PV GDP over a narrow subsidy-rate interval.

can absorb the larger fiscal envelope and generate substantial welfare and income gains, but they leave essentially no additional formalization or public organizational knowledge after withdrawal. In the small-firm environment of the median LMIC economy, expanding skill supply is therefore not sufficient to create a persistent organizational legacy. Skills rise while the policy is active, but without a corresponding increase in firms' demand for formal organization the skilled share eventually converges back toward the no-policy path.

The picture changes sharply when temporary policies are combined with distortion reform. In Panel B, the larger firms created by the permanent reform raise the payoff to formalization and make both firm- and skill-side interventions more persistent. Firm-formalization support now leaves the formalized-task share 8.63 percentage points and public organizational knowledge 0.256 above the reformed no-policy path in 2089. College scholarships also generate an organizational legacy: formalization is 1.64 percentage points higher and public knowledge increases by 0.026 because larger firms make the temporary increase in skilled labor sufficient for some firms to cross the formalization threshold. The mixed package produces the largest welfare and GDP gains among policies able to spend approximately the full 0.6-percent fiscal envelope, while direct formalization support creates the largest stock of organizational capability despite being unable to exhaust that budget.

Panel C pushes the scale margin further by restoring the manufacturing and service entry-cost shifters F_j to their U.S.-calibration values from 2026 onward, in addition to eliminating the correlated distortion. The amplification is substantial. Firm-formalization support leaves formalization 26.60 percentage points and public organizational knowledge 0.381 above the corresponding no-policy path, while even the college scholarship induces a sizeable organizational response. The comparison shows that the effectiveness of temporary capability policy depends strongly on the firm-scale environment in which it is deployed: the same temporary reduction in the cost of skills or procedure creation has much larger persistent effects when firms are sufficiently large to value formal organization.

Broad wage subsidies perform considerably worse. Manufacturing wage subsidies reduce welfare in every firm-scale environment and generate essentially no organizational legacy. Their poor performance reflects, in part, the fact that they subsidize employment in a sector from which the economy is endogenously moving away from. HSPS wage subsidies can induce some formalization once firms become sufficiently large, but their welfare and capability effects remain smaller than those of more directly targeted interventions at comparable fiscal cost. More generally, wage subsidies are poorly targeted to the model's dynamic externality: they subsidize the current flow of labor services rather than the creation of reusable organizational procedures whose diffusion changes future production possibilities.

The policy implication is therefore one of complementarity and sequencing. Reforms that allow productive firms to operate at scale generate the demand-side conditions under which temporary investments in skills and organizational formalization can leave a persistent legacy. Skill policy alone is largely transitory when firms remain too small to formalize production, while scale reform alone raises income without necessarily creating organizational knowledge. The largest persistent gains arise when firms are both large enough to value formalization and able to access the skilled labor required to undertake it.

8 Conclusion

This paper asks whether economies that move directly from agriculture into services develop the same productive capabilities as economies that first pass through large-scale manufacturing. The evidence suggests that they do not. Across countries, larger historical manufacturing peaks are followed by service economies more concentrated in high-skill producer services and by faster subsequent growth. Across U.S. counties, larger manufacturing peaks are associated with larger establishments and greater managerial and technical employment within manufacturing, while industrialization induced by national growth in high-throughput manufacturing predicts greater specialization in high-skill producer services decades later. These findings point to a persistent organizational legacy of production at scale.

I formalize this mechanism in a growth model in which larger firms have stronger incentives to replace recurring organizational problems with reusable procedures, which can subsequently diffuse across firms and sectors. Calibrated to the U.S. transition, variation in manufacturing history generates almost all of the observed cross-country gradient in the HSPS share of services and substantial shares of the corresponding gradients in tertiary education and subsequent growth. Firm scale is central to the creation of organizational capability, while knowledge diffusion transmits its effects into the service economy that follows.

However, the fact that manufacturing historically generated these capabilities does not imply that manufacturing itself should be subsidized. Manufacturing wage subsidies are poorly targeted to the underlying externality, while removing productivity-correlated distortions raises firm scale and income

but, on its own, leaves formalization and public organizational knowledge essentially unchanged. Persistent gains instead arise when reforms that allow productive firms to operate at scale are combined with temporary support for skill formation and organizational formalization. Development without industrialization is possible, but economies that bypass industrialization must still find other ways to build the capabilities that large-scale manufacturing historically created.

References

- Acemoglu, Daron and David Autor**, “Skills, tasks and technologies: Implications for employment and earnings,” in “Handbook of labor economics,” Vol. 4, Elsevier, 2011, pp. 1043–1171.
- Adão, Rodrigo, Michal Kolesár, and Eduardo Morales**, “Shift-Share Designs: Theory and Inference,” *The Quarterly Journal of Economics*, 2019, 134 (4), 1949–2010.
- Asturias, Jose, Sewon Hur, Timothy J. Kehoe, and Kim J. Ruhl**, “Firm Entry and Exit and Aggregate Growth,” *American Economic Journal: Macroeconomics*, 2023, 15 (1), 48–105.
- Atkeson, Andrew and Patrick J Kehoe**, “Modeling and measuring organization capital,” *Journal of political Economy*, 2005, 113 (5), 1026–1053.
- Barro, Robert J and Jong Wha Lee**, “A new data set of educational attainment in the world, 1950–2010,” *Journal of development economics*, 2013, 104, 184–198.
- Baumol, William J**, “Macroeconomics of unbalanced growth: the anatomy of urban crisis,” *The American economic review*, 1967, 57 (3), 415–426.
- Bento, Pedro and Diego Restuccia**, “Misallocation, establishment size, and productivity,” *American Economic Journal: Macroeconomics*, 2017, 9 (3), 267–303.
- and —, “On Average Establishment Size across Sectors and Countries,” *Journal of Monetary Economics*, 2021, 117, 220–242.
- Best, Mark and Duncan Neuhauser**, “Walter A. Shewhart, 1924, and the Hawthorne Factory,” *Quality and Safety in Health Care*, 2006, 15 (2), 142–143.
- Bloom, Nicholas, Benn Eifert, Aprajit Mahajan, David McKenzie, and John Roberts**, “Does Management Matter? Evidence from India,” *Quarterly Journal of Economics*, 2013, 128 (1), 1–51.
- , **Christos Genakos, Raffaella Sadun, and John Van Reenen**, “Management Practices Across Firms and Countries,” *Academy of Management Perspectives*, 2012, 26 (1), 12–33.
- Boppart, Timo**, “Structural change and the Kaldor facts in a growth model with relative price effects and non-Gorman preferences,” *Econometrica*, 2014, 82 (6), 2167–2196.
- Borusyak, Kirill, Peter Hull, and Xavier Jaravel**, “Quasi-Experimental Shift-Share Research Designs,” *The Review of Economic Studies*, 2022, 89 (1), 181–213.
- Buera, Francisco J and Joseph P Kaboski**, “The rise of the service economy,” *American Economic Review*, 2012, 102 (6), 2540–2569.
- , —, **Richard Rogerson, and Juan I Vizcaino**, “Skill-biased structural change,” *The Review of Economic Studies*, 2022, 89 (2), 592–625.
- Chandler, Alfred D.**, *The Visible Hand: The Managerial Revolution in American Business*, Cambridge, MA: Belknap Press of Harvard University Press, 1977.

- Chen, Xilu, Guangyu Pei, Zheng Song, and Fabrizio Zilibotti**, “Tertiarization like china,” *Annual Review of Economics*, 2023, 15, 485–512.
- Choi, Jaedo and Andrei A Levchenko**, “The long-term effects of industrial policy,” *Journal of Monetary Economics*, 2025, 152, 103779.
- Clark, Gregory**, *A farewell to alms: A brief economic history of the world*, Princeton university press, 2007.
- Comin, Diego, Danial Lashkari, and Martí Mestieri**, “Structural change with long-run income and price effects,” *Econometrica*, 2021, 89 (1), 311–374.
- Cunha, Flavio and James J. Heckman**, “The Technology of Skill Formation,” *American Economic Review*, 2007, 97 (2), 31–47.
- , —, and **Susanne M. Schennach**, “Estimating the Technology of Cognitive and Noncognitive Skill Formation,” *Econometrica*, 2010, 78 (3), 883–931.
- de Vries, Gaaitzen, Linda Arfelt, Dorothea Drees, Mareike Godemann, Calumn Hamilton, Bente Jessen-Thiesen, Ahmet Ihsan Kaya, Hagen Kruse, Emmanuel Mensah, and Pieter Wolter**, “The Economic Transformation Database (ETD): Content, Sources, and Methods,” WIDER Technical Note 2/2021, UNU-WIDER 2021.
- Duernecker, Georg and Miguel Sanchez-Martinez**, “Structural change and productivity growth in Europe—Past, present and future,” *European Economic Review*, 2023, 151, 104329.
- , **Berthold Herrendorf, and Ákos Valentinyi**, “Structural Change within the Services Sector and the Future of Cost Disease,” *Journal of the European Economic Association*, 2024, 22 (1), 428–473.
- Eckert, Fabian, Ka leung Lam, Atif R. Mian, Karsten Müller, Rafael Schwalb, and Amir Sufi**, “The Early County Business Pattern Files: 1946–1974,” NBER Working Paper 30578, National Bureau of Economic Research 2022.
- , **Teresa C. Fort, Peter K. Schott, and Natalie J. Yang**, “Imputing Missing Values in the US Census Bureau’s County Business Patterns,” NBER Working Paper 26632, National Bureau of Economic Research 2020.
- Engbom, Niklas, Hannes Malmberg, Tommaso Porzio, Federico Rossi, and Todd Schoellman**, “Economic development according to chandler,” Technical Report, National Bureau of Economic Research 2025.
- Fair, Ray C. and John B. Taylor**, “Solution and Maximum Likelihood Estimation of Dynamic Nonlinear Rational Expectations Models,” *Econometrica*, 1983, 51 (4), 1169–1185.
- Fan, Tianyu, Michael Peters, and Fabrizio Zilibotti**, “Growing like India—the unequal effects of service-led growth,” *Econometrica*, 2023, 91 (4), 1457–1494.
- Fernald, John G.**, “Productivity and Potential Output before, during, and after the Great Recession,” *NBER Macroeconomics Annual*, 2015, 29 (1), 1–51.

- Flesher, Tonya K.**, “James O. McKinsey,” *The Accounting Historians Journal*, 1985, 12 (2), 1–22.
- Fujiwara, Ippei and Kiminori Matsuyama**, “A technology-gap model of ‘prema-
ture’deindustrialization,” *American Economic Review*, 2024, 114 (11), 3714–3745.
- Galor, Oded**, “From stagnation to growth: unified growth theory,” *Handbook of economic growth*, 2005, 1, 171–293.
- **and David N Weil**, “Population, technology, and growth: From Malthusian stagnation to the demo-
graphic transition and beyond,” *American economic review*, 2000, 90 (4), 806–828.
- Giorcelli, Michela**, “The Long-Term Effects of Management and Technology Transfers,” *American Eco-
nomic Review*, 2019, 109 (1), 121–152.
- Goldsmith-Pinkham, Paul, Isaac Sorkin, and Henry Swift**, “Bartik Instruments: What, When, Why,
and How,” *American Economic Review*, 2020, 110 (8), 2586–2624.
- Graham, Laurel D.**, “Domesticating Efficiency: Lillian Gilbreth’s Scientific Management of Homemak-
ers, 1924–1930,” *Signs*, 1999, 24 (3), 633–675.
- , “Lillian Gilbreth and the Mental Revolution at Macy’s, 1925–1928,” *Journal of Management History*,
2000, 6 (7), 285–305.
- Graham, Laurel Diane**, “Lillian Moller Gilbreth’s Extensions of Scientific Management into Women’s
Work, 1924–1935.” Phd dissertation, University of Illinois at Urbana-Champaign 1992.
- Greenstein, Shane M.**, “Did Computer Technology Diffuse Quickly? Best and Average Practice in Main-
frame Computers, 1968–1983,” NBER Working Paper 4647, National Bureau of Economic Research
1994.
- Gutiérrez, Germán, Callum Jones, and Thomas Philippon**, “Entry costs and aggregate dynamics,”
Journal of Monetary Economics, 2021, 124, S77–S91.
- Haines, Michael R., Inter university Consortium for Political, and Social Research**, “Historical,
Demographic, Economic, and Social Data: The United States, 1790–2002,” Inter-university Consortium
for Political and Social Research 2010. ICPSR 02896, Version 3.
- Heathcote, Jonathan, Kjetil Storesletten, and Giovanni L Violante**, “The macroeconomic implica-
tions of rising wage inequality in the United States,” *Journal of political economy*, 2010, 118 (4), 681–722.
- Herrendorf, Berthold, Richard Rogerson, and Akos Valentinyi**, “Two perspectives on preferences
and structural transformation,” *American Economic Review*, 2013, 103 (7), 2752–2789.
- , — , **and** — , “Growth and structural transformation,” *Handbook of economic growth*, 2014, 2, 855–941.
- Hounshell, David A.**, *From the American System to Mass Production, 1800–1932: The Development of
Manufacturing Technology in the United States* number 4. In ‘Studies in Industry and Society’, Baltimore,
MD: Johns Hopkins University Press, 1984.
- Hsieh, Chang-Tai and Esteban Rossi-Hansberg**, “The industrial revolution in services,” *Journal of
Political Economy Macroeconomics*, 2023, 1 (1), 3–42.

- Huneus, Federico and Richard Rogerson**, “Heterogeneous paths of industrialization,” *Review of Economic Studies*, 2024, 91 (3), 1746–1774.
- Incoronato, Lorenzo and Salvatore Lattanzio**, “Place-based industrial policies and local agglomeration in the long run,” *Bank of Italy Temi di Discussione (Working Paper) No*, 2025, 1503.
- Johnson, H. Thomas**, “Management Accounting in an Early Integrated Industrial: E. I. duPont de Nemours Powder Company, 1903–1912,” *Business History Review*, 1975, 49 (2), 184–204.
- Jones, Charles I.**, “R&D-Based Models of Economic Growth,” *Journal of Political Economy*, 1995, 103 (4), 759–784.
- Jones, Charles I.**, “The end of economic growth? Unintended consequences of a declining population,” *American Economic Review*, 2022, 112 (11), 3489–3527.
- Jr., Alfred D. Chandler**, *Scale and Scope: The Dynamics of Industrial Capitalism*, Cambridge, MA: Harvard University Press, 1990.
- Juhász, Réka**, “Temporary protection and technology adoption: Evidence from the napoleonic blockade,” *American Economic Review*, 2018, 108 (11), 3339–3376.
- Kaplan, Robert S.**, “The Evolution of Management Accounting,” *The Accounting Review*, 1984, 59 (3), 390–418.
- Kongsamut, Piyabha, Sergio Rebelo, and Danyang Xie**, “Beyond balanced growth,” *The Review of Economic Studies*, 2001, 68 (4), 869–882.
- Kruse, Hagen, Emmanuel Mensah, Kunal Sen, and Gaaitzen De Vries**, “A manufacturing (re) naissance? Industrialization in the developing world,” *IMF Economic Review*, 2022, 71 (2), 439.
- Krusell, Per, Lee E Ohanian, José-Víctor Ríos-Rull, and Giovanni L Violante**, “Capital-skill complementarity and inequality: A macroeconomic analysis,” *Econometrica*, 2000, 68 (5), 1029–1053.
- Kydland, Finn E. and Edward C. Prescott**, “Time to Build and Aggregate Fluctuations,” *Econometrica*, 1982, 50 (6), 1345–1370.
- Lagerlöf, Nils-Petter**, “The Galor–Weil model revisited: A quantitative exercise,” *Review of Economic dynamics*, 2006, 9 (1), 116–142.
- Levey, Stanley and E. R. Jennings**, “The Use of Control Charts in the Clinical Laboratory,” *American Journal of Clinical Pathology*, 1950, 20 (11), 1059–1066.
- Mansfield, Edwin**, “Technical change and the rate of imitation,” *Econometrica: Journal of the Econometric Society*, 1961, pp. 741–766.
- McFadden, Daniel**, “Conditional Logit Analysis of Qualitative Choice Behavior,” in Paul Zarembka, ed., *Frontiers in Econometrics*, New York: Academic Press, 1974, pp. 105–142.
- McKenna, Christopher D.**, *The World’s Newest Profession: Management Consulting in the Twentieth Century* Cambridge Studies in the Emergence of Global Enterprise, Cambridge: Cambridge University Press, 2006.

- McKinsey, James O.**, *Budgetary Control*, New York: The Ronald Press Company, 1922.
- , *Managerial Accounting*, Chicago: University of Chicago Press, 1924.
- Mol, Michael J. and Julian Birkinshaw**, “The Role of External Involvement in the Creation of Management Innovations,” *Organization Studies*, 2014, 35 (9), 1287–1312.
- Ngai, L Rachel and Christopher A Pissarides**, “Structural change in a multisector model of growth,” *American economic review*, 2007, 97 (1), 429–443.
- Peretto, Pietro F. and Simone Valente**, “Growth on a Finite Planet: Resources, Technology and Population in the Long Run,” *Journal of Economic Growth*, 2015, 20 (3), 305–331.
- Peters, Michael, Youdan Zhang, and Fabrizio Zilibotti**, “Skipping the Factory: Service-Led Growth and Structural Transformation in the Developing World,” Technical Report, National Bureau of Economic Research 2026.
- Pollard, Sidney**, *The Genesis of Modern Management: A Study of the Industrial Revolution in Great Britain*, Cambridge, MA: Harvard University Press, 1965.
- Poschke, Markus**, “The firm size distribution across countries and skill-biased change in entrepreneurial technology,” *American Economic Journal: Macroeconomics*, 2018, 10 (3), 1–41.
- Prescott, Edward C and Michael Visscher**, “Organization capital,” *Journal of political Economy*, 1980, 88 (3), 446–461.
- Rodrik, Dani**, “Premature Deindustrialization,” *Journal of Economic Growth*, 2016, 21 (1), 1–33.
- Ruggles, Steven, Matt A. Nelson, Matthew Sobek, Catherine A. Fitch, Ronald Goeken, J. David Hacker, Evan Roberts, and J. Robert Warren**, “IPUMS Ancestry Full Count Data: Version 4.0,” 2024. Dataset.
- , **Sarah Flood, Matthew Sobek, Daniel Backman, Grace Cooper, Julia A. Rivera Drew, Stephanie Richards, Renae Rodgers, Jonathan Schroeder, and Kari C. W. Williams**, “IPUMS USA: Version 16.0,” 2025. Dataset.
- Schroeder, Jonathan, David Van Riper, Steven Manson, Grace Cooper, Zachary Krause, Tracy Kugler, Tsu Zhu, and Steven Ruggles**, “IPUMS National Historical Geographic Information System: Version 21.0,” 2026. Dataset.
- Shewhart, Walter A.**, *Economic Control of Quality of Manufactured Product*, New York: D. Van Nostrand Company, 1931.
- Sposi, Michael, Kei-Mu Yi, and Jing Zhang**, “The Role of Trade in Premature Deindustrialization and Industry Polarization,” *Econometrica*, 2026, 94 (5).
- U.S. Bureau of Labor Statistics**, “BLS Completes Release of Expanded Historical NAICS Basis Data,” Quarterly Census of Employment and Wages notice 2004. Last modified October 6, 2004; accessed September 11, 2026.

Westgard, James O., Patricia L. Barry, Marian R. Hunt, and Torgny Groth, “A Multi-Rule Shewhart Chart for Quality Control in Clinical Chemistry,” *Clinical Chemistry*, 1981, 27 (3), 493–501.

Yates, JoAnne, *Control Through Communication: The Rise of System in American Management*, revised edition ed., Vol. 6 of *Studies in Industry and Society*, Baltimore, MD: Johns Hopkins University Press, 1993.

A Data Sources and Variable Construction

This appendix describes the data used in the cross-country and U.S. county analyses. It defines the variables at the level at which they enter the empirical exercises and records the harmonizations required to make them comparable over time. The data used only for the quantitative model are described with the calibration in Appendix E.

A.1 Cross-country evidence

The two cross-country facts in Section 2 require different variables and therefore use complementary samples. The service-composition exercise requires a more detailed sectoral breakdown and uses the balanced Economic Transformation Database (ETD) panel. The growth exercise requires longer manufacturing histories but not the detailed service breakdown, and therefore uses the broader combined GGDC–ILOSTAT sample.

A.1.1 Detailed sectoral employment and the composition of services

The service-composition evidence uses the GGDC/UNU–WIDER Economic Transformation Database described by [de Vries et al. \(2021\)](#) and [Kruse et al. \(2022\)](#). The ETD reports annual persons engaged and value added in 12 sectors for 51 economies from 1990 through 2018. I use the employment series. The sectoral categories are agriculture, mining, manufacturing, utilities, construction, trade, transport, business services, finance, real estate, government services, and other services.

I define high-skill producer-service (HSPS) employment as employment in ETD business services, finance, and real estate. Generic market services comprise trade and other services. The two outcomes plotted in Figure 1 are

$$HSPS_{ct} = 100 \times \frac{EMP_{ct}^{Business} + EMP_{ct}^{Finance} + EMP_{ct}^{RealEstate}}{EMP_{ct}^{Total}}, \quad (38)$$

$$GenericServices_{ct} = 100 \times \frac{EMP_{ct}^{Trade} + EMP_{ct}^{Other}}{EMP_{ct}^{Total}}. \quad (39)$$

Both therefore have total employment as their denominator. Trade includes accommodation and food services, while transport and government services are excluded from both numerators.

For each of the 51 ETD economies, I construct the historical manufacturing peak using the full available employment history obtained by combining the GGDC 10-Sector Database with the ETD, giving the ETD series precedence in overlapping years. Countries are then divided at the median historical manufacturing-employment peak. The above-median group includes the country at the median and contains 26 economies; the below-median group contains 25.

For the dynamic comparison, I subtract each country’s own 1990 employment share before computing unweighted group means. The shaded bands in Figure 1 report the group mean plus or minus 1.96 cross-country standard errors.

A.1.2 Income, growth, population, education, and institutions

PPP real GDP-per-capita and population are from Gapminder’s harmonized series, which combines World Bank estimates for the modern period with historical income series. The growth outcome is average annual log GDP-per-capita growth over 2010–2019,

$$g_c = \frac{100}{9} [\log y_{c,2019} - \log y_{c,2010}]. \quad (40)$$

The corresponding service-employment regressor is the mean ILOSTAT service share over 2010–2019. Observations after 2019 are excluded from this exercise. Income-quintile fixed effects use terminal GDP-per-capita, and years since the manufacturing peak equal $2019 - P_c$.

Institutional quality is the percentile-rank measure of control of corruption from the World Bank Worldwide Governance Indicators. Initial population and initial GDP-per-capita used in the service-composition robustness regressions are measured in 1990. The initial tertiary-education control is the World Bank gross tertiary-enrollment rate measured in, or at the closest available year to, 1990. The separate education figure uses completed tertiary attainment—not enrollment—from the Barro–Lee educational attainment data ([Barro and Lee, 2013](#)). Barro–Lee observations occur at five-year intervals; the long-difference exercise retains countries with at least five observations and at least ten years between the first and last jointly observed education and income years after 1960.

A.2 U.S. county evidence

This subsection documents the sources, harmonizations, and robustness exercises underlying Section 3. The analysis combines Census of Manufactures tabulations, population-census microdata, County Business Patterns (CBP), and the Quarterly Census of Employment and Wages (QCEW). The Census of Manufactures supplies manufacturing employment and establishments; the population censuses supply classified employment and the occupations performed inside manufacturing; Census and CBP data trace HSPS service employment over the long event-study panel; and QCEW supplies the independently constructed 2000 service-composition outcome in Tables 1 and 2. Unless stated otherwise, I aggregate additive employment counts to fixed geography before forming ratios.

A.2.1 Fixed county geography

Historical county identifiers and boundary crosswalks are from IPUMS NHGIS ([Schroeder et al., 2026](#)). For historical censuses and Census of Manufactures tabulations, I assign each source-year county to the 1850 county with which it has the greatest geographic overlap and sum additive counts within the resulting fixed county. This assignment is appropriate for historical microdata because individuals cannot be located below their reported county of residence. For modern establishment data, I allocate source-county counts across fixed 1850 counties using polygon-overlap shares. Numerators and denominators are mapped separately and the ratio is calculated only after aggregation. The procedure never interpolates an employment share directly.

A.2.2 Classified employment and the manufacturing peak

County manufacturing employment, establishments, wage bills, population, and urban population come from the historical county files assembled by [Haines et al. \(2010\)](#), which compile the published population censuses and Census of Manufactures. Manufacturing employment is the Census of Manufactures average number of wage earners. Its denominator is total classified county employment: employed residents whose industries can be assigned consistently to primary industries, manufacturing, or services. Using harmonized IND1950, the three groups are

$$\begin{aligned} \text{primary industries: } & \{105, 116, 126\}, \\ \text{manufacturing: } & 306\text{--}499, \\ \text{services: } & 506\text{--}946. \end{aligned}$$

Mining, construction, unclassified employment, and industries outside these groups are excluded. This denominator provides the common economic base for peak manufacturing exposure and manufacturing HSPS.

For county c in year t , define

$$M_{ct} = \frac{L_{ct}^M}{L_{ct}^C}, \quad (41)$$

where L_{ct}^M is manufacturing employment and L_{ct}^C is classified employment. The manufacturing peak and its date are

$$M_c^{\text{peak}} = \max_{t \in \{1880, 1900, 1920, 1930, 1940\}} M_{ct}, \quad P_c = \min\{t : M_{ct} = M_c^{\text{peak}}\}. \quad (42)$$

The earliest year determines the peak when the maximum is attained more than once. The 1890 population census was largely destroyed and cannot supply the classified-employment denominator; the required county manufacturing tabulations are unavailable for 1910. Neither year enters the peak search. Manufacturing establishment scale is the logarithm of manufacturing workers per establishment in P_c . Peak exposure, establishment scale, and manufacturing HSPS therefore refer to the same county-year.

The alternative-channel variables in column 5 of Table 1 are also measured in P_c . Average manufacturing wages equal the manufacturing wage bill divided by average wage earners; the urban share is urban population divided by county population; and county size is log population. Peak-cohort fixed effects absorb common differences in nominal wages across census years. Wage bills are unavailable for 48 counties in the 1940 peak cohort. These observations receive the peak-cohort median and the regression includes a replacement indicator.

A.2.3 HSPS capability inside manufacturing

Historical occupation-by-industry counts are constructed from the IPUMS Ancestry Full Count censuses through 1940 and harmonized IPUMS USA files thereafter ([Ruggles et al., 2024, 2025](#)). I use all employed individuals rather than restricting the sample by household role. Full-count records enter as counts and sample-year records enter with their person weights. Industries and occupations are harmonized using IND1950 and OCC1950.

Manufacturing HSPS is occupation based. It counts workers performing the accounting, coordination, personnel, technical, and scientific functions in Table 9, conditional on employment in IND1950

industries 306–499. The county-year measure is

$$MHSPS_{ct} = 100 \times \frac{\text{core HSPS-occupation workers in manufacturing}_{ct}}{L_{ct}^C}. \quad (43)$$

This is the definition used in the event study and the channel-accounting table. The occupation set is fixed over time and is defined by the function performed inside the manufacturing firm, not by educational attainment in a particular census year.

Table 9: Occupation-based manufacturing HSPS classification

Function	Exact occupation content	OCC1950 codes
Accounting and record keeping	Accountants and auditors; bookkeepers; bill and account collectors; office-machine operators; shipping and receiving clerks	0, 310, 321, 341, 342
Purchasing, logistics, and coordination	Buyers and department heads in stores; buyers and shippers of farm products; purchasing agents and buyers not elsewhere classified; dispatchers and starters of vehicles	200, 201, 280, 322
Personnel administration	Personnel and labor-relations workers	72
Technical and scientific work	Draftsmen; engineers by specialty; statisticians and actuaries; surveyors; testing technicians; technicians not elsewhere classified; inspectors not elsewhere classified	35, 41–49, 83, 92, 95, 96, 533

Notes: Manufacturing is defined by IND1950 codes 306–499. The four groups jointly form the core manufacturing-HSPS measure used throughout the main analysis. Legal occupations and commercial intermediaries are not included.

A.2.4 HSPS service industries and the long event-study panel

The service-side measure is industry based. In modern data, HSPS services comprise Information; Professional, Scientific, and Technical Services; and Management of Companies and Enterprises—NAICS sectors 51, 54, and 55. These sectors supply codified information, professional and technical expertise, and organizational direction. NAICS 54 contains specialized accounting, legal, engineering, scientific, technical, and consulting suppliers; NAICS 55 captures the management and coordination of complex enterprises; and NAICS 51 captures the production and transmission of information. I exclude NAICS 56 because its two-digit aggregate combines business support with routine administrative, cleaning, security, and waste-management activities. I also exclude the broad NAICS 52–53 aggregates because they contain substantial financial intermediation, insurance distribution, property operation, and rental activity outside the organizational and technical core studied here.

The manufacturing and service measures are analogues by economic function rather than mechanically identical classifications. Occupations identify capabilities performed inside manufacturing firms; service industries identify specialized establishments through which related information, technical, and organizational functions are supplied on the market. This grouping is also the narrow U.S. analogue of the business-and-professional component of the cross-country HSPS measure in Figure 1.

The event study asks when HSPS activity becomes quantitatively important in the county economy, so its service numerator is divided by classified employment:

$$SHSPS_{ct} = 100 \times \frac{\text{employment in HSPS service industries}_{ct}}{L_{ct}^C}. \quad (44)$$

The service numerator comes from population-census industry records through 1950, imputed early CBP in 1960 and 1970, and the harmonized, suppression-imputed CBP panel of [Eckert et al. \(2020\)](#) thereafter. EFSY is useful for the event study because it supplies a long postwar series on a common industry classification and imputes confidential CBP cells subject to the adding-up restrictions in the CBP hierarchy. Table 10 records the concordance.

Table 10: HSPS service-industry measurement by data vintage

Event-study year	Source and classification	HSPS service-industry numerator
1850–1950	Population-census microdata; harmonized IND1950	Communications, advertising, accounting, business, legal, engineering, and professional and technical services: codes 578, 579, 806, 807, 808, 856, 857, 879, 898, and 899.
1960	Imputed 1959 CBP; SIC	Communications, business services, legal services, and engineering, accounting, research, and related professional services: SIC groups 4800, 7300, 8100, and 8900.
1970	Imputed 1970 CBP; SIC	The same SIC analogue used for 1960.
1980–2000	EFSY harmonized and imputed CBP; common 2012 NAICS	All mutually exclusive EFSY employment cells beginning with NAICS 51, 54, or 55.

Notes: The early CBP files are from [Eckert et al. \(2022\)](#); the suppression imputations and common-NAICS panel are from [Eckert et al. \(2020\)](#). The EFSY long panel stores employment in mutually exclusive cells at the most detailed assignable industry level, so residual aggregate cells do not duplicate their descendants. Management companies are not separately identified in the historical IND1950 and SIC classifications, and those years use the closest functional analogue. IND1950 code 808 includes some activities later assigned to NAICS 56 that cannot be separated in the historical classification.

A.2.5 The QCEW long-run service-composition outcome

Tables 1 and 2 ask whether the service sector becomes more concentrated in HSPS industries. The numerator retains the modern HSPS grouping, while the denominator is total private-service employment. I construct the 2000 endpoint independently from the public BLS QCEW files ([U.S. Bureau of Labor Statistics, 2004](#)). QCEW is particularly useful for this cross section because it directly publishes a private service-providing aggregate at the county level. By contrast, the EFSY panel contains mutually exclusive CBP industry cells, so its service denominator must be reconstructed by summing imputed components.

For county c in year t , define

$$Y_{ct}^{QCEW} = \frac{E_{c,51,t}^{\text{public}} + E_{c,54,t}^{\text{public}} + E_{c,55,t}^{\text{public}}}{E_{c,102,t}^{\text{public}} - E_{c,22,t}^{\text{public}}}, \quad (45)$$

where ownership code 5 restricts all components to private employment, code 102 is QCEW’s service-providing aggregate, and code 22 is private utilities. The denominator includes wholesale and retail trade, private transportation and warehousing, information, finance, real estate, professional and business services, private education and health, leisure, hospitality, and other private services. Private ownership excludes public administration and the U.S. Postal Service.

The main-text numerator is the sum of all public numeric records in NAICS 51, 54, and 55. This convention is observable and reproducible, but it is not an estimate of the latent employment in confidential cells. In the reconstructed 2000 file, a detailed record that fails the disclosure screen is absent; in later files, an explicitly suppressed record carries an N flag and a zero numeric field. Thus, an unpublished contribution enters the public-record sum at zero. The disclosed aggregate denominator still contains that employment. Section A.5.4 uses the first explicit-flag year to replace this convention with a QCEW-specific hierarchical imputation, and Section A.5.5 compares it with EFSY on county and commuting-zone geographies.

The unweighted mean of the main outcome is 6.36 percent because each county receives equal weight and disclosure is most consequential in small county-industry cells. Weighting the same county ratios by private-service employment produces 13.97 percent. Aggregating all available county public records before forming the ratio gives 14.02 percent, compared with 14.58 percent from the corresponding official national QCEW numerator and denominator. The close aggregate comparison indicates that most of the national HSPS employment is observed and that the disclosure gap is concentrated among small counties. It does not by itself establish that the county-level measurement error is innocuous, which is why the flag-aware and commuting-zone exercises below are central robustness checks.

A.2.6 Historical controls and estimation samples

The baseline county covariates are measured in 1850 and come from the Haines files: log population, the urban population share, the school-enrollment rate, the number of colleges, and log enslaved population plus one. The event study additionally includes predetermined 1870 employment shares in process, machinery-producing, and traditional manufacturing. Each baseline covariate is interacted with a linear calendar-time trend in the event-study specification.

The event-study sample contains 1,116 fixed-geography counties peaking in 1900, 1920, or 1940. These cohorts supply a common forty-year pre-period and have the treatment and baseline controls observed. Event time -40 is omitted. Outcome availability varies across county-event-time cells because the manufacturing and service series draw on different source years. Confidence intervals use a spatial covariance estimator allowing arbitrary covariance among observations whose county centroids lie within 100 kilometers.

Tables 1 and 2 use the same 1,487 counties. The common sample requires the QCEW 2000 outcome, the manufacturing peak, peak-year establishment scale and manufacturing HSPS, the 1850 covariates, county coordinates, all three shift-share measures, and the complete pre-1880 histories used in the IV specifications. Holding this sample fixed ensures that differences across columns do not reflect changes in county composition. Cross-sectional inference uses Conley standard errors with a 100-kilometer cutoff.

A.3 The Chandler throughput shift-share

The shift-share design begins with county-by-manufacturing-industry employment in the 1880 full-count Census of Population, harmonized to IND1950. Workers with unclassifiable industries do not enter the county-industry counts. Following [Chandler \(1990\)](#), I distinguish industries according to the production regime through which national expansion occurred. High-throughput industries are those in which continuous-process or large-batch production, or the fabrication and assembly of complex standardized machinery, required large fixed investments and extensive managerial coordination. Traditional low-throughput industries comprise wood products, textiles and apparel, printing, and leather, where

increasing output remained more readily divisible across smaller establishments. Table 11 reports the concordance and source for every broad industry. The all-industry instrument uses every classifiable manufacturing industry in IND1950 codes 306–499.

Table 11: Manufacturing industries in the Chandler throughput design

Group	Broad industry	IND1950 codes	Chandler source and classification rationale
High throughput	Glass and cement	316–317	Chandler (1990, pp. 20, 113–120) . Chandler identifies plate and flat glass as scale intensive and describes continuous-process technologies and substantial scale economies in glass and cement.
High throughput	Primary iron, steel, and nonferrous metals	336–338	Chandler (1990, pp. 20, 23, 70–71, 121–140) . Refining and smelting, including the Bessemer and open-hearth processes and electrolytic refining, required capital-intensive plants operating at high throughput.
High throughput	Nonelectrical machinery	356–358	Chandler (1990, pp. 23, 65–68, 194–205) . Agricultural, office, and industrial machinery exemplify complex standardized machinery fabricated and assembled from interchangeable parts and supported by specialized distribution and service organizations.
High throughput	Electrical machinery	367	Chandler (1990, pp. 20, 68–69, 212–220) . Electrical-equipment producers concentrated production in large plants and required technically trained organizations to design, manufacture, install, and service complex equipment.
High throughput	Transportation equipment	376–379	Chandler (1990, pp. 20, 205–212) . These codes concord to automobiles, aircraft, ships, and railroad equipment. Production required large-scale fabrication and assembly and extensive managerial coordination.
High throughput	Photographic equipment	387	Chandler (1990, pp. 20, 67) . Chandler identifies mass-produced cameras as scale intensive and discusses Eastman Kodak’s mass production of cameras and photographic film.

Continued on next page

Table 11 – continued from previous page

Group	Broad industry	IND1950 codes	Chandler source and classification rationale
High throughput	Processed foods	406–409	Chandler (1990, pp. 23, 62–65) . Meatpacking, dairy processing, canning, and grain milling relied on mechanized processing and packaging, high-volume plants, and coordinated purchasing, distribution, and national marketing.
High throughput	Beverages	418	Chandler (1990, pp. 23, 64–65) . Chandler places brewing and distilling among refining and processing industries and emphasizes the production and refrigerated-distribution networks constructed by national brewers.
High throughput	Tobacco	429	Chandler (1990, pp. 20, 23, 63–65) . Mechanized processing and packaging permitted high-volume cigarette production for national markets.
High throughput	Pulp and paper mills	456	Chandler (1990, pp. 20, 112–113) . Chandler identifies newsprint and describes pulp, newsprint, and paperboard as acquiring substantial scale economies after the transition to capital-intensive machinery.
High throughput	Chemicals, synthetic fibers, drugs, and paint	466–469	Chandler (1990, pp. 20, 23–26, 69–70, 170–193) . Chemical production is a canonical continuous- and large-batch activity in which interdependent processes generated economies of scale and scope.
High throughput	Petroleum refining and coal products	476–477	Chandler (1990, pp. 20, 23–26, 94–104) . Petroleum refining is Chandler’s canonical illustration of throughput economies: large sunk investments were profitable only when managerial coordination maintained continuous flows from inputs through refining and distribution.
High throughput	Rubber products	478	Chandler (1990, pp. 20, 105–111) . Chandler singles out tires and documents the large-batch processing, integrated distribution, and organizational investments of leading rubber producers.

Continued on next page

Table 11 – continued from previous page

Group	Broad industry	IND1950 codes	Chandler source and classification rationale
Traditional low throughput	Logging, wood products, and furniture	306–309	Chandler (1990, pp. 20, 22–23) . Chandler places lumber and furniture among traditional labor-intensive industries in which larger establishments obtained limited cost advantages over smaller ones.
Traditional low throughput	Textiles and apparel	436–439, 446, 448–449	Chandler (1990, pp. 20, 22–23) . Textiles from natural fibers and apparel are labor intensive: increasing output primarily required adding machines and workers rather than exploiting large throughput economies.
Traditional low throughput	Printing and publishing	459	Chandler (1990, pp. 20, 22–23) . Chandler includes printing and publishing among labor-intensive industries in which scale advantages were present but not substantial.
Traditional low throughput	Leather and footwear	487–489	Chandler (1990, pp. 20, 105–106) . Chandler contrasts concentrated rubber footwear with leather footwear, where scale economies remained limited and production was dispersed across establishments.

Notes: Chandler states the general distinction on pp. 22–24: high-throughput industries used capital-intensive continuous or large-batch processes, or fabricated and assembled complex standardized machinery, whereas traditional industries remained comparatively labor intensive and obtained smaller cost advantages from greater plant size. Chandler’s categories are expressed in the historical classifications used in *Scale and Scope*; the IND1950 cells are the paper’s concordance to those categories. The baseline retains only cells that can be matched unambiguously. Ambiguous adjacent industries are excluded from both groups. Page references use the printed pagination of [Chandler \(1990\)](#).

For production regime $g \in \{A, H, L\}$, let $L_{cj,1880}$ denote county c ’s employment in manufacturing industry j in 1880 and define

$$\omega_{cj,1880}^g = \frac{L_{cj,1880}}{\sum_{\ell \in g} L_{c\ell,1880}}. \quad (46)$$

Here A denotes all manufacturing industries, H the high-throughput group, and L the traditional low-throughput group. The leave-county-out national shift is

$$G_{j,-c,1880-1900} = \frac{L_{-c,j,1900}}{L_{-c,j,1880}} - 1, \quad (47)$$

and the instrument is

$$Z_c^g = \sum_{j \in g} \omega_{cj,1880}^g G_{j,-c,1880-1900}. \quad (48)$$

Within-group normalization removes mechanical variation in the initial size of a production regime from the high- and low-throughput shifters. The instruments compare counties according to the industry composition to which they were initially exposed, rather than simply according to whether high-throughput manufacturing was initially large. All three instruments use the same base year, horizon, leave-county-out rule, and normalization and are standardized within the estimation sample.

The 1880 base predates the observed manufacturing peak for 89 percent of counties and coincides with the earliest observed peak for the remainder. Removing county c from each national shift prevents its own later employment growth from entering its instrument mechanically. The preferred specification includes state and peak-cohort fixed effects; the five 1850 covariates defined above; and complete 1850, 1860, and 1870 histories of manufacturing employment and manufacturing HSPS relative to classified employment and of HSPS services relative to private-service employment. The identifying assumption is conditional: after accounting for location, initial development, and these observed pre-1880 trajectories, differential national growth across the industries represented in a county's 1880 manufacturing mix affects service composition through the local manufacturing expansion it induces.

A.4 Treatment of missing, confidential, and structurally zero observations

I distinguish a structural zero from an unpublished positive count. In full-count Census tabulations, the absence of workers in an otherwise observed county-industry cell contributes zero. In QCEW files with disclosure flags, an N-coded record is confidential and its numeric employment field is not treated as an observed zero in the flag-aware exercises below. An absent county-industry macro record is treated as a structural absence under the BLS publication convention. I do not replace an unavailable county outcome or denominator with zero. Additive counts are aggregated before ratios are calculated, and a ratio is missing whenever its denominator is unavailable or nonpositive.

The main-text 2000 QCEW numerator is deliberately a public-record measure: it sums the numeric values that BLS releases and makes no model-based allocation to omitted detailed cells. The EFSY event-study series instead uses constrained CBP imputations. The new QCEW robustness below occupies the middle ground. It preserves every disclosed QCEW count, uses explicit flags to identify unknown cells, and allocates only the residual required by disclosed QCEW parent totals. The only sample-preserving imputation among historical regressors is the peak-cohort median replacement for 48 unavailable peak-year manufacturing wages, accompanied by its own indicator.

A.5 Robustness of the U.S. county evidence

This section evaluates the measurement and identifying assumptions behind the U.S. county results. I first vary the QCEW denominator, terminal year, and manufacturing-peak window. I then examine contemporaneous population and urbanization in the event study and report Rotemberg weights, leave-one-industry-out estimates, alternative Chandler concordances, and exposure normalizations. The final exercises reconstruct confidential QCEW employment using the disclosed QCEW hierarchy, compare it directly with EFSY's CBP-based imputation at both county and commuting-zone levels, organize inference at the level of the national industry shocks, and vary the conditioning set and spatial cutoff.

Table 12: Channel accounting across QCEW endpoints

	2000	2001	2010
<i>Panel A. Baseline</i>			
Peak manufacturing employment share	0.048*** (0.010)	0.043*** (0.009)	0.042*** (0.009)
<i>Panel B. Add manufacturing scale and HSPS</i>			
Peak manufacturing employment share	0.015 (0.012)	0.014 (0.011)	0.018* (0.010)
Log manufacturing establishment scale at peak	0.0041* (0.0023)	0.0033 (0.0021)	0.0024 (0.0025)
Manufacturing HSPS / classified employment at peak	1.927*** (0.535)	1.884*** (0.508)	1.671*** (0.506)
Attenuation of peak coefficient (%)	68.3	67.8	56.9
Observations	1,487	1,511	1,509

Notes: Every outcome uses the QCEW public-record numerator and disclosed service-aggregate denominator. Each column uses the maximum available common sample for that endpoint. Specifications otherwise match the baseline and joint-channel columns of Table 1. Standard errors are Conley with a 100-kilometer cutoff. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 13: Outcome endpoint and peak-window robustness

	Full sample 2000	Interior peaks 2000	Full sample 2010	Interior peaks 2010
All-industries IV	0.133 (0.105)	0.198* (0.111)	0.092 (0.090)	0.188* (0.105)
First-stage Conley F	14.67	16.14	13.77	14.36
Anderson–Rubin p -value	0.203	0.093	0.330	0.100
High-throughput IV	0.283** (0.118)	0.318*** (0.121)	0.241** (0.120)	0.273** (0.117)
First-stage Conley F	12.87	12.57	11.00	10.29
Anderson–Rubin p -value	0.009	0.008	0.051	0.028
Observations	1,487	1,013	1,509	1,031

Notes: The 2000 and 2010 outcomes use the same QCEW public-record construction. Interior-peak samples exclude counties whose observed manufacturing peak occurs at either boundary of the 1880–1940 window. All other variables and specifications match Table 2. Standard errors are Conley with a 100-kilometer cutoff. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.5.1 QCEW endpoint and peak-window robustness

Table 12 moves the public-record QCEW outcome from 2000 to 2001 and 2010. The baseline peak coefficient remains positive in every year. Adding manufacturing scale and manufacturing HSPS attenuates it by 68 percent in 2000, 68 percent in 2001, and 57 percent in 2010. Manufacturing HSPS remains large and precisely estimated at every endpoint.

Table 13 reports the IV results. The high-throughput estimate is 0.241 in 2010, with a conventional p -value of 0.045 and an Anderson–Rubin p -value of 0.051. Excluding counties whose observed peaks occur at the 1880 or 1940 boundaries raises the high-throughput estimate to 0.318 in 2000 and 0.273 in 2010. The result is therefore not generated by counties whose true peak may lie outside the observed window.

Table 14: QCEW denominator and public-record bridge

Year	Construction	Estimate	SE	<i>p</i> -value	AR <i>p</i> -value
2000	Public numerator; disclosed service aggregate	0.283	0.118	0.016	0.009
2000	Public numerator; sum of public components	0.188	0.123	0.126	0.117
2000	Public numerator; aggregate including utilities	0.281	0.117	0.017	0.009
2001	Public numerator; disclosed service aggregate	0.279	0.116	0.016	0.013
2010	Public numerator; disclosed service aggregate	0.241	0.120	0.045	0.051
2010	Public numerator; sum of public components	0.138	0.127	0.275	0.289
2010	Public numerator; aggregate including utilities	0.239	0.119	0.045	0.052

Notes: Each row reports the high-throughput IV with otherwise identical treatment, controls, sample rule, and 100-kilometer Conley inference. The component-sum construction adds only public two-digit values. The aggregate-denominator construction retains all employment incorporated by BLS in the disclosed private service-providing total. The final construction in each endpoint shows that subtracting published private utilities is immaterial.

The disclosed service aggregate is preferable to a denominator formed by summing public two-digit records. The aggregate retains confidential non-HSPS employment that belongs in the denominator, whereas the component sum removes confidential employment from both sides of the ratio. Table 14 shows that including private utilities is immaterial. It also shows that the public-record result is present in 2001, when the underlying file records explicit suppression flags. Section A.5.4 uses those flags rather than leaving their numeric fields at zero.

A.5.2 High- and low-throughput manufacturing peaks

Population growth and urbanization may accompany industrialization even if they do not explain why some counties subsequently develop HSPS services. I examine this distinction by partitioning manufacturing into Chandler high-throughput industries and all remaining industries, a mutually exhaustive group labeled low throughput. For each county, I calculate the maximum employment share of each group separately over 1880, 1900, 1910, 1920, 1930, and 1940. I then estimate a separate event study around each type-specific peak. The high-throughput regression interacts its peak magnitude with event time while conditioning on the low-throughput peak magnitude interacted with the same event-time windows; the low-throughput regression reverses the roles. This design does not equate a larger share of one type at the county-wide manufacturing peak with a larger manufacturing peak overall. It asks whether counties with larger peaks of each manufacturing type exhibit the same subsequent population, urbanization, manufacturing-HSPS, and service-HSPS patterns. More specifically I run the following specification:

$$Y_{ct} = \alpha_c^g + \lambda_{s(c),t}^g + \mu_{k_{ct}^g}^g + \sum_{b \in \mathcal{B}} \mathbf{1}\{k_{ct}^g \in b\} \left[\beta_b^g \frac{P_c^g}{0.10} + \rho_b^g \frac{P_c^g}{0.10} \right] + X_c' \Gamma^g \frac{t - 1900}{10} + \varepsilon_{ct}^g,$$

Where, the terms α_c^g , $\lambda_{s(c),t}^g$, and $\mu_{k_{ct}^g}^g$ are county, state-by-calendar-year, and exact decennial event-time fixed effects, respectively. X_c contains the predetermined population, urbanization, schooling, college, and enslaved-population controls measured in 1850 and the 1870 process-, machinery-, and traditional-industry shares; Γ^g allows each to have a linear trend per calendar decade. The error is ε_{ct}^g . I estimate the equation separately for each manufacturing type and each outcome. The plotted coefficient β_b^g is the association with a 10-percentage-point larger peak of type g in window b , conditional on the

other type's peak magnitude through ρ_b^g . The other type's peak *year* does not set the event clock. Standard errors use the 100-kilometer spatial covariance estimator. Figure 8 uses uniform 20-year event-time windows for both types and every outcome. The event-time specification ends at +60 because this is the last horizon at which every eligible peak cohort can contribute.³³ Service numerators and denominators come from a common employment universe within each vintage: harmonized Census records through 1950, imputed early CBP in 1960 and 1970, and harmonized EFSY CBP thereafter; this is the same panel as Figure 3.

Both types of manufacturing peak predict local development, being associated with increasing urbanization and log-population both before and after their peak. Over event times +10 through +40, a 10-percentage-point larger high-throughput peak is associated with 0.375 log points more population and a 6.33-percentage-point larger urban share. The corresponding low-throughput associations are 0.143 log points and 2.32 percentage points; both are statistically distinguishable from zero. Thus, population growth and urbanization are not exclusive to Chandler-type manufacturing. However, the HSPS outcomes differ sharply across manufacturing types. In the same +10 to +40 window, the manufacturing-HSPS share is 0.831 percentage points higher following a larger high-throughput peak and 0.396 points lower following a larger low-throughput peak. In services HSPS the contrast is even more informative: at +50/+60, the high-throughput coefficient is 0.429 percentage points ($p = 0.006$), whereas the low-throughput coefficient is -0.095 points ($p = 0.362$). Their difference is 0.525 points ($p = 0.003$). These results are consistent with high-throughput industrialization being associated with HSPS work first within manufacturing and later in service industries. They also show why broad population growth or urbanization alone is an incomplete account of the HSPS pattern: low-throughput peaks predict both, but no HSPS response.

A.5.3 Identification from initial industry shares

The shift-share instrument has complementary shares and shocks interpretations. Under the shares interpretation of Goldsmith-Pinkham et al. (2020), the national industry-growth rates transform the predetermined 1880 county industry shares into a single instrument, and the 2SLS coefficient is a weighted combination of the just-identified estimates associated with those shares. This representation makes clear which industries supply identifying variation. Figure 9 reports the corresponding Rotemberg weights after residualizing the treatment, shares, and outcome with respect to the complete Table 2 conditioning set.

Twenty-six high-throughput industry shocks are usable in the estimation sample. Blast furnaces and steel works, grain milling, tobacco, primary nonferrous metals, and other primary iron and steel carry the largest absolute weights. The largest absolute Rotemberg weight is 19.8 percent; the three largest account for 49.8 percent and the five largest for 70.4 percent. Negative weights account for 10.2 percent of total absolute weight. The inverse-Herfindahl effective number of Rotemberg weights is 8.6, while the exposure-weighted effective number of national shocks is 4.7. The design is therefore concentrated, but it is not identified by one industry alone.

Figure 10 verifies this point directly by removing each high-throughput industry in turn and rebuilding both stages. Every second-stage coefficient remains positive. Twenty-seven of the 28 reported leave-one-out concordance estimates remain significant at 5 percent; the only exception is the removal of grain

³³ For the high-throughput service regression, support falls from 1,417 county observations at +60 to 226 at +70 and six at +100.




Figure 8: Local development and HSPS employment around separate manufacturing-type peaks. The blue path is aligned on the county’s high-throughput manufacturing-employment peak and the red path on its separately calculated low-throughput peak. Each point is the coefficient associated with a 10-percentage-point larger peak share in the indicated manufacturing type, conditional on the magnitude of the other type’s peak interacted with every event-time window. The windows are $-30/-20$, $-10/0$, $+10/+20$, $+30/+40$, and $+50/+60$; event time -40 is omitted. Panels A and B report log county population and the urban population share. Panel C reports manufacturing HSPS employment divided by manufacturing employment; Panel D reports HSPS-service employment divided by classified private employment. The high- and low-peak regressions use different eligible county sets because each requires an observation at -40 relative to its own peak: 1,419 and 1,158 counties, respectively, for the service outcome. Both specifications include county, state-by-calendar-year, and exact decennial event-time fixed effects and predetermined baseline-covariate linear trends. Bars are 95-percent confidence intervals based on a 100-kilometer spatial covariance estimator.




Figure 9: Rotemberg weights in the high-throughput design. Absolute Rotemberg-weight shares are calculated on the QCEW-2000 estimation sample after partialling out the state and peak-cohort fixed effects, five 1850 controls, and complete 1850–1870 histories used in Table 2. Colors distinguish positive from negative weights. The underlying machine-readable table reports signed weights, first stages, and just-identified estimates for every usable industry.

milling, which lowers first-stage precision but leaves a positive estimate of 0.287. Removing primary metals, steel, or transportation equipment does not eliminate the result.

The within-group weights in equation (46) sum to one wherever the instrument is defined. The benchmark shifter therefore asks which high-throughput industries a county initially specialized in, conditional on having a classifiable high-throughput base; it does not use the initial size of that base as an omitted exposure share. Table 15 nevertheless reports economy-scaled alternatives in which the incomplete high-throughput share is included separately, as required under the quasi-experimental exposure framework of [Borusyak et al. \(2022\)](#). It also alters the mapping of debatable sectors. The estimate remains positive and statistically significant when processed foods, machinery, or transportation equipment are removed separately, under a narrower continuous-process classification, and under a broader classification. Removing all processed foods weakens the first stage and makes the estimate imprecise; that specification discards several industries that Chandler explicitly associates with high-throughput production.

The comparison that most directly tests the mechanism holds the instrument construction fixed and changes the production regime supplying the national shifts. Table 16 enters the high-throughput and comparison-group shifters jointly. High-throughput national growth predicts both a larger manufacturing peak and a larger long-run HSPS-service share. Growth in all remaining manufacturing industries or in the narrower traditional low-throughput set does not predict the service outcome. Equality of the second-stage coefficients is rejected at 10 percent against all non-high-throughput industries and at 5




Figure 10: Leave-one-industry-out high-throughput estimates. Each row removes the labeled IND1950 industry from the high-throughput set, reconstructs the leave-county-out shifter and its within-group weights, and re-estimates the first and second stages on the Table 2 sample. Points are 2SLS estimates; bars are 95-percent confidence intervals based on 100-kilometer Conley standard errors. The vertical line marks the benchmark estimate.

percent against the traditional group. A generic local-demand or agglomeration account would need to explain why national growth outside high-throughput manufacturing generates industrial exposure without the same long-run service-composition response.

Table 15: Alternative throughput classifications and exposure normalizations

Specification	IV estimate	SE	<i>p</i> -value	First-stage <i>F</i>	AR <i>p</i> -value
<i>Panel A. Throughput classification</i>					
Benchmark Chandler set	0.283	0.118	0.016	12.87	0.009
Exclude processed foods	0.325	0.257	0.207	3.63	0.159
Exclude machinery	0.189	0.098	0.055	28.98	0.041
Exclude processed foods and machinery	0.366	0.175	0.036	9.19	0.011
Exclude transportation equipment	0.251	0.102	0.014	16.38	0.008
Narrow continuous/process set	0.183	0.078	0.018	29.69	0.011
Broader Chandler set	0.267	0.120	0.026	12.45	0.014
<i>Panel B. Exposure normalization</i>					
Within-high shares, benchmark	0.283	0.118	0.016	12.87	0.009
Within-high shares; control high-throughput share of manufacturing	0.322	0.135	0.017	11.36	0.009
Economy-scaled exposure; control included-share sum	0.206	0.080	0.010	40.29	0.015
Within-high shares; control high-throughput share of county employment	0.822	0.674	0.223	1.59	0.016
Economy-scaled exposure; county-employment denominator	−0.419	1.313	0.749	0.38	0.645

Notes: Every row instruments the same peak-manufacturing treatment and uses the QCEW-2000 outcome, sample rule, fixed effects, and predetermined controls in Table 2. The within-high specifications normalize 1880 industry employment by the county's total employment in the relevant high-throughput set. The economy-scaled specification normalizes by the stated broader base and includes the sum of included exposure shares as a separate control. First-stage *F* statistics and conventional inference use a 100-kilometer Conley covariance estimator; AR entries are the corresponding reduced-form tests.

Table 16: High-throughput and comparison-industry contrasts

Joint instrument set	High throughput	Comparison group	<i>p</i> -value, equality
<i>Panel A. First stage: peak manufacturing employment share</i>			
High throughput versus all non-high-throughput industries	1.928	0.193	0.014
High throughput versus traditional low-throughput industries	1.928	0.041	0.011
<i>Panel B. Reduced form: 2000 HSPS share of private services</i>			
High throughput versus all non-high-throughput industries	0.546	−0.049	0.053
High throughput versus traditional low-throughput industries	0.547	−0.074	0.021

Notes: Entries are percentage-point effects of a one-standard-deviation shifter from regressions that include both standardized group shifters simultaneously. Every regression uses the QCEW-2000 sample and the complete Table 2 conditioning set. The final column tests equality of the two coefficients using a 100-kilometer Conley covariance estimator. Panel B reports reduced forms because each group shifter predicts the common manufacturing-peak treatment with a different strength.

A.5.4 Flag-aware QCEW imputation

The public-record outcome is well-known to suppress data on employment from smaller counties to protect confidentiality of businesses. To address whether this may systematically bias the estimates, I perform two additional robustness checks: imputing suppressed data and aggregating to commuting zones. First I use the flags available beginning in 2001 which provide employment bins for suppressed data in QCEW to estimate the outcome variable in data suppressed counties. QCEW is preferable to EFSY for this specific purpose because it reports, in addition to a denominator capturing all private service employment, an aggregate employment count of NAICS industries 54-55 and 56 (parent code 1024). This permits a more precise estimate for the dependent variable than would be possible from EFSY. The disclosed county private service-providing total, code 102, supplies the top adding-up constraint, and disclosed service supersectors 1021–1029 supply the next set of margins. Two of these supersectors are especially informative for the outcome. Code 1022 equals Information (NAICS 51), while code 1024 is the higher-level Professional and Business Services aggregate that equals NAICS 54, 55, and 56. Consequently, whenever NAICS 56 is public, NAICS 54 plus 55 is observed exactly as code 1024 less NAICS 56 even if the separate 54 and 55 records are confidential. When NAICS 56 is also confidential, the disclosed 1024 total still tightly constrains the joint residual that can be assigned to 54, 55, and 56. Public three-digit descendants provide additional lower bounds whenever a two-digit parent is confidential.

After imposing these restrictions, I allocate only the residual within each disclosed parent. The allocation uses the corresponding state-industry employment shares. National-industry and equal-share allocations are reported as sensitivities. The allocation never uses the shift-share instrument, the manufacturing peak, the long-run outcome, or any regression residual. I implement both zero- and one-worker lower bounds for an explicitly confidential macro cell. The lower-bound choice is immaterial because the disclosed parent margins determine nearly all allocated employment. QCEW and CBP draw on different administrative systems and reference periods, and their treatment of auxiliary establishments is not identical. An EFSY value is therefore not an imputation of the missing QCEW record: it is an estimate of the corresponding CBP employment cell. EFSY remains a valuable independent-source robustness exercise, but the QCEW hierarchy is the appropriate source for reconstructing the latent QCEW numerator used in the main analysis.

Across 3,194 located counties, the preferred reconstruction implies an employment-weighted HSPS share of 14.38 percent, compared with 14.53 percent in the official national QCEW aggregates. The state- and national-share allocations differ by less than 0.01 percentage point in the aggregate, and changing the lower bound from zero to one worker changes the high-throughput estimate by less than 0.001. Unallocated hierarchy margins and published-child overages total 1,504 and 2,680 workers, respectively, and are negligible relative to national service employment.

The regression results are similarly stable across economically informed allocations. The high-throughput coefficient is 0.187 (SE 0.100, $p = 0.061$, AR $p = 0.066$) using state-industry shares and 0.185 (SE 0.099, $p = 0.063$, AR $p = 0.066$) using national-industry shares. An equal split, which deliberately ignores the observed sector composition at both the state and national levels, gives a smaller and less precise estimate of 0.159 (SE 0.100, $p = 0.113$, AR $p = 0.121$). I use state shares as the preferred allocation because they allow the division of a disclosed local parent to reflect regional industry composition while remaining external to the county's treatment and outcome.

Table 17: Disclosure flags and QCEW-imputation audit, 2001

Sector	County records	Confidential records	Confidential with public local parent	Confidential with public three-digit descendant
NAICS 51	3,129	583	0	36
NAICS 54	3,197	948	826	0
NAICS 55	2,160	1,153	1,124	0
<i>Aggregate reconciliation</i>				
Preferred state-share allocation	14.3825%		employment-weighted HSPS share	
National-share allocation	14.3788%		employment-weighted HSPS share	
Equal-share allocation	14.4228%		employment-weighted HSPS share	
Official national QCEW aggregates	14.5272%		employment-weighted HSPS share	

Notes: A public local parent is QCEW supersector 1022 for Information or 1024 for Professional and Business Services. The latter equals NAICS 54, 55, and 56. Published cells remain fixed. State, national, and equal shares are alternative priors used only to split an otherwise unidentified residual among confidential children of a disclosed parent. Official national employment is 12,179,566 in NAICS 51, 54, and 55 and 83,839,631 in private services net of utilities.

Table 18 gives the county results. The first row reproduces Table 2. EFSY's fully imputed native-2000 CBP outcome is positive but does not reproduce conventional county-level significance: the high-throughput estimate is 0.103 with a standard error of 0.080 and an Anderson–Rubin p -value of 0.202. The QCEW-specific flag-aware estimate on the other hand is 0.187 with a standard error of 0.100. Its conventional p -value is 0.061 and its Anderson–Rubin p -value is 0.066. This discrepancy between the point estimates obtained using EFSY and QCEW as sources for the outcome variable is to be expected because QCEW reports a separate parent (code 1024) which contains the sum of NAICS 54-55-56, thus providing an additional adding up constraint and a more precise estimate of the data suppressed portion of the outcome variable. The all-industry first stage remains strong, with a Conley F statistic of 13.55, but its second-stage estimate is 0.062 and its Anderson–Rubin p -value is 0.519. Averaging the flag-aware outcomes over 2001–2003 produces the same contrast: the all-industry estimate is 0.065 and insignificant, while the high-throughput estimate is 0.164 and significant at 10 percent under both the conventional and Anderson–Rubin tests.

The last two rows hold the year, denominator, controls, instruments, and 1,501-county sample fixed and change only the numerator source. The QCEW reconstruction gives 0.180, while the EFSY numerator gives 0.138 and is not statistically significant. The EFSY-minus-QCEW difference is -0.043 with a standard error of 0.098 and a p -value of 0.662. The data do not reject equality of the two estimates.

The flag-aware outcome also preserves the channel result. At the county level, the baseline peak coefficient is 0.029. Adding manufacturing scale and manufacturing HSPS reduces it to 0.012, an attenuation of 59 percent; manufacturing HSPS remains positive and precisely estimated at 1.260. At the commuting-zone level, the point estimate on the peak is already less precise, but the joint specification again leaves a small residual peak coefficient while both scale and manufacturing HSPS remain positive. EFSY gives a 75-percent attenuation at the commuting-zone level in both 2000 and 2001.

Table 18: County IV results under alternative disclosure treatments

Outcome construction	OLS	All-industry IV	F_{all}	AR p_{all}	High-throughput IV	F_{high}	AR p_{high}
QCEW public records, 2000	0.029** (0.012)	0.133 (0.105)	14.67	0.203	0.283** (0.118)	12.87	0.009
EFSY-imputed CBP, native 2000	0.028*** (0.010)	0.031 (0.094)	14.67	0.747	0.103 (0.080)	12.87	0.202
Flag-aware QCEW, state shares, 2001	0.018* (0.009)	0.062 (0.093)	13.55	0.519	0.187* (0.100)	11.25	0.066
Flag-aware QCEW, 2001–2003 average	0.018** (0.009)	0.065 (0.085)	14.30	0.458	0.164* (0.091)	13.00	0.075
Flag-aware QCEW, matched 2001 sample	0.017* (0.009)	0.065 (0.095)	13.17	0.511	0.180* (0.100)	11.07	0.076
EFSY numerator / QCEW denominator, matched 2001 sample	0.034*** (0.011)	−0.013 (0.113)	13.17	0.910	0.138 (0.102)	11.07	0.186
Observations	1,487; 1,487; 1,511; 1,497; 1,501; 1,501, respectively						

Notes: Every specification matches Table 2. The matched 2001 rows use the same QCEW service denominator and differ only in the NAICS 51, 54, and 55 numerator. F statistics use 100-kilometer Conley first-stage standard errors; AR entries are the corresponding Conley reduced-form tests. The EFSY-minus-QCEW high-throughput coefficient in the matched sample is −0.043 (SE 0.098, $p = 0.662$). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 19: Channel accounting under imputation and commuting-zone aggregation

	QCEW public county, 2000	QCEW flag-aware county, 2001	QCEW flag-aware CZ, 2001	EFSY native CZ, 2000	EFSY harmonized CZ, 2001
<i>Panel A. Baseline</i>					
Peak manufacturing employment share	0.048*** (0.010)	0.029*** (0.008)	0.022 (0.017)	0.038** (0.016)	0.037** (0.017)
<i>Panel B. Add manufacturing scale and HSPS</i>					
Peak manufacturing employment share	0.015 (0.012)	0.012 (0.011)	−0.012 (0.020)	0.010 (0.022)	0.009 (0.023)
Log manufacturing establishment scale at peak	0.0041* (0.0023)	0.0016 (0.0021)	0.0050* (0.0030)	0.0028 (0.0033)	0.0028 (0.0034)
Manufacturing HSPS / classified employment at peak	1.927*** (0.535)	1.260*** (0.433)	1.382* (0.814)	1.729** (0.834)	1.693** (0.837)
Attenuation of peak coefficient (%)	68.3	58.7	153.3	74.8	74.6
Observations	1,487	1,511	391	391	391

Notes: The first two columns use fixed 1850 counties. The last three aggregate additive historical and modern employment counts to 1990 commuting zones before forming shares. CZ fixed effects use the dominant historical state and the CZ peak cohort. The QCEW flag-aware columns use the preferred state-share allocation. Specifications otherwise reproduce the baseline and joint-channel columns of Table 1. The attenuation percentage exceeds 100 when the residual peak coefficient changes sign. Standard errors are Conley with a 100-kilometer cutoff. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.5.5 EFSY imputation and commuting-zone aggregation

Commuting zones provide an additional check for the sensitivity of the estimate to data suppression practices because, given commuting zones are larger, data suppression is much less of an issue at the commuting zone level. To perform this robustness check, I assign county data in 2000 and 1850 to commuting zones using a weighted area overlap method. I then recompute all outcomes, controls and the shift-share instruments at the commuting-zone level using leave-commuting-zone-out national growth.

Table 20 makes two comparisons. First, the 2001 rows put the flag-aware QCEW and EFSY outcomes on the same 391 commuting zones. The QCEW estimate is 0.324 (SE 0.141, $p = 0.022$, AR $p = 0.005$). The EFSY estimate is 0.204 (SE 0.125); the conventional p -value is 0.102 and the weak-instrument-robust

Table 20: QCEW and EFSY IV estimates at the commuting-zone level

Outcome construction	OLS	All-industry IV	F_{all}	AR p_{all}	High-throughput IV	F_{high}	AR p_{high}
Flag-aware QCEW, 2001	0.004 (0.020)	−0.274 (1.176)	0.17	0.736	0.324** (0.141)	9.03	0.005
EFSY harmonized, 2001	0.021 (0.020)	−0.537 (1.632)	0.17	0.541	0.204 (0.125)	9.03	0.094
EFSY numerator / QCEW denominator, 2001	0.044* (0.026)	−1.015 (2.804)	0.17	0.314	0.248* (0.141)	9.03	0.070
EFSY native classification, 2000	0.028 (0.018)	−0.640 (1.877)	0.17	0.428	0.279** (0.127)	9.03	0.012
EFSY harmonized, 2000	0.023 (0.020)	−0.362 (1.309)	0.17	0.669	0.248* (0.130)	9.03	0.035
EFSY harmonized, 2010	0.018 (0.018)	0.020 (0.839)	0.17	0.981	0.279** (0.135)	9.03	0.009

Commuting zones

391

Notes: All rows use the same standard 1990 commuting zones, historical treatment, controls, fixed effects, and commuting-zone shift-share instruments. The QCEW row aggregates the preferred county-level state-share reconstruction. The EFSY 2001 row uses an EFSY numerator and EFSY private-service denominator; the following row changes only the numerator while retaining the QCEW denominator. The all-industry instrument is weak after aggregation and its CZ second-stage estimates are not interpreted. Coefficient and first-stage standard errors are Conley with a 100-kilometer cutoff. AR entries report the corresponding reduced-form test. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

reduced form rejects at 10 percent (AR $p = 0.094$). Holding the QCEW denominator fixed and changing only the numerator raises the EFSY estimate to 0.248. The EFSY-minus-QCEW numerator difference is −0.076 (SE 0.080, $p = 0.343$), so the two imputed estimates are not statistically distinguishable. The discrepancy between coefficients obtained from regressions using EFSY and QCEW to compute the dependent variable is thus smaller at the commuting zone level, where data suppression is much less of an issue.

Second, the table reports the EFSY outcomes used in the pre-existing endpoint exercises. The native-2000 estimate is 0.279 (SE 0.127, $p = 0.028$, AR $p = 0.012$); the harmonized-2000 estimate is 0.248 (SE 0.130, $p = 0.056$, AR $p = 0.035$); and the harmonized-2010 estimate is 0.279 (SE 0.135, $p = 0.039$, AR $p = 0.009$). EFSY therefore does not reproduce a positive and conventionally significant estimate at the county level, but it does recover the high-throughput result after employment is aggregated to commuting zones. This geography dependence is consistent with confidentiality being concentrated in small county-industry cells and with HSPS activity serving labor markets larger than counties.

The difference between QCEW and EFSY is concentrated in the two sectors for which institutional coding and local margins differ most. In the 2000 county bridge, QCEW and EFSY give similar high-throughput coefficients for NAICS 51, 0.072 and 0.077. The QCEW coefficients for NAICS 54 and 55 are 0.141 and 0.070, whereas the corresponding EFSY coefficients are 0.034 and 0.003. EFSY's harmonization is built to recover CBP employment consistently across industry vintages; the QCEW allocation is built to recover the QCEW concept used by the main outcome. Because the two systems differ in their reference period, establishment frame, treatment of auxiliary establishments, and available hierarchy margins, neither is a mechanical correction of the other.

The QCEW-specific reconstruction is preferable for auditing the main QCEW result for three reasons. It retains the same employment concept and annual-average reference period as the main outcome; it uses explicit QCEW confidentiality flags rather than inferring which absent records are confidential; and it exploits the disclosed code-1024 parent, which contains NAICS 54 and 55 plus NAICS 56 and therefore supplies a direct local adding-up restriction on the target numerator.

Table 21: Industry-shock inference for the shift-share estimates

Outcome and inference	IV estimate	SE	<i>p</i> -value	First-stage <i>F</i>	AR <i>p</i> -value	Industries
<i>Panel A. High-throughput instrument</i>						
QCEW public, 2000: county-level Conley	0.283	0.118	0.016	12.87	0.009	26
QCEW public, 2000: exact leave-out shocks	0.283	0.159	0.088	5.26	0.062	26
QCEW public, 2000: common national shocks	0.284	0.132	0.041	3.94	0.047	26
Flag-aware QCEW, 2001: county-level Conley	0.187	0.100	0.061	11.25	0.066	26
Flag-aware QCEW, 2001: exact leave-out shocks	0.187	0.121	0.135	5.02	0.065	26
Flag-aware QCEW, 2001: common national shocks	0.188	0.106	0.089	3.83	0.034	26
Flag-aware QCEW, 2001–2003: county-level Conley	0.164	0.091	0.069	13.00	0.075	26
Flag-aware QCEW, 2001–2003: exact leave-out shocks	0.164	0.108	0.142	5.39	0.082	26
<i>Panel B. All-industry instrument, QCEW public outcome</i>						
County-level Conley	0.133	0.105	0.202	14.67	0.203	57
Exact leave-out shocks	0.133	0.223	0.551	1.59	0.623	57
Common national shocks	0.133	0.141	0.349	3.29	0.422	57
<i>Panel C. Concentration diagnostics for high-throughput shocks</i>						
Exposure-weighted effective number of shocks				4.71		
Inverse-Herfindahl effective number of Rotemberg weights				8.62		
Largest / three largest / five largest absolute weights				19.8% / 49.8% / 70.4%		
Negative share of total absolute weight				10.2%		

Notes: County-level rows use the 100-kilometer Conley estimator reported elsewhere. Exact leave-out rows retain the county-specific leave-out shock and cluster the equivalent shock-level score by IND1950 industry. Common-shock rows use one leave-sample-out national growth shock per industry. Their first-stage *F* is the shock-level statistic. *p*-values use finite-industry degrees of freedom; AR entries are the corresponding reduced-form tests. The bounded 95-percent exact-shock AR confidence set for the main high-throughput estimate is $[-0.023, 2.108]$; for the flag-aware outcome it is $[-0.014, 2.181]$.

A.5.6 Identification and inference from national industry shocks

Under the shocks interpretation of [Borusyak et al. \(2022\)](#), the identifying assumptions apply to national industry-growth shocks rather than to county exposure shares. The relevant sampling variation is therefore at the industry level. I implement two versions. The exact version retains each leave-county-out county–industry shock and clusters the equivalent shock-level score by industry. The common-shock version replaces the very small county-specific leave-out differences with one national shock per industry and estimates the equivalent collapsed regression. Following [Adão et al. \(2019\)](#), both versions account for the fact that counties with similar industry shares are exposed to common disturbances.

Table 21 shows the results. In the exact leave-out transformation, the shock-level coefficient is algebraically the county-level 2SLS coefficient. What changes is the estimated uncertainty. For the public-record QCEW outcome, the exact high-throughput estimate is 0.283 with an industry-clustered standard error of 0.159. The conventional *t* test rejects zero at 10 percent ($p = 0.088$), and the weak-instrument-robust Anderson–Rubin test also rejects at 10 percent ($p = 0.062$). The common-national-shock approximation gives 0.284, a standard error of 0.132, and an AR *p*-value of 0.047. The flag-aware QCEW and 2001–2003 averaged outcomes give smaller but still positive estimates; their exact-shock AR *p*-values are 0.065 and 0.082.

The same exercise also makes the design’s limits transparent. There are 26 usable high-throughput shocks, and the exposure-weighted effective count is approximately 4.7. The exact shock-clustered first-stage *F* is 5.26 for the main outcome and 5.02 for the flag-aware outcome, below the county-level Conley statistics. For the all-industry shifter, shock-level first-stage statistics are also low and neither the reduced form nor the second stage is distinguishable from zero. Thus the national-shocks interpretation supports a positive high-throughput effect at the 10-percent level, but it cannot deliver the same precision as an analysis that treats 1,487 counties as the independent sampling units.

Table 22: High-throughput exposure and pre-1880 changes

Pre-period outcome	Coefficient	SE	<i>p</i> -value
Manufacturing share, 1850–1860 change	−0.0027	0.0018	0.126
Manufacturing share, 1860–1870 change	0.0047	0.0017	0.007
Manufacturing HSPS, 1850–1860 change	0.000012	0.000007	0.100
Manufacturing HSPS, 1860–1870 change	0.000006	0.000007	0.399
HSPS services, 1850–1860 change	0.151	0.171	0.378
HSPS services, 1860–1870 change	−0.225	0.122	0.064

Notes: Each row regresses the indicated change on the standardized high-throughput shifter, state fixed effects, the five 1850 controls, and the complete histories of the other two outcome families, using the 1,487-county QCEW-2000 sample. The tested outcome’s own history is omitted from the conditioning set. Standard errors are Conley with a 100-kilometer cutoff. Peak-cohort fixed effects are excluded because the peak is realized after the 1880 exposure shares.

A.5.7 Predetermined balance, controls, and spatial inference

The high-throughput shifter is not unconditionally unrelated to earlier industrial development. Table 22 regresses each pre-1880 outcome on the standardized high-throughput shifter, the five 1850 covariates, state fixed effects, and the other two families of pre-period histories. The shifter predicts the 1860–1870 change in manufacturing intensity and, at the 10-percent level, the corresponding change in HSPS services; it does not predict the 1850–1860 change in manufacturing, manufacturing HSPS, or services, or the 1860–1870 change in manufacturing HSPS. The level correlations are also nonzero. These facts are why the preferred specification conditions on the complete observed 1850–1870 paths of manufacturing, manufacturing HSPS, and HSPS services rather than claiming unconditional balance of the 1880 industry mix.

Table 23 then varies the conditioning set and the spatial cutoff. Omitting the pre-1880 histories reduces the high-throughput coefficient from 0.283 to 0.184 rather than creating it. Adding proximity to coal, waterpower, ports, the Ohio and Mississippi rivers, and the Great Lakes raises the estimate to 0.320. Adding peak-year wages, population, and urbanization lowers first-stage precision and produces a positive but imprecise estimate of 0.242. Adding the manufacturing employment share in 2000 leaves a positive estimate of 0.196 that is significant at 5 percent. This last result rules out the mechanical explanation that high-throughput exposure changes service composition simply because it leaves more manufacturing employment in place at the endpoint. These variables are post-1880 and can be caused by industrialization, so both rows are channel-blocking exercises rather than preferred identification specifications. The point estimate is invariant to the Conley radius by construction, while the 95-percent uncertainty widens at long cutoffs. It remains significant at 5 percent through 300 kilometers and at 10 percent at 500 kilometers; the AR test likewise rejects at 10 percent throughout.

Taken together, the exercises preserve the empirical structure emphasized in the main text. Manufacturing scale and manufacturing HSPS account for most of the long-run relationship between peak industrialization and service composition. In the event-time sequence, higher peaks in both high-throughput and low-throughput manufacturing are associated with increasing urbanization and population, yet only high-throughput manufacturing predicts the later composition of services, providing suggestive evidence that the mechanism operates in manufacturing industries characterized specifically by high-throughput production rather than general local development. The event-sequence also shows that higher manufacturing peaks do not predict larger service sectors overall: they only predict larger HSPS service sectors

Table 23: Conditioning-set and spatial-inference robustness

Specification	IV estimate	SE	<i>p</i> -value	First-stage <i>F</i>	AR <i>p</i> -value
Preferred controls, 100 km	0.283	0.118	0.016	12.87	0.009
No peak-cohort fixed effects	0.298	0.115	0.010	13.46	0.004
Omit 1850–1870 histories	0.184	0.047	< 0.001	45.26	< 0.001
Add historical geography controls	0.320	0.120	0.008	13.43	0.002
Add peak-year wages, urbanization, and population	0.242	0.197	0.218	5.45	0.188
Add 2000 manufacturing employment share	0.196	0.096	0.042	14.81	0.042
Preferred controls, 200 km	0.283	0.114	0.013	8.37	0.018
Preferred controls, 300 km	0.283	0.134	0.034	5.20	0.038
Preferred controls, 500 km	0.283	0.170	0.096	2.83	0.067

Notes: Every row uses the QCEW-2000 outcome and high-throughput instrument. Historical-geography controls are log distance to coal, waterpower, the nearest port, the Ohio River, and the Mississippi River and an indicator for Great Lakes access. The peak-year row includes average manufacturing wages, the urban population share, log population, and the missing-wage indicator. The endpoint-manufacturing row adds manufacturing employment divided by total classified employment in 2000. Rows 1–6 use a 100-kilometer cutoff; rows 7–9 vary the spatial radius. AR entries are the reduced-form tests computed with the same covariance estimator.

in total county employment. The main IV results are robust to the terminal year, boundary-peak exclusions, alternative Chandler mappings, leave-one-industry-out reconstruction, exposure normalization, historical geography, and wider spatial cutoffs. The estimates are also not driven by data suppression practices which systematically change the outcome. A QCEW-specific flag-aware reconstruction, which imputes missing employment at the county level, recovers the same high-versus-all contrast at the 10-percent level. EFSY yields a positive but imprecise county estimate and positive, statistically significant commuting-zone estimates in 2000 and 2010, which are statistically indistinguishable from the corresponding QCEW estimates. Finally, industry-shock inference is necessarily less precise with 26 high-throughput shocks, but its exact Anderson–Rubin test rejects zero at the 10-percent level. The combined evidence therefore supports a specific conclusion: national growth in production regimes that required large-scale coordination generated both deeper local industrialization and a persistent reallocation of services toward the functions used to organize complex production; broad manufacturing growth alone did not.

B Additional Cross-Country Evidence

This appendix examines whether manufacturing history remains associated with the composition and growth of service economies after manufacturing has receded. I document the timing and size of recorded manufacturing peaks, relate those peaks to HSPS employment and subsequent income growth, and examine both relationships by peak decade.

B.1 When do countries reach their manufacturing-employment peaks?

A country’s recorded peak is its maximum manufacturing-employment share over the available annual history; exact ties are assigned to the earliest year. Table 24 reports country counts and unweighted mean peak shares for the 174-country growth-by-decade sample. Its manufacturing histories combine

Table 24: Timing and Size of Recorded Manufacturing-Employment Peaks

Peak decade	Number of countries	Mean peak share (%)
1950s	3	25.97
1960s	4	29.35
1970s	7	20.67
1980s	4	23.01
1990s	108	15.76
2000s	14	12.29
2010s	34	11.25
All countries	174	15.45

Notes: The sample comprises the 174 countries in Table 27. The manufacturing-employment share uses total employment as its denominator. Peak shares are maxima over each country’s available history; exact ties are assigned to the earliest year. Means give equal weight to each country. Recorded peak dates need not identify historical turning points: 79 of the 108 countries assigned to the 1990s peak at their first available observation, and 88 have a recorded peak in 1991. The 173-country baseline growth sample excludes South Sudan because its initial World Bank income classification is unavailable.

GGDC 10-sector and Gapminder employment data. The service-composition analysis uses a separate 51-country GGDC–ETD sample.

Recorded peaks generally decline in size across later cohorts: the mean is 15.76 percent in the 1990s, 12.29 percent in the 2000s, and 11.25 percent in the 2010s. Of the 174 countries, 140 (80.5 percent) peak before 2010. The concentration in the 1990s partly reflects data coverage: 79 of that decade’s 108 countries peak at their first observation. These dates identify observed-series maxima, which may differ from historical turning points.

B.2 Manufacturing peaks and HSPS employment

I first relate manufacturing history to the expansion of business services, finance, and real estate, which together constitute HSPS. The analysis retains all 51 ETD countries used in Figure 1. Manufacturing peaks combine the full GGDC 10-sector and ETD histories, with ETD preferred in overlapping country-years. GGDC extends the histories of 32 countries; the remaining 19 use their full available ETD histories.

Let H_{ct} denote HSPS employment as a percentage of total employment, $\Delta H_c = H_{c,2018} - H_{c,1990}$, and P_c the historical manufacturing peak in percentage points. The most extensive change and terminal-level specifications are

$$\Delta H_c = \alpha + \beta(P_c/10) + \gamma \Delta \ln y_c + \theta S_{c,1990} + \delta_{r(c)} + \lambda_{g(c,1990)} + \varepsilon_c, \quad (49)$$

$$H_{c,2018} = a + b(P_c/10) + \vartheta S_{c,2018} + d_{r(c)} + \ell_{g(c,1990)} + u_c. \quad (50)$$

Here $\Delta \ln y_c = \ln y_{c,2018} - \ln y_{c,1990}$, and S_{ct} is non-government service employment as a percentage of total employment. GDP per capita comes from the Maddison Project Database 2023. Region effects distinguish Africa, the Americas, and Asia. Initial income-group effects use the four World Bank categories based on 1990 GNI per capita (FY1992). The two sets of fixed effects enter additively.

Table 25 introduces income growth in column (2), both sets of fixed effects in column (3), and the 1990 service share in column (4). For terminal HSPS employment, column (6) adds both sets of fixed effects

Table 25: Manufacturing peaks and HSPS employment: changes and terminal levels

	HSPS share change, 1990–2018				HSPS share, 2018		
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Manufacturing peak (10 pp)	2.867*** (0.507)	2.937*** (0.501)	2.173*** (0.782)	2.283** (0.949)	5.962*** (0.729)	3.255*** (0.648)	2.116** (1.003)
GDP-per-capita gain ($\Delta \ln y$)		−1.202 (0.732)	−0.176 (1.029)	−0.145 (1.079)			
Service share, 1990 (pp)				−0.021 (0.068)			
Service share, 2018 (pp)							0.172* (0.093)
Region and income-group FE	No	No	Yes	Yes	No	Yes	Yes
Observations	51	51	51	51	51	51	51
R^2	0.530	0.550	0.617	0.618	0.647	0.817	0.842

Notes: All columns use the same 51 ETD economies. HSPS is business services, finance, and real estate employment as a percentage of total employment. Columns (1)–(4) use the 1990–2018 change; columns (5)–(7) use the 2018 level. The manufacturing peak is the maximum employment share over the full available GGDC 10-sector-ETD history, with ETD preferred in overlapping years. Its coefficient is reported per 10-percentage-point larger peak. GDP-per-capita gain is $\ln(y_{2018}) - \ln(y_{1990})$, using the Maddison Project Database 2023. The service-share controls use 1990 in column (4) and 2018 in column (7). They measure employment in trade/accommodation/food, transport, business services, finance, real estate, and other services as a percentage of total employment; government services are excluded from their numerator. Their coefficients are per one-percentage-point higher service share in the indicated year. Region and income-group effects are included together as additive fixed effects. Region effects use Africa, Americas, and Asia; initial income-group effects use the World Bank low-, lower-middle-, upper-middle-, and high-income classifications based on 1990 GNI (FY1992). All regressions include an intercept. HC3 standard errors are in parentheses; * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

and column (7) adds the 2018 service share. All regressions use equal country weights and HC3 standard errors.

A ten-percentage-point larger manufacturing peak is associated with a 2.87-percentage-point greater increase in HSPS employment over 1990–2018. The estimate is 2.28 after including income growth, region and initial income-group effects, and the initial service share. For the 2018 HSPS share, the coefficient falls from 5.96 without controls to 2.12 with fixed effects and the contemporary service share. Both final estimates remain significant at the five-percent level.

Disaggregating service-employment growth.

For each service industry j , I estimate $\Delta s_{cj} = \alpha_j + \beta_j(M_c^{\max}/0.10) + \varepsilon_{cj}$, where Δs_{cj} is the 1990–2018 change in its share of total employment, in percentage points, and $M_c^{\max}$ is the historical manufacturing peak expressed as a fraction. Each regression uses all 51 countries and no additional controls. Government services are excluded from the industries examined; total employment remains the denominator.

Business services have the largest coefficient: 1.59 percentage points per ten-percentage-point larger peak, followed by finance (0.67) and real estate (0.61). Business services account for approximately 56




Figure 11: Manufacturing peaks and service-industry employment growth. Points report coefficients from separate regressions of the 1990–2018 change in each industry’s share of total employment on historical peak manufacturing employment. Coefficients are expressed in percentage points per ten-percentage-point larger manufacturing peak. Industries are ordered by their estimated coefficients. Manufacturing peaks use the full available GGDC 10-sector–ETD histories, with ETD preferred in overlapping years. All regressions include an intercept, use the same 51 economies, and assign equal weight to each country. Bars show pointwise 95 percent confidence intervals based on HC3 standard errors and the residual-degrees-of-freedom t distribution. Government services are excluded from the displayed industries.

percent of the aggregate HSPS association. The ETD category includes information and communication, professional, scientific, and technical activities, and administrative and support services; business-organization services are not separately identified. The remaining coefficients are negative: -0.13 for other services, -0.59 for transport, and -2.00 for trade, accommodation, and food services.

B.3 GDP-per-capita growth and peak manufacturing employment

I next relate manufacturing history to average annual GDP-per-capita growth over 2010–2019. To distinguish the historical peak from manufacturing’s remaining size, the most extensive specification includes both contemporary sectoral shares:

$$g_c = \alpha + \alpha_{r(c)} + \alpha_{i(c)} + \beta^m (M_c^{\max}/0.10) + \beta^s (\bar{S}_{c,2010:2019}/0.10) + \delta (\bar{M}_{c,2010:2019}/0.10) + \gamma I_c + \varepsilon_c. \quad (51)$$

The outcome g_c is mean annual log GDP-per-capita growth multiplied by 100. Employment shares are fractions of total employment, averaged over 2010–2019 where indicated. The fixed effects distinguish

Table 26: Manufacturing Peaks and Growth in Service Economies

	All countries					Services > 50%	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Peak manufacturing share (10 pp)	0.305* (1.67)	0.953*** (4.54)	0.733*** (2.70)	0.629** (2.53)	0.398 (0.92)	0.645** (1.99)	1.113** (2.36)
Service-employment share (10 pp)		-0.538*** (-4.70)	-0.311 (-1.41)	-0.572*** (-2.96)	-0.540*** (-2.68)	-1.045** (-2.59)	-1.313** (-2.41)
Manufacturing-employment share (10 pp)					0.354 (0.70)		-0.924 (-1.16)
Observations	173	173	173	173	173	104	104
R^2	0.012	0.137	0.236	0.360	0.362	0.278	0.288
Region and income-group FE	No	No	Yes	Yes	Yes	Yes	Yes
Institutional quality control	No	No	No	Yes	Yes	Yes	Yes

Notes: The dependent variable is average annual GDP-per-capita growth over 2010–2019, in percentage points. Column (1) includes only the manufacturing peak and an intercept. Columns (1)–(5) use the same 173-country sample; columns (6)–(7) use the 104 countries whose average service-employment share over 2010–2019 exceeds 50 percent. The service-employment share and the manufacturing-employment share are averages over 2010–2019; all 173 countries have manufacturing observations in each of these ten years. The peak manufacturing share is the maximum over the available historical series. All employment shares use total employment as the denominator, and their coefficients are reported per ten-percentage-point increase. Columns (5) and (7) add the average manufacturing-employment share to the specifications in columns (4) and (6), respectively. Region and income-group fixed effects enter together as separate additive sets in columns (3)–(7). Income groups are the World Bank classifications based on 2009 GNI per capita (FY2011); the full sample contains 40 low-, 50 lower-middle-, 37 upper-middle-, and 46 high-income countries. South Sudan is omitted because its 2009 classification is unavailable. Columns (4)–(7) control for the World Bank control-of-corruption percentile rank. All regressions include an intercept. HC1 robust t statistics are in parentheses; tests use the residual-degrees-of-freedom t distribution. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

regions and the four World Bank income groups based on 2009 GNI per capita (FY2011). Institutional quality I_c is the World Bank control-of-corruption percentile rank.

Columns (1)–(5) of Table 26 successively add the service share, both sets of fixed effects, institutional quality, and the contemporary manufacturing share to the peak-only regression. Columns (6)–(7) repeat columns (4)–(5) for economies with an average service-employment share above 50 percent.

In the full sample, controlling for contemporary manufacturing reduces the peak coefficient from 0.63 to 0.40, with $p = 0.356$ for the latter. Among economies where services employ a majority of workers, the corresponding coefficient rises from 0.64 to 1.11 and remains significant ($p = 0.020$). In this group, a ten-percentage-point larger historical peak is therefore associated with 1.11 percentage points faster annual growth after accounting for both contemporary sectoral shares. The difference from the full-sample coefficient has not been established statistically; alternative service-share thresholds were also examined.

B.4 Service composition and growth after earlier manufacturing peaks

I next examine whether the growth and service-composition associations are present among countries whose manufacturing peaks occurred well before the outcome period. Let D_{cd} indicate the decade of country c 's recorded peak. For each outcome Y_c , I estimate

$$Y_c = a + \sum_{d \in \mathcal{D}} \beta_d D_{cd} (M_c^{\max} / 0.10) + X_c' \Gamma + e_c. \quad (52)$$

Each regression has a common intercept and decade-specific manufacturing-peak slopes. Table 27 presents the same two specifications for each outcome: one without additional controls and one adding log GDP per capita in 2018 and region fixed effects. The regressions include neither separate decade intercepts nor a linear peak-year control. All coefficients are scaled per ten-percentage-point higher manufacturing peak, and HC3 standard errors are used throughout.

Columns (1)–(2) use average annual GDP-per-capita growth over 2010–2019 for 174 countries. Columns (3)–(4) use the 2018 HSPS share of non-government service employment for the 51 ETD countries. For the latter outcome, the numerator is business, finance, and real-estate employment; the denominator additionally includes trade, accommodation and food services, transport, and other services. Government is excluded from both. Unlike the earlier HSPS regressions, this outcome measures composition within services rather than a share of total employment.

For countries with recorded peaks in the 1990s, a ten-percentage-point larger manufacturing peak is associated with 0.56 percentage points faster annual GDP-per-capita growth without controls ($p = 0.008$) and 0.63 points after accounting for 2018 income and region ($p = 0.044$). For these 108 countries, average manufacturing employment fell from 10.92 to 10.45 percent of total employment between 2010 and 2019, while services rose from 52.75 to 56.64 percent. The positive growth association is therefore present among countries whose recorded peaks preceded the growth window and whose employment continued, on average, to shift toward services. The 2010s coefficients are larger, at 1.28 and 1.05, with $p = 0.008$ and $p = 0.051$, respectively.

The service-composition results show a related pattern. Among the 11 ETD countries with 1990s peaks, the estimated association with the 2018 HSPS share is 9.10 percentage points without controls ($p < 0.001$) and 5.11 points with income and region controls ($p = 0.099$). For the 20 countries with 2010s peaks, neither coefficient is statistically significant. Income and region adjustment thus weakens the earlier-peak composition association, while the corresponding growth association remains significant at the five-percent level.

The decade results complement the contemporary manufacturing-share comparison in Table 26. That table accounts for manufacturing’s remaining employment share; this table establishes that positive associations are present among countries with substantially earlier recorded peaks. Growth in 2010–2019 occurs 11–29 years after a 1990s peak, and 2018 service composition is measured 19–28 years later. Peak dates remain subject to the coverage limits in Table 24; the smallest decade groups also contain few countries.

B.5 Timing and size of county manufacturing peaks

Figure ?? summarizes the timing and size of county manufacturing peaks. For each county, I define the manufacturing peak as in equation (1): the maximum of manufacturing employment relative to total classified county employment over 1880, 1900, 1920, 1930, and 1940. The years 1890 and 1910 are excluded because the required population and manufacturing tabulations are unavailable, respectively. The figure reports the average peak manufacturing-employment share and the number of counties attaining their peak in each observed census year.

The distribution shows substantial variation in the timing of industrialization. Some counties peak around the turn of the twentieth century, while others peak only in the interwar period. This variation provides the staggered timing used in the event-study analysis.

Table 27: Historical Manufacturing Peaks, Subsequent Growth, and Service Composition

Peak decade	GDP-per-capita growth, 2010–2019		HSPS share of services, 2018	
	(1)	(2)	(3)	(4)
1950s	0.085 (0.299)	0.340 (0.406)	6.078*** (1.176)	3.846 (2.297)
1960s	0.062 (0.139)	0.174 (0.289)	—	—
1970s	0.169 (0.293)	0.226 (0.380)	8.751*** (0.956)	5.147* (3.027)
1980s	−0.146 (1.027)	−0.120 (1.056)	7.822*** (1.353)	3.717 (2.383)
1990s	0.559*** (0.208)	0.625** (0.308)	9.098*** (2.211)	5.107* (3.022)
2000s	0.179 (0.471)	0.167 (0.568)	10.425*** (3.307)	5.787 (4.404)
2010s	1.278*** (0.475)	1.054* (0.535)	2.603 (1.634)	−0.045 (2.247)
Log GDP per capita, 2018	No	Yes	No	Yes
Region fixed effects	No	Yes	No	Yes
Observations	174	174	51	51
R^2	0.076	0.115	0.701	0.754

Notes: Each column is a pooled, unweighted OLS regression with a common intercept and decade-specific slopes on the historical manufacturing-employment peak. Coefficients are per ten-percentage-point higher peak. Columns (1) and (3) include no additional controls; columns (2) and (4) add log GDP per capita in 2018 and region fixed effects. No linear peak-year control, separate decade intercepts, service-share control, or other covariates are included. Income is measured in 2018 in both controlled columns, using Gapminder for the growth sample and Maddison Project Database 2023 for the ETD sample, consistent with their underlying datasets. Growth is the mean annual log GDP-per-capita change over 2010–2019, multiplied by 100. HSPS composition is business, finance, and real-estate employment as a percentage of non-government service employment in 2018; government is excluded from numerator and denominator. Growth columns use 174 countries and GGDC–Gapminder manufacturing histories; composition columns use 51 ETD countries and GGDC–ETD histories. Both specifications within each pair retain the same countries. No ETD country has a recorded peak in the 1960s. HC3 standard errors are in parentheses throughout; inference uses residual-degrees-of-freedom t distributions and pointwise significance levels. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C Historical examples of producer-service codification

The key distinction in the model is between a function and the mode in which that function is performed. Work organization, quality control, inventory management, accounting, procurement, and personnel management exist in many economies. In low-skill mode, these functions are carried out through local experience, tacit knowledge, and informal supervision. In high-skill mode, the same functions are transformed into records, procedures, metrics, job categories, and decision rules. Manufacturing and industrial procurement are natural settings for this transformation because many of the relevant tasks are repeated, measurable, and costly to perform informally. Once codified, however, the resulting routines

can travel into service organizations such as department stores, hospitals, banks, retailers, offices, and universities.

Table 28 summarizes the mapping from the model to several historical examples.

Table 28: Examples of producer-service tasks moving from low-skill to high-skill mode

Producer-service task	Low-skill mode	Industrial codification	Later service sector use
Work organization	Foreman judgment, local rules of thumb, informal supervision	Motion study, workflow design, timing of tasks, written procedures, training systems	Department stores, offices, hospitals, and other service organizations use workflow design, scheduling, clerical organization, and personnel practices
Quality control	Inspection by eye, informal rejection of defective output	Testing protocols, defect records, statistical process control, control charts	Clinical laboratories and hospitals use control charts and statistical quality-control routines to monitor measurement processes
Accounting and managerial control	Bookkeeping and transaction recording	Cost accounting, budgeting, return-on-investment calculations, divisional reporting, performance metrics	Retail chains, department stores, banks, hospitals, universities, and other multi-unit service organizations use budgeting and management-control systems
Inventory, procurement, and logistics	Visual stock checks, ad hoc ordering, relationship-based purchasing	Stock-control systems, procurement planning, supplier coordination, delivery scheduling	Retail chains, wholesalers, hospitals, restaurants, and distribution firms use inventory and logistics routines
Personnel and training	Informal apprenticeship, personal hiring, on-the-job learning	Job analysis, hiring rules, training manuals, personnel departments, performance evaluation	Department stores, hotels, hospitals, offices, and other service organizations use formal hiring and training systems

Work organization: from factory motion study to department stores. A first example is work organization. In low-skill mode, the task of organizing work is performed by foremen and supervisors using experience and local knowledge. Manufacturing made this task easier to codify because factory

operations were repeated and measurable: worker motions, delays, output, and defects could be observed, timed, compared, and redesigned. The scientific-management tradition transformed work organization into a high-skill producer-service function involving motion study, workflow design, written procedures, and training systems.

Lillian Gilbreth provides a useful historical example of diffusion across sectors. Frank and Lillian Gilbreth developed motion-study methods within the scientific-management tradition, initially oriented toward industrial work. Historical accounts emphasize that Lillian Gilbreth later adapted motion study and industrial psychology to service and household settings, including department-store employees and consumers (Graham, 1992, 1999). In particular, her work at Macy's between 1925 and 1928 applied scientific-management and personnel-management ideas to hiring, training, employee handling, clerical organization, and the organization of work within the department store (Graham, 2000). This is the mechanism in the model: manufacturing and industrial work make a producer-service task easier to codify, and the resulting routine can later be used by service organizations.

Quality control: from industrial process control to clinical laboratories. A second example is quality control. In low-skill mode, quality is monitored through direct inspection or the judgment of experienced workers. Manufacturing made quality problems measurable because defects, tolerances, downtime, and output variation could be recorded. This gave rise to statistical process control. Walter Shewhart's work at Bell Laboratories in the 1920s introduced the control chart, and his later book formalized statistical control as a method for monitoring variation in production processes (Shewhart, 1931; Best and Neuhauser, 2006).

Once codified, the same routine was portable outside manufacturing. Clinical laboratories later adopted statistical quality-control methods to monitor the reliability of laboratory measurements. Levey and Jennings introduced statistical quality-control charts for clinical laboratories in 1950, adapting the logic of process control to repeated measurement in medical testing (Levey and Jennings, 1950; Westgard et al., 1981). The task changed from informal inspection to a high-skill routine based on records, thresholds, and statistical decision rules. This example maps directly into the model: codified organizational knowledge K_t makes it easier for a service organization to adopt a high-skill version of a task that was first formalized in an industrial setting.

Accounting and managerial control: from industrial cost control to retail management. A third example is accounting and managerial control. In low-skill mode, accounting records transactions and keeps the books. Large industrial firms and industrial procurement systems created a different problem: managers needed to compare plants, products, inventories, costs, divisions, and capital allocations. This transformed accounting into a high-skill producer-service activity: cost accounting, budgeting, return-on-investment analysis, divisional reporting, and management control. The historical literature on management accounting emphasizes the development of cost accounting and managerial control in large industrial and transportation enterprises, the influence of scientific management, and the later use of formal budgeting and performance measures in large corporations such as DuPont and General Motors (Johnson, 1975; Kaplan, 1984; Chandler, 1977).

The same managerial routines later became portable tools for service firms. James O. McKinsey helped professionalize budgeting and managerial accounting as a consulting service through books such as McKinsey (1922) and McKinsey (1924). In 1935, McKinsey completed a management consulting engagement for Marshall Field's, a large Chicago department-store chain, and then became its chairman and chief

executive (Flesher, 1985). The relevant diffusion is not a machine or a product, but a producer-service routine: a way of measuring, budgeting, and managing a complex organization. In the language of the model, management accounting is a task that moves from bookkeeping in low-skill mode to codified high-skill managerial control.

D Model Proofs and Equilibrium Details

This appendix provides the formal equilibrium definition, derives the reduced establishment problem used in the main text, and proves Propositions 1, 2, and 3.

D.1 Equilibrium details

Producer-service supply. Formal and informal producer services are supplied competitively using skilled and unskilled labor, respectively. If $\mathcal{L}_t^H$ and $\mathcal{L}_t^L$ denote the corresponding employment levels, output is

$$Y_t^H = A_t^H h_s \mathcal{L}_t^H, \quad Y_t^L = A_t^L h_u \mathcal{L}_t^L. \quad (53)$$

Because w_t^s and w_t^u are prices per efficiency unit, perfect competition implies $p_t^H = w_t^s/A_t^H$ and $p_t^L = w_t^u/A_t^L$. Formal-service output is used in gross production, recurring coordination, and the creation and adaptation of procedures. Informal-service output is used in gross production and recurring coordination. The two producer-service productivities follow the value-added productivity law in (29) with exponent ϕ_c .

Aggregate state and prices. The beginning-of-period aggregate state is

$$\mathcal{X}_t \equiv (L_t, S_t, K_t, \{A_{jt}, Z_{jt}, \mu_{jt}\}_{j \in \mathcal{J}}, A_t^H, A_t^L, \{e_{i,t}\}_{i \in \{a,m,s\}}). \quad (54)$$

Here $\mu_{jt}(a, s)$ is the measure of active sector- j establishments over permanent ability and inherited formalization frontiers. Establishment mass, employment, and the employment-weighted formal share are

$$M_{jt} = \int d\mu_{jt}(a, s), \quad N_{jt} = \int n_{jt}^*(a, s) d\mu_{jt}(a, s), \quad x_{jt} = \frac{1}{N_{jt}} \int n_{jt}^*(a, s) q(s) d\mu_{jt}(a, s). \quad (55)$$

The aggregate x_{jt} summarizes the establishment distribution; each establishment's productive and recurring producer-service demands depend on its own inherited share $q(s)$.

Let

$$\mathcal{P}_t \equiv (\{p_{jt}, p_{jt}^c\}_{j \in \mathcal{J}}, p_t^H, p_t^L, w_t^s, w_t^u), \quad p_{mt} \equiv 1, \quad (56)$$

where p_{jt} and p_{jt}^c are producer and consumer prices and manufacturing output is the numeraire. Let $\mathcal{Q}_t$ collect household allocations, establishment input and organizational policies, entry, and producer-service production.

Perfect-foresight competitive equilibrium. Given an initial state $\mathcal{X}_0$ and deterministic paths for exogenous fundamentals and government policies, a perfect-foresight competitive equilibrium is a sequence

$$\{\mathcal{P}_t, \mathcal{Q}_t, \mathcal{X}_{t+1}\}_{t=0}^{\infty} \quad (57)$$

such that, at every date t , the following conditions hold.

1. *Households.* Households choose sectoral consumption, fertility, and education optimally subject to their budget and time constraints. Children choose whether to acquire skill according to the random-utility rule in (27). Households take current prices and dividends as given and correctly anticipate future wages and skill premia.
2. *Final-good establishments.* For every sector j and state (a, s) , an incumbent chooses current labor, materials, and output-linked producer services to solve (14). It chooses its next-period frontier $u \geq s$ according to (21). The inherited share $q(s)$ governs current production and recurring coordination; the chosen frontier affects production from $t + 1$ onward.
3. *Producer-service suppliers.* Formal and informal suppliers maximize profits under (53). Their outputs equal the sums of output-linked demand, recurring organizational demand, and, for formal services, the inputs used to create and adapt procedures.
4. *Entry.* Potential entrants pay the setup cost before drawing $a \sim G$ and enter with $s = 0$. Entry and establishment counts satisfy (23).
5. *Market clearing and resource feasibility.* Agricultural, manufacturing, and consumer-service output is allocated to household consumption, intermediate use, and entry setup. Formal and informal producer-service markets clear separately. Demand for skilled and unskilled efficiency units equals the corresponding household supply. Establishment profits net of organizational investment, entry costs, and taxes are distributed to households, and the government budget balances.
6. *Aggregation and laws of motion.* Individual allocations aggregate according to (55). Let $u_{jt}^*(a, s)$ be the organizational policy implied by the Bellman equation. For every measurable set B of establishment states,

$$\begin{aligned} \mu_{j,t+1}(B) &= (1 - \delta_f) \int \mathbb{1}\{(a, u_{jt}^*(a, s)) \in B\} d\mu_{jt}(a, s) \\ &\quad + I_{j,t+1} \int \mathbb{1}\{(a, 0) \in B\} dG(a). \end{aligned} \quad (58)$$

Public knowledge follows (22); population, skills, and productivities follow (28)–(29).

7. *Perfect foresight.* Forecasts of future aggregate prices and states coincide with their equilibrium realizations: for every $k \geq 1$,

$$\hat{\mathcal{P}}_{t+k|t} = \mathcal{P}_{t+k}, \quad \hat{\mathcal{X}}_{t+k|t} = \mathcal{X}_{t+k}. \quad (59)$$

Perfect foresight concerns aggregate variables. Establishments still face idiosyncratic exit risk, and entrants draw heterogeneous ability. With a continuum of establishments, those risks aggregate according to their known distributions and generate no aggregate uncertainty.

D.2 Proof of Proposition 1

I first derive the reduced static establishment problem used in the proof. Conditional cost minimization in the CES labor nest gives

$$w_{jt}^v = \left[\alpha_j^{\sigma_L} (w_t^s)^{1-\sigma_L} + (1 - \alpha_j)^{\sigma_L} (w_t^u)^{1-\sigma_L} \right]^{\frac{1}{1-\sigma_L}}. \quad (60)$$

Because the labor aggregator is homogeneous of degree one, headcount is $n = \nu_{jt}\ell$, where

$$\nu_{jt} = \frac{\alpha_j^{\sigma_L}}{h_s} \left(\frac{w_{jt}^v}{w_t^s} \right)^{\sigma_L} + \frac{(1 - \alpha_j)^{\sigma_L}}{h_u} \left(\frac{w_{jt}^v}{w_t^u} \right)^{\sigma_L}. \quad (61)$$

Recurring organizational cost per unit of the labor composite is therefore

$$c_{jt}^o(q) = \frac{\nu_{jt}}{N} [qp_t^H + r(1 - q)p_t^L]. \quad (62)$$

Fix (j, t, a, s) and write $q = q(s)$. Before conditional inputs are eliminated, the establishment solves

$$g_{jt}^*(a, s) = \max_{\ell, m^a, m^m, H^p, L^p \geq 0} \left\{ p_{jt} Z_{jt} (A_{jt} a \ell^\eta)^{\theta_{jt}} (m^a)^{\gamma_{jt}^a} (m^m)^{\gamma_{jt}^m} [P^p(H^p, L^p; q)]^{\gamma_{jt}^P} \right. \\ \left. - w_{jt}^v \ell - p_{at} m^a - p_{mt} m^m - p_t^H H^p - p_t^L L^p - c_{jt}^o(q) \ell \right\}. \quad (63)$$

For fixed P^p , CES duality implies expenditure $\omega_t^P(q) P^p$, where

$$\omega_t^P(q) = [q(p_t^H)^{1-\sigma_P} + (1 - q)(p_t^L)^{1-\sigma_P}]^{\frac{1}{1-\sigma_P}}.$$

Next fix value added $v = A_{jt} a \ell^\eta$ and optimize over (m^a, m^m, P^p) . If $\mathcal{R}$ denotes gross revenue at the conditional optimum, the first-order conditions imply

$$p_{at} m^a = \gamma_{jt}^a \mathcal{R}, \quad p_{mt} m^m = \gamma_{jt}^m \mathcal{R}, \quad \omega_t^P(q) P^p = \gamma_{jt}^P \mathcal{R}. \quad (64)$$

Substitution into the gross-output technology gives

$$\mathcal{R} = \left[p_{jt} Z_{jt} \left(\frac{\gamma_{jt}^a}{p_{at}} \right)^{\gamma_{jt}^a} \left(\frac{\gamma_{jt}^m}{p_{mt}} \right)^{\gamma_{jt}^m} \left(\frac{\gamma_{jt}^P}{\omega_t^P(q)} \right)^{\gamma_{jt}^P} \right]^{\frac{1}{\theta_{jt}}} v. \quad (65)$$

Conditional expenditure on materials and productive producer services is $(1 - \theta_{jt})\mathcal{R}$, so revenue net of these inputs is $\theta_{jt}\mathcal{R} = p_{jt}^v(q)v$. Hence the static problem reduces to

$$g_{jt}^*(a, q) = \max_{\ell \geq 0} \{ p_{jt}^v(q) A_{jt} a \ell^\eta - [w_{jt}^v + c_{jt}^o(q)] \ell \}. \quad (66)$$

The unique interior solution satisfies

$$\ell_{jt}^*(a, q) = \left[\frac{\eta p_{jt}^v(q) A_{jt} a}{w_{jt}^v + c_{jt}^o(q)} \right]^{\frac{1}{1-\eta}}, \quad n_{jt}^*(a, q) = \nu_{jt} \ell_{jt}^*(a, q). \quad (67)$$

Proof of Proposition 1. By the envelope theorem,

$$\begin{aligned}\frac{\partial g_{jt}^*(a, q)}{\partial q} &= v_{jt}^*(a, q) \frac{\partial p_{jt}^v(q)}{\partial q} - \ell_{jt}^*(a, q) \frac{\partial c_{jt}^o(q)}{\partial q} \\ &= v_{jt}^*(a, q) \frac{\partial p_{jt}^v(q)}{\partial q} + \frac{n_{jt}^*(a, q)}{\bar{N}} (rp_t^L - p_t^H),\end{aligned}\tag{68}$$

where the second equality uses

$$\frac{\partial c_{jt}^o(q)}{\partial q} = \frac{\nu_{jt}}{\bar{N}} (p_t^H - rp_t^L).$$

The first-order condition for ℓ implies

$$\eta p_{jt}^v(q) v_{jt}^*(a, q) = [w_{jt}^v + c_{jt}^o(q)] \frac{n_{jt}^*(a, q)}{\nu_{jt}}.$$

Substituting into (68) gives

$$\frac{\partial g_{jt}^*(a, q)}{\partial q} = n_{jt}^*(a, q) \left[\frac{w_{jt}^v + c_{jt}^o(q)}{\eta \nu_{jt}} \frac{\partial \log p_{jt}^v(q)}{\partial q} + \frac{rp_t^L - p_t^H}{\bar{N}} \right] = n_{jt}^*(a, q) \Psi_{jt}(q).$$

Thus, whenever $\Psi_{jt}(q) > 0$, the marginal operating gain from formalization is positive and proportional to employment.

It remains to establish the converse effect from formalization to desired scale. Differentiating (67) gives

$$\frac{\partial \log n_{jt}^*(a, q)}{\partial q} = \frac{1}{1 - \eta} \left[\frac{\partial \log p_{jt}^v(q)}{\partial q} + \frac{\nu_{jt}}{\bar{N}} \frac{rp_t^L - p_t^H}{w_{jt}^v + c_{jt}^o(q)} \right].\tag{69}$$

If $\Psi_{jt}(q) > 0$, then

$$\frac{\partial \log p_{jt}^v(q)}{\partial q} > -\frac{\eta \nu_{jt}}{\bar{N}} \frac{rp_t^L - p_t^H}{w_{jt}^v + c_{jt}^o(q)}.$$

Hence, when $rp_t^L > p_t^H$, the bracket in (69) is strictly positive. Therefore

$$\frac{\partial \log n_{jt}^*(a, q)}{\partial q} > 0.$$

This proves both statements of the proposition. □

D.3 Proof of Proposition 2

I first derive the household decision rules used in the proof. Define

$$\zeta_t \equiv 1 + \frac{d_t}{\bar{w}_t}, \quad \chi_t \equiv \tau_e(S_t).$$

For a household of employment origin i , the time cost of one child receiving education $e_{i,t}$ is

$$T_{i,t}(e_{i,t}) = \tau_i + \chi_t e_{i,t}.$$

Conditional on the aggregate growth and skill environment, the first-order condition for fertility gives

$$n_{i,t} = \frac{\gamma \zeta_t}{\tau_i + \tau_e(S_t)e_{i,t}}. \quad (70)$$

For an interior education choice, the first-order condition is

$$\frac{g_t}{(e_{i,t} + \tau_h)(e_{i,t} + \tau_h + g_t)} = \frac{\tau_e(S_t)}{\tau_i + \tau_e(S_t)e_{i,t}},$$

which rearranges to

$$\tau_e(S_t)(e_{i,t} + \tau_h)^2 = g_t[\tau_i - \tau_e(S_t)\tau_h].$$

Hence

$$e_{i,t}^* = \left[\sqrt{g_t \left(\frac{\tau_i}{\tau_e(S_t)} - \tau_h \right)} - \tau_h \right]_+, \quad (71)$$

where $[x]_+ \equiv \max\{x, 0\}$. Aggregate fertility is

$$\bar{n}_t = \sum_{i \in \{a, m, s\}} \lambda_{i,t} n_{i,t}. \quad (72)$$

Proof of Proposition 2. Fix $S \in (0, 1)$ and define

$$t \equiv \tau_e(S), \quad a \equiv \gamma \zeta(S) - \tau_s, \quad b \equiv \tau_s - t\tau_h.$$

The maintained interiority conditions imply $a > 0$ and $b > 0$. Equations (71), (70), and (29) reduce in the terminal service economy to

$$g(L) = \left[1 + \lambda_A(SL)^{\phi_c} \right]^{1/\Delta} - 1, \quad (73)$$

$$e(L) = \left[\sqrt{\frac{b}{t} g(L)} - \tau_h \right]_+, \quad (74)$$

$$n(L) = \frac{\gamma \zeta(S)}{\tau_s + te(L)}. \quad (75)$$

Population therefore follows $L_{t+1} = T(L_t)$, where $T(L) = Ln(L)^{1/\Delta}$.

The function $g(L)$ is continuous and strictly increasing from zero to infinity. Hence $e(L)$ is continuous and nondecreasing, while $n(L)$ is continuous and nonincreasing. Moreover,

$$n(0) = \frac{\gamma \zeta(S)}{\tau_s} > 1, \quad \lim_{L \rightarrow \infty} n(L) = 0.$$

There is therefore a unique $L^\infty(S) > 0$ satisfying $n(L^\infty(S)) = 1$.

To establish global stability, first consider an interval in which education is interior. Let $\epsilon_g(L) = d \log g(L) / d \log L$. Concavity of $[1 + x]^{1/\Delta} - 1$ implies $0 < \epsilon_g(L) \leq \phi_c$. From (74),

$$\frac{de(L)}{d \log L} = \frac{1}{2}[e(L) + \tau_h]\epsilon_g(L).$$

Since $b > 0$ implies $t[e(L) + \tau_h] < \tau_s + te(L)$,

$$-\frac{d \log n(L)}{d \log L} = \frac{t}{\tau_s + te(L)} \frac{de(L)}{d \log L} < \frac{\epsilon_g(L)}{2} \leq \frac{\phi_c}{2}.$$

Therefore

$$\frac{d \log T(L)}{d \log L} = 1 + \frac{1}{\Delta} \frac{d \log n(L)}{d \log L} > 1 - \frac{\phi_c}{2\Delta} > 0.$$

Where education is at its corner, $n(L)$ is locally constant and $T(L)$ remains strictly increasing. Thus T is globally increasing. If $L < L^\infty(S)$, then $L < T(L) < L^\infty(S)$; if $L > L^\infty(S)$, then $L > T(L) > L^\infty(S)$. Hence the population sequence is monotone and bounded and converges to $L^\infty(S)$ from every $L_0 > 0$, while fertility converges to replacement.

At the fixed point, (75) implies

$$e^\infty(S) = \frac{\gamma\zeta(S) - \tau_s}{\tau_e(S)} = \frac{a}{t}.$$

Using the interior education rule then gives

$$\begin{aligned} g^\infty(S) &= \frac{[a + t\tau_h]^2}{tb} \\ &= \frac{[\gamma\zeta(S) - \tau_s + \tau_e(S)\tau_h]^2}{\tau_e(S)[\tau_s - \tau_e(S)\tau_h]} = \mathcal{G}(S), \end{aligned}$$

which proves (30). Inverting (73) at $g(L^\infty) = \mathcal{G}(S)$ gives (31).

Finally, since $t'(S) = -(1 - \tau_e^s) < 0$, logarithmic differentiation yields

$$\begin{aligned} \frac{\mathcal{G}'(S)}{\mathcal{G}(S)} &= \frac{2\gamma\zeta'(S)}{a + t\tau_h} \\ &\quad + t'(S) \left[\frac{2\tau_h}{a + t\tau_h} - \frac{1}{t} + \frac{\tau_h}{\tau_s - t\tau_h} \right]. \end{aligned} \tag{76}$$

The first term is nonnegative when $\zeta'(S) \geq 0$. Since $t'(S) < 0$, the second term is strictly positive whenever the bracketed expression is negative. Rearranging this condition gives

$$\tau_h = 0 < \frac{\tau_s[\gamma\zeta(S) - \tau_s]}{\tau_e(S)[2\gamma\zeta(S) - \tau_s]},$$

Hence $\mathcal{G}'(S) > 0$ under the conditions stated in the proposition. $\square$

D.4 Proof of Proposition 3

Proof of Proposition 3. Consider first the terminal service economy. For a stationary service establishment, the adaptation cost in (19) and the discounted recurring-coordination gain imply

$$q_f(K) = \min \left\{ K, \left[1 - \frac{1}{e\mathcal{A}_f} \right]_+ \right\}.$$

The employment-weighted formalized-task share,

$$x(K) = \frac{1}{N^s} \int n_f q_f(K) d\mu^s(f),$$

is therefore weakly increasing in public organizational knowledge. Together with the producer-service demands in (11), the skill-supply rule (27), and labor-market clearing, this implies

$$K_{A,3} > K_{B,3} \iff S_A^\infty > S_B^\infty.$$

Proposition 2 then gives

$$g_A^\infty > g_B^\infty \iff K_{A,3} > K_{B,3}. \quad (77)$$

It remains to determine when the manufacturing path generates the larger public stock. Under the period-2 restrictions used in the proposition, the creation cutoff and diffusion law imply

$$K_{A,3} > K_{B,3} \iff \bar{e}_m^o n_m > \bar{e}_c^o n_c.$$

Thus the comparison depends on education embodied in organizational labor and on the firm scale over which newly created procedures are used.

Taking $\tau_h \rightarrow 0$ in (71) gives

$$\frac{\bar{e}_m^o}{\bar{e}_c^o} = \left(\frac{g_m \tau_m}{g_c \tau_s} \right)^{1/2}.$$

With $\Delta = 1$, (29) implies

$$\frac{g_m}{g_c} = (S_1 L_1)^{\phi_m - \phi_c}.$$

Under specialization and symmetry, establishment value is proportional to employment and entry is inversely related to average establishment size. Combining these relationships with (23), under the assumed equality of price and operating-surplus terms across the two paths, gives

$$\frac{n_m}{n_c} = \left(\frac{F_m}{F_c} \right)^{\frac{1}{1+\phi_E}}.$$

Consequently,

$$\frac{\bar{e}_m^o n_m}{\bar{e}_c^o n_c} = \left[\frac{\tau_m}{\tau_s} (S_1 L_1)^{\phi_m - \phi_c} \right]^{1/2} \left(\frac{F_m}{F_c} \right)^{\frac{1}{1+\phi_E}}.$$

The manufacturing path therefore generates more public organizational knowledge if and only if

$$\frac{\tau_m}{\tau_s} (S_1 L_1)^{\phi_m - \phi_c} \left(\frac{F_m}{F_c} \right)^{\frac{2}{1+\phi_E}} > 1,$$

which is exactly condition (32). Combining this result with (77) proves the proposition. $\square$

Table 29: Calibration moment blocks

Block	Rows	Measurement
Sectoral employment paths	48	Three sectors at 16 benchmark years
Manufacturing peak share	1	Maximum manufacturing employment share
Manufacturing peak timing	1	Interval target, 1930–1960
Fertility path	16	Aggregate fertility at benchmark years
Sectoral fertility ratio	1	Agriculture relative to services
GDP-per-capita path	18	U.S. series normalized to one in 1850
Sectoral VA productivity growth	3	Annual means, 1970–2010
Aggregate skilled share	1	Employment share, 2019
Sectoral HSPS shares	3	Means over 2015–2020
Aggregate HSPS share	1	Sector-weighted accounting aggregate
Within-VA skilled shares	3	Means over 2015–2020
Skilled earnings premium	1	Skilled/unskilled earnings, 2019
Sectoral firm size	3	Employees per employer firm, 2019
Total	100	

E Calibration Details

This appendix describes the construction of the calibration targets, the parameters assigned outside the simulated-moments search, the numerical implementation, and the local identification and robustness exercises. The main text reports the parameter estimates and summarizes the fit.

E.1 Calibration moments and measurement

The calibration uses 100 moment rows collected in 13 blocks. For moment k in block b , the normalized discrepancy is

$$r_{bk}(\Theta) = \frac{m_{bk}(\Theta) - \hat{m}_{bk}}{\max\{|\hat{m}_{bk}|, 0.01\}}. \quad (78)$$

The objective averages squared discrepancies within each block before applying the block weights:

$$\mathcal{L}(\Theta) = \frac{\sum_b \omega_b K_b^{-1} \sum_{k=1}^{K_b} r_{bk}(\Theta)^2}{\sum_b \omega_b}. \quad (79)$$

This construction prevents a path from receiving more weight merely because it contains more observations. Table 29 lists the rows in each block. The peak-date target is an interval: dates from 1930 to 1960 receive zero timing loss, with a quadratic penalty outside that interval.

Producer-service employment. The distinction between producer and consumer services is based on use rather than the supplier’s industry. I combine the 2019 IPUMS USA microdata with the 2019 BEA Summary Use Table. Within agriculture, manufacturing, and consumer-service industries, occupations are first divided between sectoral value-added work and organizational or technical producer-service work. Producer-service workers are then classified by skill, with high skill defined as at a bachelor’s degree. For industries that primarily supply intermediate services, employment is allocated across destination sectors in proportion to BEA intermediate-use expenditures. The share sold directly to households

is assigned to consumer-service value added. The IPUMS–BEA crosswalk and occupational classification are included in the replication files.

The allocation imposes two restrictions. First, workers supplying different buyers from the same industry have the same skill composition. Second, a dollar of output from that industry has the same average labor requirement regardless of the purchaser. These restrictions are needed because available input-output tables do not link a transaction to the characteristics of the workers who produced it. Applying the same rule to model and data yields the three HSPS shares and the three within-value-added skill shares used in the calibration. The aggregate HSPS moment uses the same sector weights in model and data and is therefore an accounting check on the sectoral construction, not an independent fourth source of variation.

Firm scale. The 2019 scale targets are employees per employer *firm*, not employees per establishment, from the Census Statistics of U.S. Businesses. The targets are 7.435 workers in agriculture, 49.694 in manufacturing, and 20.454 in broad services. The model’s producing unit is interpreted as the empirical employer firm throughout this comparison. Including all service industries makes the service target conservative relative to a narrower employer-sector definition. Firm scale in the historical initial condition is anchored to the 1880 Census of Manufactures. The modern scale moments primarily discipline the three entry costs; the common entry-supply curvature and the dispersion of firm ability are assigned externally.

Productivity, skill, and demand. Sectoral productivity targets are mean annual growth rates of value added per worker from 1970 through 2010. GDP-per-capita is normalized to one in 1850, so its path identifies growth rather than a level unit. The terminal skill target is the employment share with college education, and the skilled earnings premium is the ratio of skilled to unskilled earnings. Sectoral employment shares are observed at 16 dates between 1850 and 2010. Their levels and transition paths, together with the manufacturing peak, discipline the nonhomothetic demand system.

E.2 Assigned parameters and normalizations

Table 30 collects the parameters fixed outside the 19-dimensional calibration. Some parameters set units or reference normalizations, while others are assigned using external empirical evidence or maintained at model-specific benchmark values. Where the literature disciplines a closely related rather than identical object, I use it to anchor the relevant empirical magnitude rather than interpret the reported value as a direct estimate. The generation length is 20 years, following Lagerlöf (2006); the elasticity of substitution across final goods is 0.30, as in Comin et al. (2021); and the two skill-substitution elasticities are set to 1.20, consistent with the estimates surveyed by Krusell et al. (2000) and Acemoglu and Autor (2011). The input-output structure begins with value-added production in 1850 and converges gradually to the modern BEA coefficients by 2025, allowing intermediate producer services to emerge endogenously over the transition.

The organizational-scale benchmark $\bar{N} = 25$ implies that recurring coordination absorbs 5.42 percent of employment in agriculture, 5.04 percent in manufacturing, and 5.19 percent in consumer services in 2019. This scale is consistent in order of magnitude with the organization-capital evidence in Atkeson and Kehoe (2005). The external evidence disciplines the magnitude of recurring organizational inputs rather than providing a direct estimate of $\bar{N}$.

Table 30: Assigned parameters, external anchors, and normalizations

Parameter	Value	Interpretation or source
<i>A. Units and reference normalizations</i>		
$p_m; h_u; \beta_\pi$	1; 1; 1	Manufacturing-price, unskilled-efficiency, and skill-index units
$\log A_{j0}$	0	Initial value-added productivity normalization
τ_e^u	1	Education-time unit
$\epsilon_m; w_m$	1; 2.5	Manufacturing-demand reference values
$\kappa; \xi_m$	1; 1	Task-complexity scale and manufacturing-adaptability normalization
<i>B. Parameters disciplined by external evidence</i>		
β	0.96	Annual discount factor; Kydland and Prescott (1982)
δ_f	0.04	Annual establishment exit rate; Asturias et al. (2023)
η	0.85	Span-of-control curvature; Atkeson and Kehoe (2005)
$\bar{N}$	25	Organizational-scale benchmark; recurring coordination absorbs about 5 percent of employment; Atkeson and Kehoe (2005)
σ_a	0.16	Dispersion of firm ability, calculated from World Bank Enterprise Survey data
ϕ_E	1.552	Entry-supply curvature; Gutiérrez et al. (2021)
σ	0.30	Elasticity of substitution across final goods; Comin et al. (2021)
$\sigma_L; \sigma_P$	1.20; 1.20	Substitution across skill types; Krusell et al. (2000) ; Acemoglu and Autor (2011)
Δ	20	Generation length in years; Lagerlöf (2006)
λ_A	0.015	Productivity-growth scale; Fernald (2015)
h_s	1.60	Skilled-efficiency scale, benchmarked to the U.S. college earnings premium; Acemoglu and Autor (2011)
δ_K	0.09	Annual diffusion benchmark; Mansfield (1961) ; Greenstein (1994)
<i>C. Model-specific maintained benchmarks</i>		
σ_{skill}	0.10	Scale of the smooth education-choice margin; McFadden (1974)
τ_e^s	0.50	Skill-education cost complementarity; Cunha and Heckman (2007) .
$\tau_m^n; \tau_s^n$	1; 1	Common child-time cost outside agriculture ^a
ρ	0.40	Adaptation cost relative to creation cost; sensitivity reported below
γ	0.05	Child-preference weight, benchmarked to an education-expenditure share of approximately 5 percent
ϵ_a	0.01	Near-subsistence agricultural-demand benchmark; Kongsamut et al. (2001)

Notes: Panel A contains unit choices and reference normalizations. Panel B contains parameters fixed before the joint calibration using external empirical evidence. Panel C contains model-specific benchmark values that are not directly estimated from the calibration moments. The citations in Panel C motivate the relevant functional form or empirical scale and are not interpreted as one-for-one estimates of the reported parameter.

The ability dispersion σ_a is calculated from the pooled, employment-weighted dispersion of firm employment in the World Bank Enterprise Surveys. The knowledge parameter δ_K governs the conversion of exposure to privately created routines into additions to the public stock of organizational knowledge; it is not a physical depreciation rate.

^a [Cunha and Heckman \(2007\)](#) motivate the education-environment complementarity. I set $\tau_e^s = 0.5$ as a conservative benchmark: at the calibrated skilled share of $S \simeq 0.34$, it implies that the time cost of education is only about 17 percent lower than in a fully unskilled environment.

^b Setting $\tau_m^n = \tau_s^n = 1$ shuts down systematic differences in the fixed time cost of children between manufacturing and service households. This is consistent with the absence of a persistent manufacturing-service fertility gap in the U.S. historical data and with the developing-country evidence discussed by [Peters et al. \(2026\)](#).

The diffusion parameter. Let E_t denote the employment-weighted exposure to privately created routines that have not yet entered the public stock. The maintained accumulation law is

$$K_{t+1} = K_t + \min\{1 - K_t, \delta_K E_t\}. \quad (80)$$

Thus δ_K converts current exposure into an annual addition to the stock of publicly usable organizational knowledge.

The empirical literature provides a useful scale for this speed, although it does not measure the object in equation (80) one-for-one. In his study of 12 major process innovations across four industries, [Mansfield \(1961\)](#) reports that the time required for one-half of firms to adopt ranged from 0.9 to 15 years, with a mean of 7.8 years. Under a constant annual closing rate, a half-adoption time of T years corresponds to $d(T) = 1 - 2^{-1/T}$. Mansfield's mean therefore implies $d(7.8) = 0.085$, while the slow end of his range implies $d(15) = 0.045$. [Greenstein \(1994\)](#) finds that average practice in mainframe computing advanced by 14.4 percent per year from 1968 to 1983 and lagged best practice by at most six to seven years. We summarize these results with a slow annual benchmark of 0.04 and a fast benchmark of 0.14, and set $\delta_K = 0.09$, their arithmetic midpoint. The mean Mansfield half-adoption time independently maps to essentially the same value.

This mapping is intentionally approximate. Mansfield measures interfirm adoption of capital-embodied process innovations; Greenstein measures the installed performance of computer hardware; the model concerns the public availability of organizational routines after firms create them. The two studies are therefore used to discipline the annual scale of diffusion, not as direct structural estimates of δ_K . Section E.6 accordingly varies δ_K from 0.01 to 0.15 and also changes the functional form of the accumulation law.

The other four externally assigned parameters in the last part of Table 30 are held fixed before the joint search. The calibration then jointly chooses

$$\Theta = \{\tau, \tau_a^n, \phi_a, \phi_m, \phi_c, g_{GO}, c_{col}, \alpha_a, \alpha_m, \alpha_c, \epsilon_c, w_a, w_c, R, \xi_a, \xi_c, F_a, F_m, F_c\}. \quad (81)$$

The moments identify these parameters jointly rather than through exact one-parameter/one-moment assignments. Broadly, fertility moments discipline τ and τ_a^n ; the GDP and sectoral productivity paths discipline ϕ_j and g_{GO} ; skill shares and the earnings premium discipline c_{col} and α_j ; structural-change paths discipline the demand parameters; HSPS shares discipline R , ξ_a , and ξ_c ; and the three firm-size targets discipline F_j . In particular, ξ_c is a jointly calibrated parameter, not a normalization.

E.3 Numerical implementation

The model is solved annually from 1790 through 2289. Extending the terminal horizon beyond the reported outcomes prevents the 2019 equilibrium and the policy transitions from being driven by terminal conditions. The exact calibration uses nine firm-productivity types and 101 organizational-task grid points. Entry, firm scale, education, formalization, and the public knowledge stock are solved jointly.

All strictly positive calibrated parameters are searched in logs. The search uses multiple starting points and damped local updates, followed by an exact-grid evaluation of every retained candidate. A candidate is admissible only if the equilibrium converges on the exact grid and no 2019 sectoral firm-size error exceeds three percent. The selected vector has $\mathcal{L} = 0.011878$ and a maximum firm-size error of 0.08 percent. From the cold start, the fixed point is accepted after 30 iterations. The last update in the discrete formalization block is 0.00646, below the explicit three-grid-cell epsilon-equilibrium tolerance of 0.03; the corresponding 2019 update is 0.00171. Education-budget and size-accounting residuals are zero to machine precision.




Figure 12: Targeted calibration fit. Bars compare the calibrated model with the corresponding U.S. targets. HSPS shares are means over 2015–2020, firm size is measured in 2019, and productivity growth is the annual average over 1970–2010.

To assess numerical stability, I solve the exact-grid equilibrium from 12 independent cold starts. All 12 initializations converge to the same equilibrium, up to the numerical tolerance used by the solver. The calibration and cross-country counterfactuals are likewise solved independently on the exact grid. Finite-difference derivatives are computed around the converged benchmark equilibrium, with the benchmark solution used as the initialization for nearby perturbations to improve computational efficiency and numerical comparability.

E.4 Detailed fit

Figure 12 shows the targeted modern moments most closely connected to the organizational mechanism. Panel (a) compares sectoral HSPS shares, panel (b) compares sectoral firm size, and panel (c) compares value-added productivity growth. The model HSPS shares are 0.325, 0.461, and 0.399 in agriculture, manufacturing, and services, compared with targets of 0.325, 0.461, and 0.399. The largest discrepancy is 0.04 percentage points. Model firm sizes are 7.44, 49.72, and 20.45 workers, all within 0.08 percent of their targets. Annual model productivity growth is 3.85, 2.98, and 1.31 percent, compared with 3.71, 3.17, and 1.42 percent in the data.

The model also generates meaningful cross-firm variation in formalization without targeting dispersion or introducing sector-specific dispersion parameters. The employment-weighted standard deviations of formalization intensity are 0.176 in agriculture, 0.187 in manufacturing, and 0.049 in services.




Figure 13: Identification diagnostics. Panel (a) reports singular values of the exact weighted-moment Jacobian. Panel (b) evaluates the exact objective over ξ_c while holding the other 18 calibrated parameters at their reported values.

The corresponding data-based proxies are 0.229, 0.205, and 0.211. Thus the model reproduces 76.8 percent of the agricultural dispersion, 91.1 percent of the manufacturing dispersion, and 23.1 percent of the service dispersion. None of these dispersion moments enters the calibration objective.

E.5 Local identification and uncertainty

Let $r(\theta)$ denote the vector of weighted residuals and let $\theta = \log \Theta$ denote the 19-dimensional vector of calibrated parameters. I compute the Jacobian

$$J = \frac{\partial r(\theta)}{\partial \theta'} \quad (82)$$

using exact-grid central differences around the frozen benchmark equilibrium. Every perturbation uses the same benchmark-equilibrium anchor. For parameter p and step h , the numerical derivative is

$$J_{\cdot p}(h) = \frac{r(\theta + he_p) - r(\theta - he_p)}{2h}. \quad (83)$$

I use $h = 0.025$ and $h = 0.05$. All 77 required exact simulations—one baseline and four perturbations for each of 19 parameters—converge.

The maintained 100-by-19 objective Jacobian is algebraically full column rank at both steps. At the preferred five-percent step its condition number is 952.1; its effective rank is 17 at a one-percent relative singular-value cutoff and 19 at a 0.1-percent cutoff. I also form a smooth, nonredundant Jacobian by dropping the locally flat peak-timing penalty, the aggregate-HSPS accounting identity, and one redundant sectoral employment share at each date. This 82-by-19 Jacobian is algebraically full column rank, with condition number 789.9. Its effective rank is again 17 at a one-percent cutoff and 19 at a 0.1-percent cutoff. The calibration therefore has two weak local directions, but no exact local redundancy among the 19 calibrated parameters. Every parameter retains at least 14.0 percent of its derivative norm after projection on the other 18 columns. The identified subspace is stable across the two finite-difference steps except for the weakest singular direction.

Table 31: Parameter-level local identification diagnostics

Parameter	Estimate	Unique share (%)	HC1 relative SE (%)
τ	0.03236	17.92	4.86
τ_a^n	0.44017	30.66	5.92
ϕ_a	0.99840	15.49	10.15
ϕ_m	0.91919	14.04	10.38
ϕ_c	0.69046	15.97	10.86
gGO	0.00318	61.33	24.31
c_{col}	0.56821	35.90	5.36
α_a	0.24338	81.42	2.81
α_m	0.33001	84.90	2.96
α_c	0.40995	82.84	2.62
ϵ_c	1.01048	24.87	30.15
w_a	0.65105	34.49	12.44
w_c	3.82326	25.57	39.53
R	22.78696	33.06	5.33
ξ_a	3.56402	64.17	88.56
ξ_c	6.87739	70.22	491.31
F_a	41.24479	80.00	48.09
F_m	277.23883	69.56	60.27
F_c	275.73211	91.00	22.85

Notes: The unique shares use the exact, smooth, nonredundant 82-by-19 Jacobian. HC1 standard errors are computed in log-parameter space and translated to levels by the delta method.

The service-adaptability parameter illustrates the distinction between local identification and precision. The exact objective falls from 0.014807 at $\xi_c = 1$ to 0.011901 at $\xi_c = 3$ and reaches 0.0118833 at the jointly calibrated value $\xi_c = 6.877$. The objective is nearly flat within the saturation region: it is 0.0118837 at $\xi_c = 6.5$, 0.0118832 at 7, and 0.0118829 at 8. The service formalization moments therefore reject low adaptability and place the economy near the saturation region, but do not precisely locate ξ_c within that region. This is reflected in its large local standard error.

Table 31 reports parameter-level diagnostics from the smooth, nonredundant system. The unique share is the norm of a Jacobian column after projection on the other columns, divided by its original norm. The final column reports the HC1 residual-based minimum-distance standard error in percent of the estimate. These quantities summarize local curvature and residual fit conditional on the target vector; they are not sampling confidence intervals for the heterogeneous empirical sources.

For formal SMM inference, let Ω denote the sampling covariance matrix of the empirical moments and W the calibration weight matrix. The corresponding sandwich covariance is

$$\text{Avar}(\hat{\theta}) = (J'WJ)^{-1}J'W\Omega WJ(J'WJ)^{-1}. \quad (84)$$

Because the targets combine historical time series, administrative aggregates, and microdata moments, a sampling implementation of Ω must preserve serial dependence and cross-moment covariance within each source. The rank, objective-slice, and residual-based uncertainty exercises establish local numerical identification of the 19-parameter calibration without treating the heterogeneous moment rows as a common random sample.

E.6 Robustness and Sensitivity

This subsection examines which assumptions matter for the quantitative relationship between manufacturing and subsequent development. Two exercises are considered: a broad sensitivity analysis covering 18 alternative specifications, and a more detailed exploration of the knowledge-diffusion parameter

Table 32: Sensitivity of cross-country slopes

Specification	GDP level	GDP growth	Tertiary share	HSPS share	Population growth
Adaptation cost: -25%	0.8	1.4	3.3	1.8	0.6
Adaptation cost: +25%	0.1	2.5	3.6	0.4	0.5
Knowledge diffusion: -25%	2.4	7.7	11.3	4.6	2.3
Knowledge diffusion: +25%	-4.4	-11.2	-20.8	-8.4	-4.6
Routine repetition: -25%	-7.0	16.3	4.2	-16.4	-8.2
Routine repetition: +25%	1.8	-11.1	-9.0	5.5	3.0
Span of control: 0.80	15.1	13.0	27.2	0.8	13.7
Span of control: 0.90	-10.4	-9.8	-12.7	2.6	-12.0
Producer-service substitution: 0.90	0.1	0.2	0.0	0.2	0.0
Producer-service substitution: 1.50	-0.4	-0.9	-1.3	-0.8	-0.4
Firm-ability dispersion: -25%	5.0	17.5	27.3	6.9	4.0
Firm-ability dispersion: +25%	-2.3	-5.6	-4.4	-2.1	-1.6
Education complementarity: $\tau_e^s = 0.25$	3.5	-2.5	-11.6	-2.0	-9.1
Education complementarity: $\tau_e^s = 0.75$	-2.0	-0.5	8.8	1.6	8.3
Initial knowledge: $K_0 = 0.10$	-5.0	-9.2	-24.4	-3.1	-6.0
Equal sector weights	-16.7	-13.4	-38.5	-24.0	-20.1
No knowledge spillovers	-30.4	-33.8	-77.9	-56.6	-34.3
Exponential-hazard accumulation	-0.1	0.5	0.6	0.0	-0.2

Notes: Entries are percentage changes in the absolute magnitude of the cross-country slope relative to the benchmark. Positive entries indicate a stronger relationship; negative entries indicate attenuation. All slopes retain their benchmark sign. Outcomes are log GDP per capita in 2019, annual GDP-per-capita growth over 2010–2019, average tertiary-education and HSPS shares over 2010–2019, and cumulative log population growth over 1850–2019. Each specification uses the same five manufacturing-demand economies. Parameters not listed in a row remain fixed. The benchmark and 18 alternatives comprise 95 economies, all of which satisfy the numerical convergence criterion. This broader exercise uses a coarser numerical approximation than Table 33; all percentage changes are calculated relative to the benchmark computed with the same approximation.

δ_K . In each exercise, parameters not being varied remain at their calibrated values; the model is not recalibrated. Cross-country slopes are estimated across the same five economies, which differ in their manufacturing-demand weights. The signs of all five headline relationships are preserved throughout: a higher peak manufacturing employment share is associated with higher GDP per capita, faster GDP-per-capita growth, larger tertiary-education and HSPS shares, and lower population growth. The magnitudes are most sensitive to the extent of knowledge diffusion and to parameters governing firm scale. They are substantially less sensitive to adaptation costs, producer-service substitution, and the functional form of knowledge accumulation.

Table 32 reports the percentage change in each slope relative to the benchmark. Across the 18 specifications, all 90 slopes retain their benchmark sign. Median absolute changes are 2.9 percent for GDP levels, 8.5 percent for GDP growth, 10.2 percent for tertiary education, 2.4 percent for HSPS, and 4.3 percent for population growth. Excluding the two exercises that eliminate knowledge spillovers or replace employment weights with equal sector weights, the largest absolute changes are 15.1, 17.5, 27.3, 16.4, and 13.7 percent, respectively.

Knowledge spillovers account for the largest changes. Eliminating diffusion reduces the GDP-level and GDP-growth slopes by 30.4 and 33.8 percent, respectively, but attenuates the tertiary-education slope by 77.9 percent and the HSPS slope by 56.6 percent. The magnitude of the negative population-growth slope falls by 34.3 percent. Weighting sectoral knowledge creation equally, rather than by employment,

has smaller but still substantial effects: the education and HSPS slopes decline by 38.5 and 24.0 percent, compared with 16.7 and 13.4 percent for GDP levels and growth. Thus knowledge diffusion is particularly important for the education and organizational-formalization relationships. The persistence of the signs when diffusion is removed also shows that these relationships are not generated by knowledge spillovers alone.

Parameters governing firm scale affect magnitudes without overturning the results. Varying span of control changes the GDP-level slope by between -10.4 and 15.1 percent and the tertiary-education slope by between -12.7 and 27.2 percent. Reducing firm-ability dispersion raises the GDP-growth slope by 17.5 percent and the education slope by 27.3 percent. Routine repetition matters more for growth and HSPS than for GDP levels: a 25 percent reduction raises the growth slope by 16.3 percent while lowering the HSPS slope by 16.4 percent. Initial knowledge also matters primarily for education. Setting $K_0 = 0.10$ reduces the tertiary-education slope by 24.4 percent, compared with 5.0 percent for GDP levels and 3.1 percent for HSPS. These responses indicate that the quantitative strength of the education channel is more sensitive than the GDP-level relationship to several features of firms' organizational environment.

The education-environment complementarity parameter has moderate effects. The benchmark sets $\tau_e^s = 0.5$ in $\tau_e(S) = 1 - S(1 - \tau_e^s)$; lower values therefore imply a stronger reduction in education costs as the skilled share rises. Setting $\tau_e^s = 0.25$ reduces the tertiary-education slope by 11.6 percent, whereas setting $\tau_e^s = 0.75$ raises it by 8.8 percent. The corresponding changes in the magnitude of the population-growth slope are -9.1 and 8.3 percent. GDP-level, GDP-growth, and HSPS slopes change by no more than 3.5 percent in either specification. The benchmark complementarity assumption therefore affects the strength of the education and demographic relationships, but has little effect on the income and HSPS relationships over this range.

Other assumptions have little quantitative influence. Varying adaptation costs by 25 percent changes every slope by less than 4 percent, and the GDP-level slope by less than 1 percent. Changing producer-service substitution to either 0.90 or 1.50 changes no slope by more than 1.3 percent. Finally, replacing equation (80) with an exponential-hazard law having the same first-order flow changes every slope by less than 0.7 percent. The results are therefore insensitive to the hard cap in the benchmark accumulation equation. Within the ranges considered, adaptation costs have modest effects, while producer-service substitution and the choice between these two accumulation laws are quantitatively negligible.

The more detailed diffusion exercise varies δ_K from 0.01 to 0.15 in increments of 0.01. Table 33 reports slopes associated with a 10-percentage-point increase in peak manufacturing employment. The responses are not monotonic: several slopes reach their largest values near the benchmark and decline on either side. Nevertheless, the main sensitivity is to very slow diffusion, particularly for the skilled share and HSPS.

Within the empirically motivated interval $[0.04, 0.14]$, the GDP-level slope ranges from 13.8 percent below to 0.6 percent above its value at the $\delta_K = 0.09$ benchmark; the GDP-growth slope ranges from 6.9 percent below to 5.0 percent above. The corresponding ranges are -27.8 to 4.1 percent for tertiary education, -22.4 to 1.6 percent for HSPS, and -16.2 to 0.6 percent for the magnitude of the population-growth slope. Extending the range to $[0.01, 0.15]$ increases the largest attenuations to 28.2 percent for GDP levels, 30.1 percent for GDP growth, 73.7 percent for tertiary education, 49.8 percent for HSPS, and 32.8 percent for population growth. Every sign is preserved. As in the broader sensitivity analysis, slow knowledge diffusion weakens the education and organizational-formalization relationships more sharply than the GDP relationship.

Table 33: Cross-country slopes by the rate of knowledge diffusion

δ_K	Log GDPpc	GDPpc growth	College share	HSPS share
0.01	0.1312	0.00112	0.0119	0.0333
0.02 [†]	0.1230	0.00105	0.0081	0.0291
0.03 [†]	0.1322	0.00128	0.0164	0.0364
0.04 [†]	0.1477	0.00139	0.0222	0.0450
0.05 [†]	0.1604	0.00151	0.0281	0.0520
0.06	0.1676	0.00157	0.0308	0.0558
0.07	0.1677	0.00155	0.0308	0.0564
0.08	0.1725	0.00156	0.0320	0.0588
0.09	0.1714	0.00150	0.0307	0.0579
0.10	0.1708	0.00148	0.0299	0.0574
0.11	0.1693	0.00147	0.0303	0.0569
0.12	0.1682	0.00145	0.0297	0.0564
0.13	0.1667	0.00143	0.0287	0.0558
0.14	0.1652	0.00140	0.0278	0.0553
0.15	0.1636	0.00137	0.0265	0.0541

Notes: Each entry is the OLS slope associated with a 10-percentage-point increase in the endogenous peak manufacturing employment share across five manufacturing-demand economies. All other parameters remain at their calibrated values. GDP per capita is measured in 2019; growth and education and HSPS shares refer to 2010–2019. Growth rates and shares are expressed as fractions. The population-growth slope, omitted from the table, is negative at every value of δ_K . [†]In one economy in each marked row, the discrete formalization decision continues to oscillate beyond the numerical tolerance, although the residuals for knowledge, employment, and establishment counts are below 8.5×10^{-4} . The reported values retain the same solution procedure across specifications. As an additional check, using a neighboring economy’s solution as the starting point at $\delta_K = 0.02$ changes the four reported slopes to (0.1188, 0.00105, 0.0082, 0.0291) and preserves all five signs. The precise magnitudes in the marked rows should therefore be interpreted with this numerical qualification.

E.7 Interpreting the manufacturing demand counterfactual

The counterfactuals in the main text vary the manufacturing demand shifter w_m to generate economies with different peak manufacturing employment shares. This is a reduced-form way to change the relative pull of manufacturing during the transition while holding technology, fertility and education preferences, and production primitives fixed.

This counterfactual has an equivalent price-wedge representation. Final demand over agriculture, manufacturing, and consumer services implies relative demand terms of the form

$$K_a = \left(\frac{p_a}{p_m} \right)^{1-\sigma} Y^{(1-\sigma)(k_a-1)} \frac{w_a}{w_m^{k_a}}, \quad K_c = \left(\frac{p_c}{p_m} \right)^{1-\sigma} Y^{(1-\sigma)(k_c-1)} \frac{w_c}{w_m^{k_c}},$$

where $k_j \equiv \varepsilon_j / \varepsilon_m$. Changing the manufacturing demand weight from w_m to $f_m w_m$ therefore changes the relative demand terms to

$$K_a(f_m) = K_a(1) f_m^{-k_a}, \quad K_c(f_m) = K_c(1) f_m^{-k_c}.$$

Now consider an economy with the original w_m , but in which households face effective final-good prices

$$\tilde{p}_{a,t} = (1 + \tau_a) p_{a,t}, \quad \tilde{p}_{m,t} = p_{m,t}, \quad \tilde{p}_{c,t} = (1 + \tau_c) p_{c,t}.$$

The wedge economy generates the same relative demand schedules as the $f_m w_m$ economy if

$$(1 + \tau_a)^{1-\sigma} = f_m^{-k_a}, \quad (1 + \tau_c)^{1-\sigma} = f_m^{-k_c}.$$

Equivalently,

$$\boxed{1 + \tau_a = f_m^{-k_a/(1-\sigma)}, \quad 1 + \tau_c = f_m^{-k_c/(1-\sigma)}}. \quad (85)$$

Thus, the w_m counterfactual admits two equivalent representations: a manufacturing-demand shifter and a system of effective final-good prices. The equivalence applies to household relative-demand schedules, conditional on income and producer prices. It does not require interpreting w_m as a primitive preference parameter. Instead, w_m provides a parsimonious reduced-form representation of forces that alter the relative demand for manufacturing during structural transformation. Table 34 reports the corresponding price wedges for the selected counterfactual economies in Figure 7.

The asymmetry between the two wedges follows directly from the calibrated demand system. For agriculture, the relevant exponent is

$$-\frac{k_a}{1-\sigma} = -0.0143,$$

so even large changes in w_m imply small movements in the effective agricultural price. For consumer services, the corresponding exponent is

$$-\frac{k_c}{1-\sigma} = -1.4435.$$

The service-price mapping is therefore substantially more responsive and becomes nonlinear as f_m falls. For example, the $w_m = 1.16$ counterfactual implies $1 + \tau_c = 3.03$: under the normalization in the table, the effective consumer-service price is approximately three times its benchmark level. At $w_m = 1.80$, the corresponding multiplier is 1.61.

Table 34: Manufacturing-demand counterfactuals and equivalent final-consumption wedges

w_m	f_m	Peak manufacturing employment (%)	Agriculture wedge τ_a (%)	Consumer-services wedge τ_c (%)
0.60 [†]	0.240	11.19	2.06	684.67
1.16	0.464	15.92	1.10	202.97
1.80	0.720	20.62	0.47	60.67
2.50 (benchmark)	1.000	24.93	0.00	0.00
2.56	1.024	25.27	-0.03	-3.37
3.49	1.396	29.88	-0.48	-38.22
4.69	1.876	34.56	-0.89	-59.67
6.10	2.440	38.80	-1.27	-72.41

Notes: The calibrated benchmark is $w_m^* = 2.5$, and $f_m \equiv w_m/w_m^*$. The wedges are computed from $1 + \tau_j = f_m^{-k_j/(1-\sigma)}$, using $\sigma = 0.30$, $k_a = 0.0100$, and $k_c = 1.0105$. The normalization holds the effective manufacturing price fixed. Positive values raise the effective final-consumption price relative to manufacturing; negative values lower it. Peak manufacturing employment is endogenous. The wedges are demand-equivalent shadow prices, not statutory tax rates.

These magnitudes should be interpreted as broad shadow-price equivalents. They summarize any combination of market access, distribution costs, markups, regulation, product availability, relative quality, and other pecuniary or nonpecuniary forces that generate the same relative final-demand schedules. The counterfactuals do not introduce government revenue or tax rebates, and the implied wedges play no independent role in the calibration. The economic variation is disciplined by the endogenous manufacturing peaks; the wedge representation provides a transparent interpretation of the reduced-form demand variation used to generate them.

E.8 Response of the transition to policy

In Figure 14, I compare the paths of skills, GDP-per-capita, welfare, formalization and public organizational knowledge under the different policy instruments considered in Section 7. The figure shows that policy interventions can have long-lasting effects when combined with the removal of productivity-correlated distortions, but otherwise remain mostly transitory level-shifts rather than placing the economy on a different long-run growth path.




Figure 14: Unexpected policy withdrawal: scale and organizational absorption. All temporary programs operate from 2026 through 2045 and are withdrawn unexpectedly at the start of 2046. The shaded region denotes the policy period and the vertical dotted line denotes withdrawal. The left column retains the LMIC distortion block; the right column includes the permanent distortion reform. Each path is a change relative to the no-temporary-policy path in the same distortion environment. The continuation inherits the full firm distribution, installed routines, private localized knowledge, and public organizational-knowledge stock. The figure uses the 81-point task grid and displays outcomes through 2070.